



# Troubleshooting and Debugging

Omar Ismail

Solutions Developer, Google  
Cloud





---

# Agenda

Course Intro

Monitoring

Logging and Error Reporting

**Troubleshooting and Debugging**


Performance

Testing and CI/CD

Reliability

Flex Templates


Course Summary



---


# Troubleshooting and Debugging

## Agenda



---

# Troubleshooting workflow




Troubleshooting  
workflow

- Checking for errors
- Looking for anomalies in the Job Metrics tab

---


# Troubleshooting workflow



Troubleshooting  
workflow

- Checking for errors
- Looking for anomalies in the Job Metrics tab


# Checking for errors




Dataflow





Jobs

[+ CREATE JOB FROM TEMPLATE](#)









Filter jobs



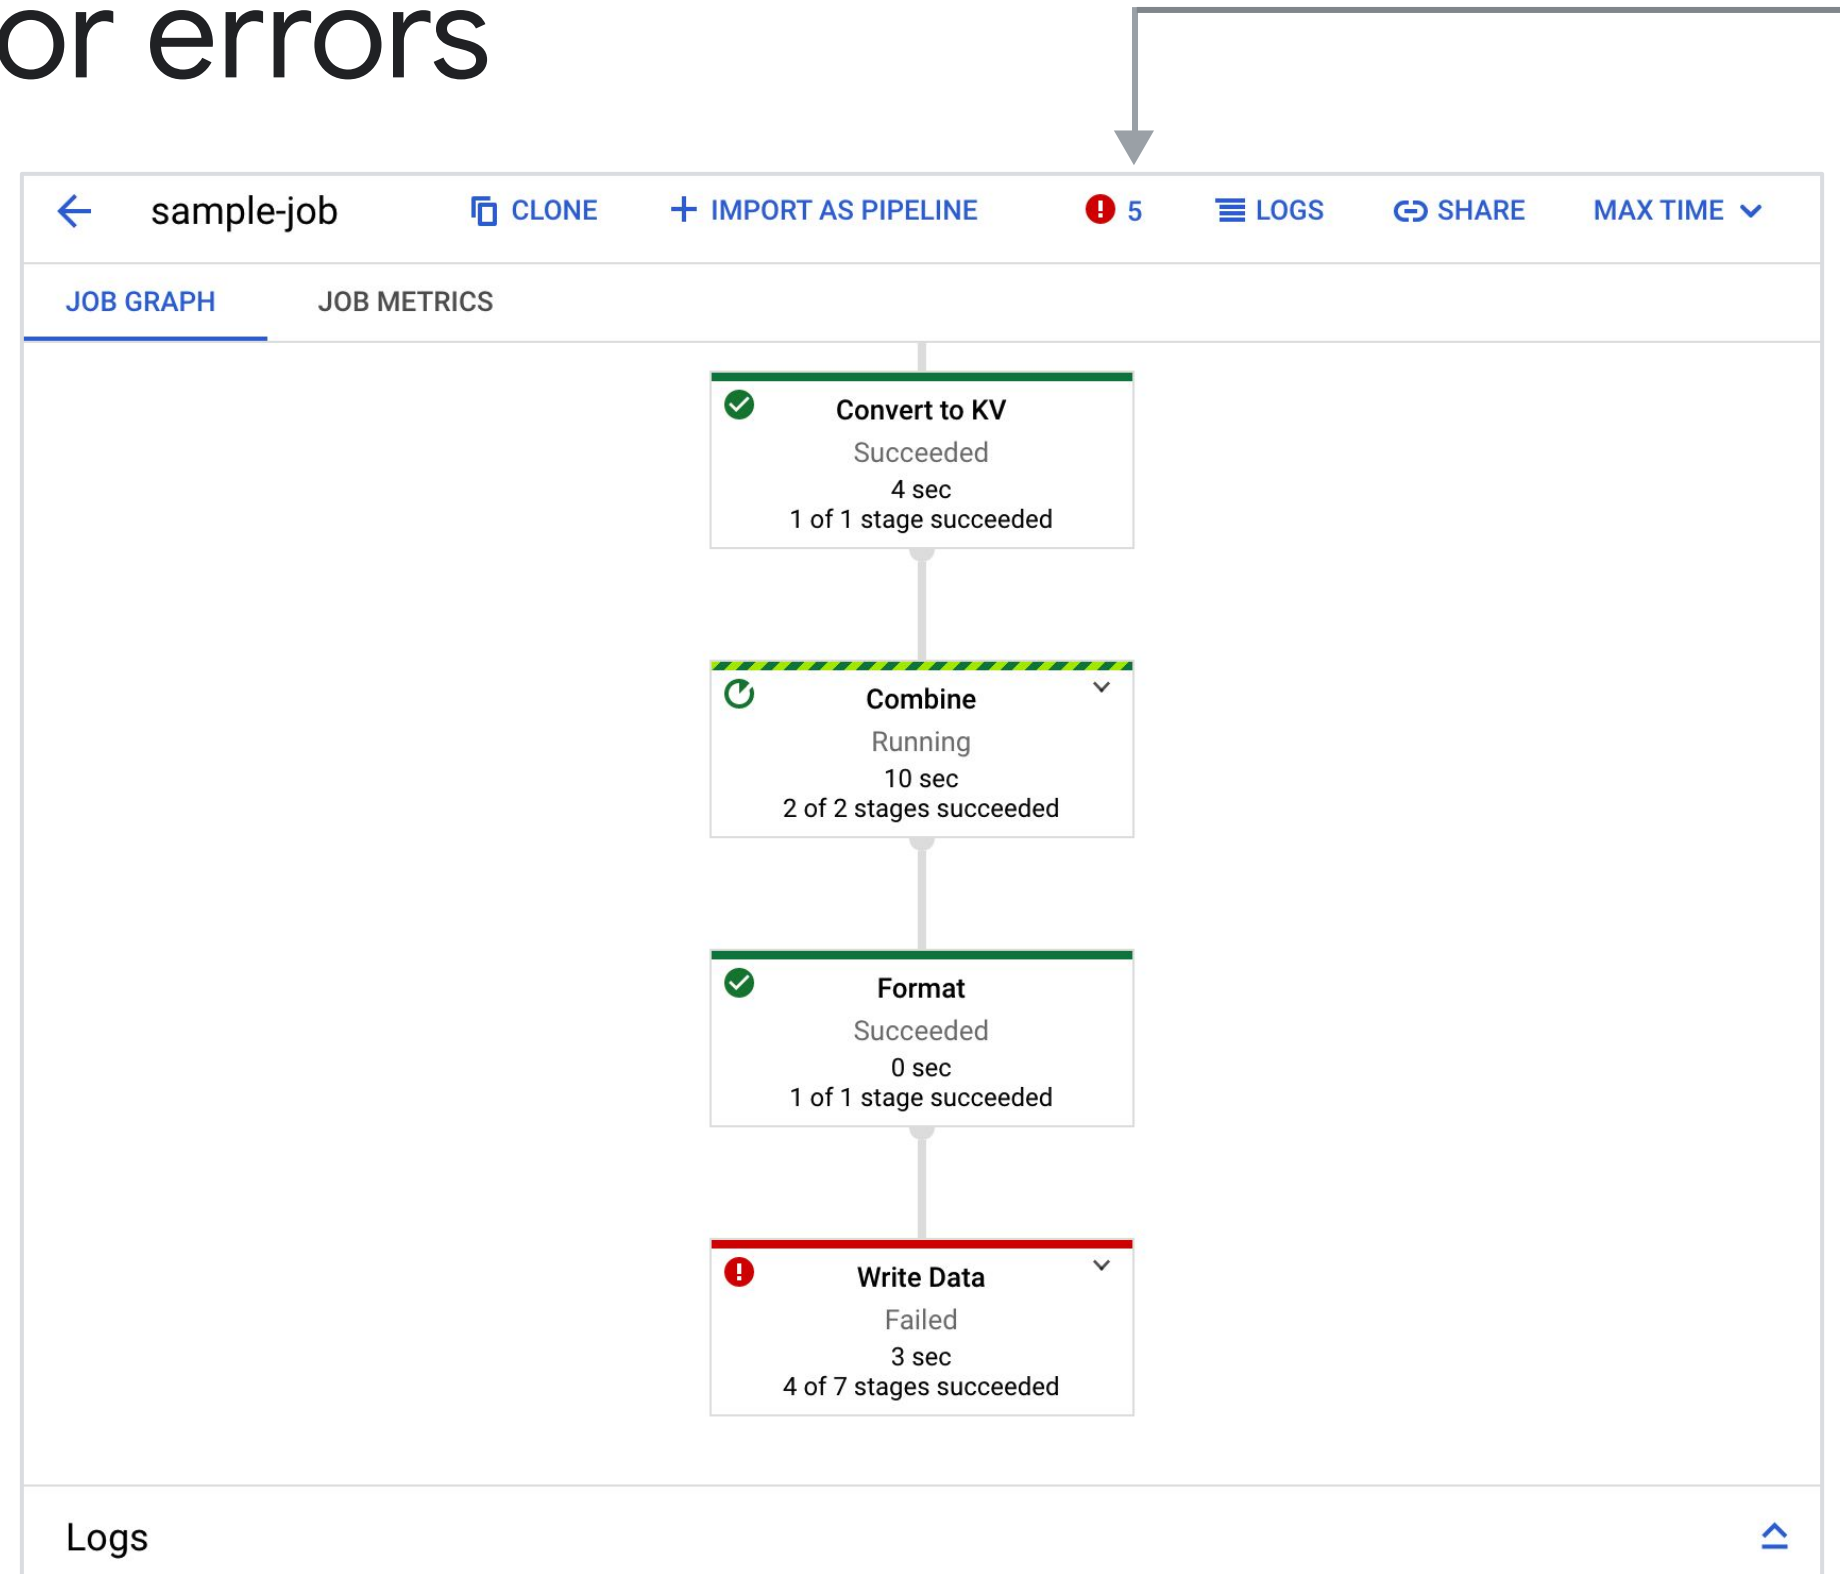
Name	Type	End time	Elapsed time	Start time	Status	SDK version	ID	Region
 wordcount4	Batch	—	7 sec	Nov 21, 2018, 3:22:45 PM	Running	2.7.0	2018-11-21_15_22_45-	us-central1
 wordcount3	Batch	Nov 21, 2018, 1:10:23 PM	3 min 29 sec	Nov 21, 2018, 1:06:54 PM	Succeeded	2.7.0	2018-11-21_13_06_53-	us-central1
 wordcount2	Batch	Nov 21, 2018, 12:57:05 PM	3 min 56 sec	Nov 21, 2018, 12:53:09 PM	Failed	2.7.0	2018-11-21_12_53_08-	us-central1
 wordcount1	Batch	Nov 21, 2018, 12:55:25 PM	3 min 36 sec	Nov 21, 2018, 12:51:49 PM	Succeeded	2.7.0	2018-11-21_12_51_48-	us-central1

# Checking for errors

<div> Dataflow   Jobs <a href="#">+ CREATE JOB FROM TEMPLATE</a></div>								
<div><div> Filter jobs</div><div>?</div></div>								
Name	Type	End time	Elapsed time	Start time	Status	SDK version	ID	Region
 wordcount4	Batch	—	7 sec	Nov 21, 2018, 3:22:45 PM	Running	2.7.0	2018-11-21_15_22_45- XXXXXXXXXXXXXXXXXXXX	us-central1
 wordcount3	Batch	Nov 21, 2018, 1:10:23 PM	3 min 29 sec	Nov 21, 2018, 1:06:54 PM	Succeeded	2.7.0	2018-11-21_13_06_53- XXXXXXXXXXXXXXXXXXXX	us-central1
 wordcount2	Batch	Nov 21, 2018, 12:57:05 PM	3 min 56 sec	Nov 21, 2018, 12:53:09 PM	Failed	2.7.0	2018-11-21_12_53_08- XXXXXXXXXXXXXXXXXXXX	us-central1
 wordcount1	Batch	Nov 21, 2018, 12:55:25 PM	3 min 36 sec	Nov 21, 2018, 12:51:49 PM	Succeeded	2.7.0	2018-11-21_12_51_48- XXXXXXXXXXXXXXXXXXXX	us-central1

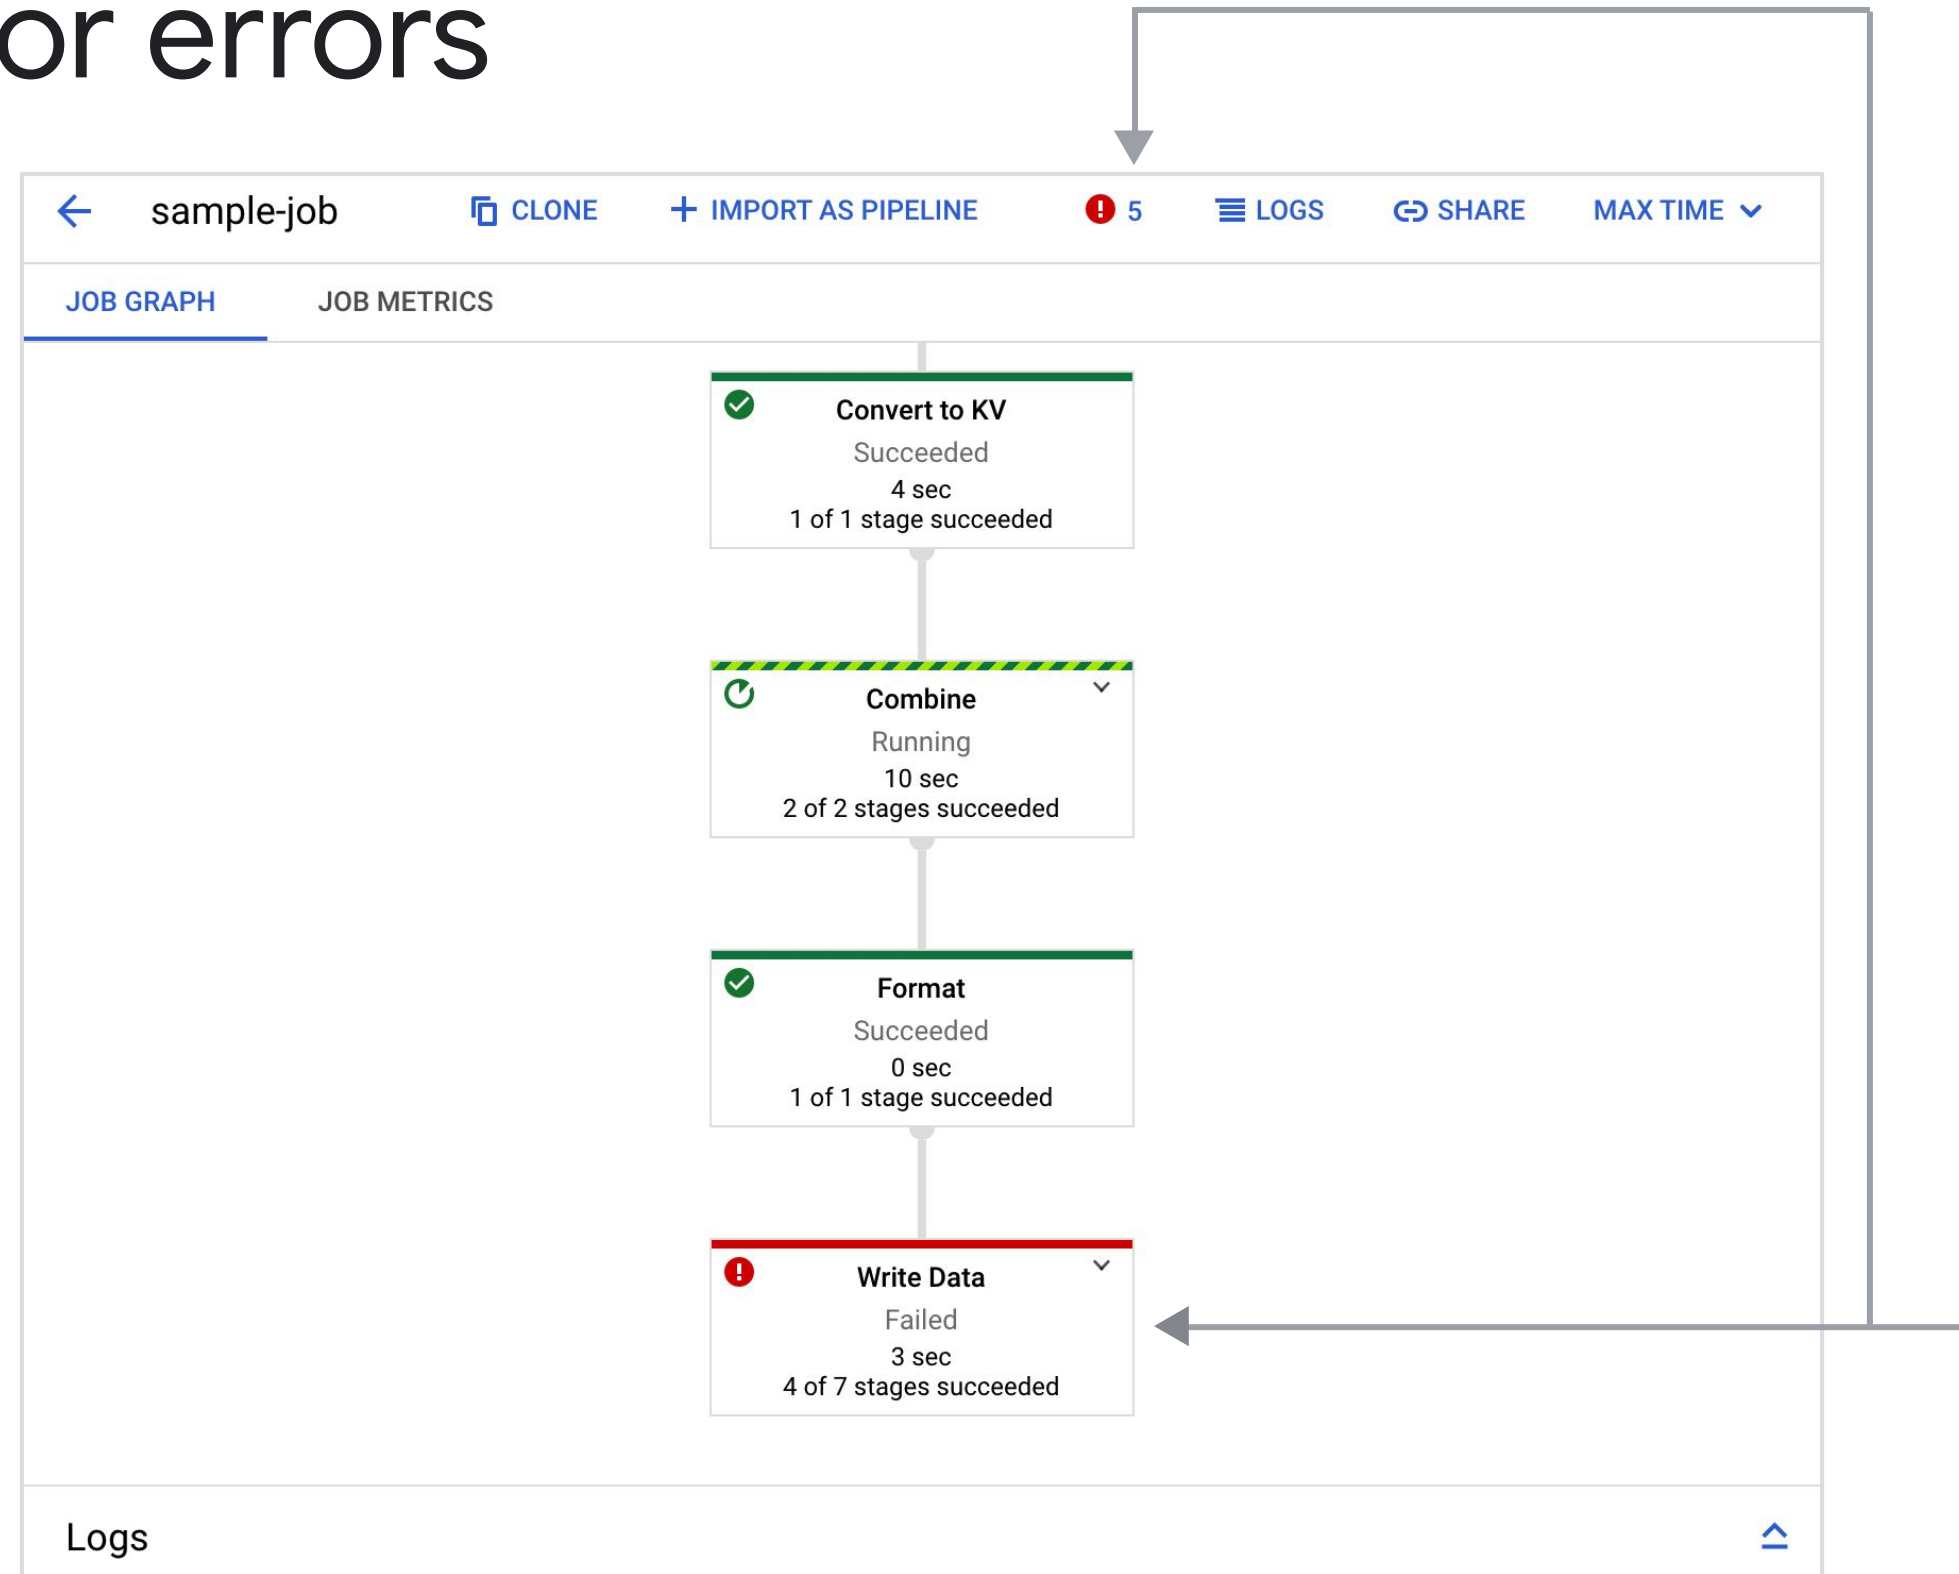


# Checking for errors



Checking  
for errors

# Checking for errors



Checking  
for errors

# Checking for errors

sample-job

CLONE

5

LOGS

SHARE

MAX TIME

JOB GRAPH

JOB METRICS

1 of 1 stage succeeded

Write Data

Failed

3 sec

4 of 7 stages succeeded

Logs

JOB LOGS

WORKER LOGS

DIAGNOSTICS

Showing 50 messages

Info

Filter

Filter logs

2021-02-16 13:48:02.861 CST

Executing operation Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/WaitForDes...

2021-02-16 13:48:05.261 CST

Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2....

2021-02-16 13:48:05.906 CST

Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2....

2021-02-16 13:48:07.603 CST

Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2....

2021-02-16 13:48:10.354 CST

Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2....

2021-02-16 13:48:10.714 CST

Finished operation Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/WaitForDest...

2021-02-16 13:48:10.797 CST

Workflow failed. Causes: S19:Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/W...

2021-02-16 13:48:11.061 CST

Stopping worker pool...

2021-02-16 13:50:15.803 CST

Worker pool stopped.

No newer entries found matching current filter.

# Checking for errors

sample-job

CLONE

5

LOGS

SHARE

MAX TIME

JOB GRAPH

JOB METRICS

1 of 1 stage succeeded

Write Data

Failed

3 sec

4 of 7 stages succeeded

Logs

JOB LOGS

WORKER LOGS

DIAGNOSTICS

Showing 50 messages

Info

Filter

Filter logs

↺

↗

⛶

▶ 2021-02-16 13:48:02.861 CST Executing operation Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/WaitForDes...

▶ 2021-02-16 13:48:05.261 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2.7/...

▶ 2021-02-16 13:48:05.906 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2.7/...

▶ 2021-02-16 13:48:07.603 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2.7/...

▶ 2021-02-16 13:48:10.354 CST Error message from worker: Traceback (most recent call last): File "/usr/local/lib/python2.7/...

▶ 2021-02-16 13:48:10.714 CST Finished operation Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/WaitForDest...

▶ 2021-02-16 13:48:10.797 CST Workflow failed. Causes: S19:Write Data/BigQueryBatchFileLoads/WaitForDestinationLoadJobs/W...

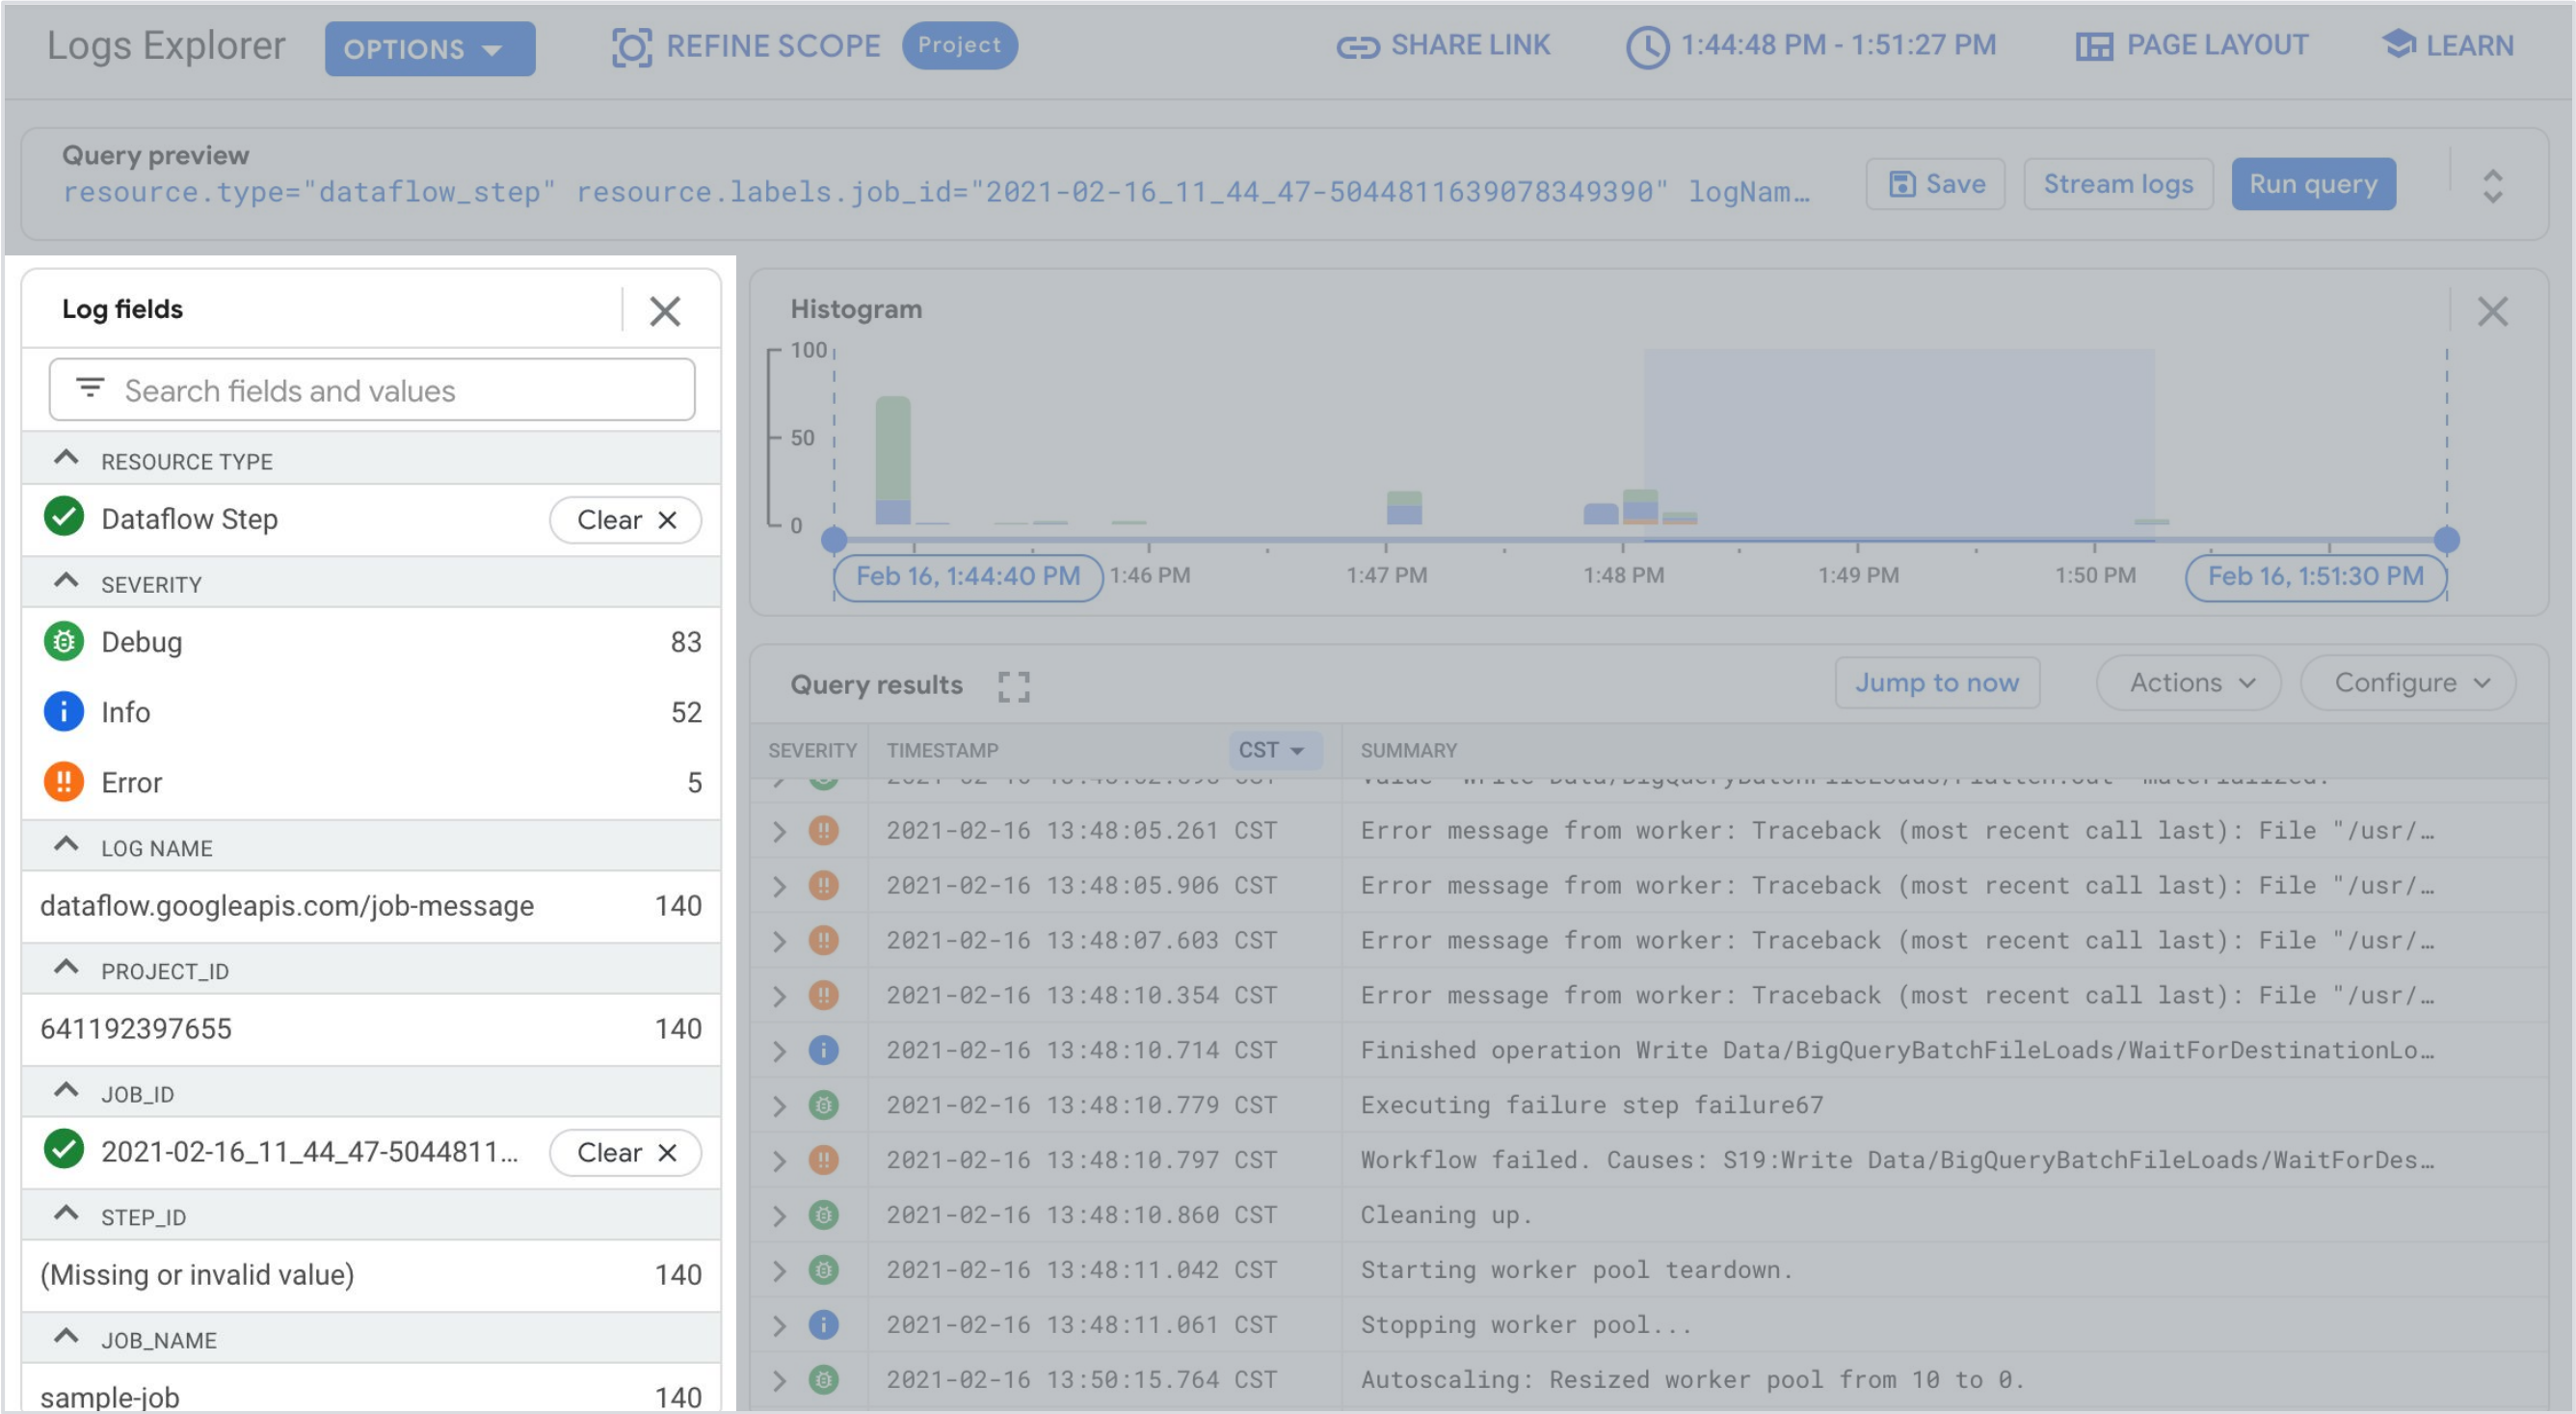
▶ 2021-02-16 13:48:11.061 CST Stopping worker pool...

▶ 2021-02-16 13:50:15.803 CST Worker pool stopped.

▶ No newer entries found matching current filter.


View full logs

# Checking for errors



---

# Troubleshooting workflow




Troubleshooting  
workflow

- Checking for errors
- Looking for anomalies in the Job Metrics tab




# Looking for anomalies in the Job Metrics tab

Streaming pipelines



# Looking for anomalies in the Job Metrics tab


## Streaming pipelines





# Looking for anomalies in the Job Metrics tab

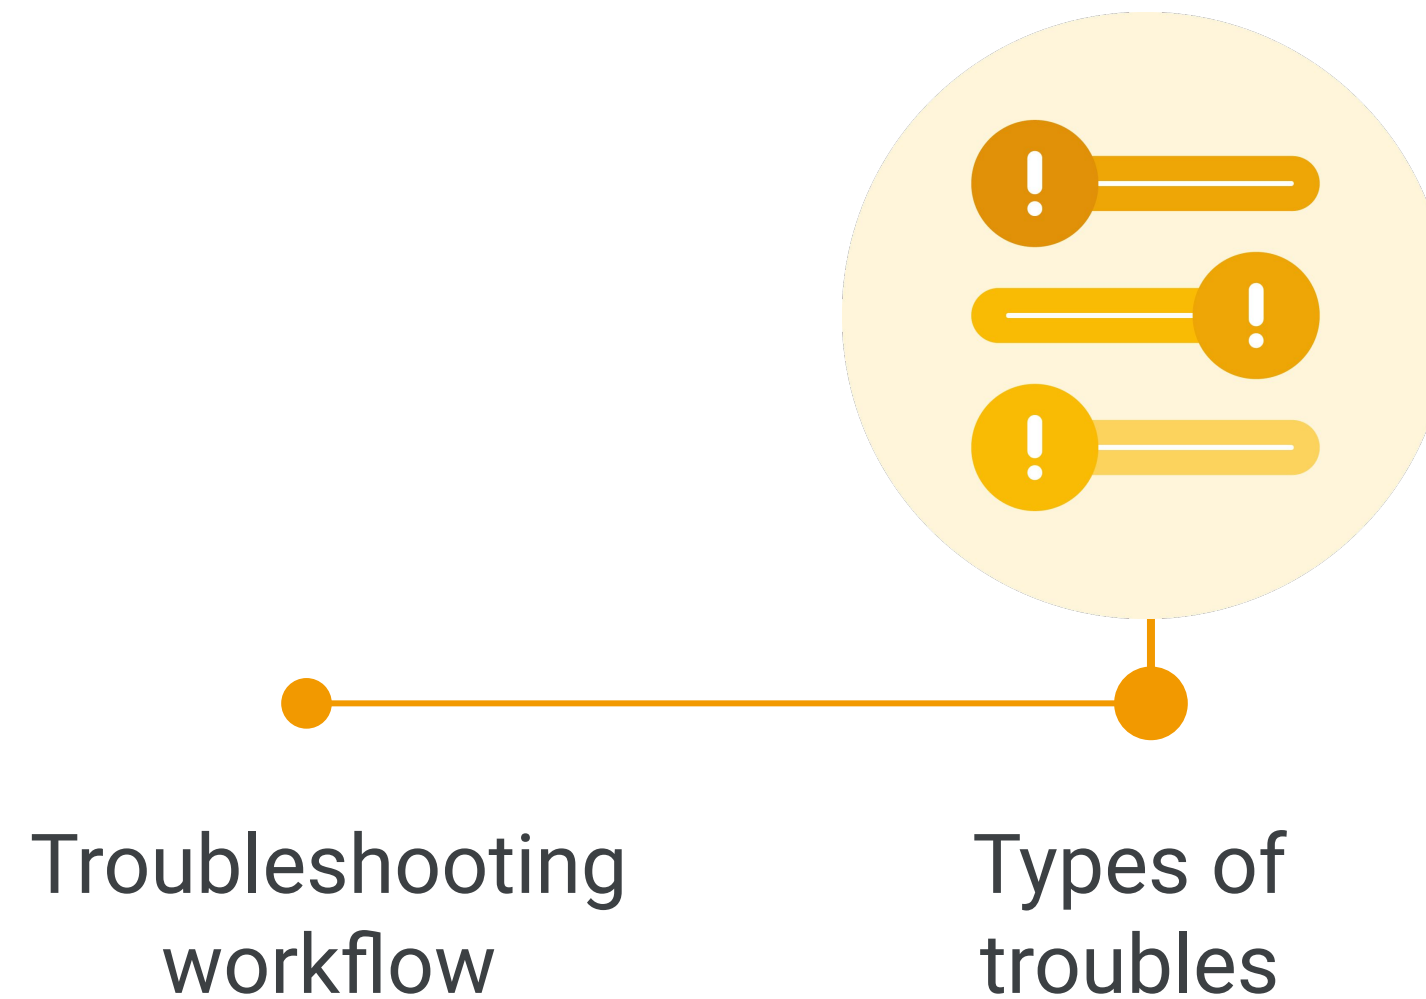
All Dataflow jobs



---


# Troubleshooting and Debugging

## Agenda



---

# Types of troubles




Types of  
troubles

- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

---

# Types of troubles



Types of  
troubles

- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

---

# Failure building the pipeline

These errors occur when Apache Beam is building the pipeline:

- Validating "Beam Model" aspects of the pipeline
- Validating input/output specifications
- You can reproduce it with the direct runner, in a unit test

---


# Failure building the pipeline

```
java.lang.IllegalStateException: GroupByKey cannot be applied to non-bounded PCollection
in the GlobalWindow without a trigger. Use a Window.into or Window.triggering transform
prior to GroupByKey.
```

```
    at org.apache.beam.sdk.transforms.GroupByKey.applicableTo (GroupByKey.java:153)
    at org.apache.beam.sdk.transforms.GroupByKey.expand (GroupByKey.java:185)
    at org.apache.beam.sdk.transforms.GroupByKey.expand (GroupByKey.java:107)
    at org.apache.beam.sdk.Pipeline.applyInternal (Pipeline.java:537)
    at org.apache.beam.sdk.Pipeline.applyTransform (Pipeline.java:471)
    at org.apache.beam.sdk.values.PCollection.apply (PCollection.java:357)
    at ...
```

---

# Types of troubles



Types of  
troubles

- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

---

# Failure starting the pipeline on Dataflow

Once the Dataflow service has received your pipeline's graph, the service will attempt to validate your job:

- Making sure the service can access your job's associated Cloud Storage buckets for file staging and temporary output.





---

# Failure starting the pipeline on Dataflow

Once the Dataflow service has received your pipeline's graph, the service will attempt to validate your job:

- Making sure the service can access your job's associated Cloud Storage buckets for file staging and temporary output.
- Checking for the required permissions in your Google Cloud project.



---

# Failure starting the pipeline on Dataflow

Once the Dataflow service has received your pipeline's graph, the service will attempt to validate your job:

- Making sure the service can access your job's associated Cloud Storage buckets for file staging and temporary output.
- Checking for the required permissions in your Google Cloud project.
- Making sure the service can access input and output sources, such as files.




---

# Failure starting the pipeline on Dataflow

```
apitools.base.py.exceptions.HttpForbiddenError: HttpError accessing
<https://dataflow.googleapis.com/v1b3/projects/xyz/locations/us-central1/jobs?alt=json>:
response: <{'status': '403', 'content-length': '288', 'x-xss-protection': '0',
'x-content-type-options': 'nosniff', 'transfer-encoding': 'chunked', 'vary': 'Origin,
X-Origin, Referer', 'server': 'ESF', '-content-encoding': 'gzip', 'cache-control': 'private',
'date': 'Tue, 10 Dec 2019 17:56:00 GMT', 'x-frame-options': 'SAMEORIGIN', 'alt-svc':
'quic=":443"; ma=2592000; v="46,43",h3-Q050=":443"; ma=2592000,h3-Q049=":443";
ma=2592000,h3-Q048=":443"; ma=2592000,h3-Q046=":443"; ma=2592000,h3-Q043=":443"; ma=2592000',
'content-type': 'application/json; charset=UTF-8'}>, content <{
  "error": {
    "code": 403,
    "message": "(43abf11c2446750): Could not create workflow; user does not have write
access to project: xyz Causes: (43abf11c244683b): Permission 'dataflow.jobs.create'
denied on project: 'xyz'",
    "status": "PERMISSION_DENIED"
  }
}
```

---

# Types of troubles



Types of  
troubles

- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

# Failure during pipeline execution

## Logs

JOB LOGSWORKER LOGS⚠️ DIAGNOSTICS

Showing 78 messagesSeverityErrorFilterFilter logs

!!	2021-03-08 11:30:53.696 CST	Uncaught exception:
!!	2021-03-08 11:30:53.696 CST	Uncaught exception:
!!	2021-03-08 11:30:53.697 CST	Uncaught exception:
!!	2021-03-08 11:30:53.698 CST	Uncaught exception:
!!	2021-03-08 11:30:53.698 CST	Uncaught exception:
!!	2021-03-08 11:30:53.701 CST	Uncaught exception:
!!	2021-03-08 11:30:53.704 CST	Uncaught exception:
!!	2021-03-08 11:30:53.707 CST	Uncaught exception:



# Failure during pipeline execution

Log fields

Search fields and values

RESOURCE TYPE

Dataflow Step

Clear X

SEVERITY

Error

Clear X

LOG NAME

dataflow.googleapis.com/worker

Clear X

PROJECT ID

JOB ID

STEP ID

(Missing or invalid value) 32

JOB NAME

REGION

us-central1 32

Query results

Jump to now

Actions

Configure

SEVERITY	TIMESTAMP	CST	SUMMARY
> !!	2021-03-08 11:22:40.458	CST	"Uncaught exception: "
✓ !!	2021-03-08 11:22:40.458	CST	"Uncaught exception: "

{

insertId: "776189048441790998:16632:0:753343"

jsonPayload: {

exception:

"java.lang.RuntimeException: org.apache.beam.sdk.util.UserCodeException: java.lang.OutOfMemoryError: unable to create native thread: possibly out of memory or process/resource limits reached

at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowsParDoFn\$1.output(GroupAlsoByWindowsParDoFn.java:184)

at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowFnRunner\$1.outputWindowedValue(GroupAlsoByWindowFnRunner.java:107)

at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.ReduceFnRunner.lambda\$onTrigger\$1(ReduceFnRunner.java:1057)

at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.ReduceFnContextFactory\$OnTriggerContextImpl.output(ReduceFnContextFactory.java:438)

at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.SystemReduceFn.onTrigger(SystemReduceFn.java:125)

at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.ReduceFnRunner.onTrigger(ReduceFnRunner.java:1060)

at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.ReduceFnRunner.onTimers(ReduceFnRunner.java:768)

at org.apache.beam.runners.dataflow.worker.StreamingGroupAlsoByWindowViaWindowSetFn.processElement(StreamingGroupAlsoByWindowViaWindowSetFn.java:95)

at org.apache.beam.runners.dataflow.worker.StreamingGroupAlsoByWindowViaWindowSetFn.processElement(StreamingGroupAlsoByWindowViaWindowSetFn.java:42)

at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowFnRunner.invokeProcessElement(GroupAlsoByWindowFnRunner.java:120)

at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowFnRunner.processElement(GroupAlsoByWindowFnRunner.java:73)

at org.apache.beam.runners.dataflow.worker.repackaged.org.apache.beam.runners.core.LateDataDroppingDoFnRunner.processElement(LateDataDroppingDoFnRunner.java:80)

at org.apache.beam.runners.dataflow.worker.GroupAlsoByWindowsParDoFn.processElement(GroupAlsoByWindowsParDoFn.java:134)

at org.apache.beam.runners.dataflow.worker.util.common.worker.ParDoOperation.process(ParDoOperation.java:44)

at org.apache.beam.runners.dataflow.worker.util.common.worker.OutputReceiver.process(OutputReceiver.java:49)

at org.apache.beam.runners.dataflow.worker.util.common.worker.ReadOperation.runReadLoop(ReadOperation.java:201)

at org.apache.beam.runners.dataflow.worker.util.common.worker.ReadOperation.start(ReadOperation.java:159)

at org.apache.beam.runners.dataflow.worker.util.common.worker.MapTaskExecutor.execute(MapTaskExecutor.java:77)

at org.apache.beam.runners.dataflow.worker.StreamingDataflowWorker.process(StreamingDataflowWorker.java:1363)

at org.apache.beam.runners.dataflow.worker.StreamingDataflowWorker.access\$1100(StreamingDataflowWorker.java:153)

at org.apache.beam.runners.dataflow.worker.StreamingDataflowWorker\$7.run(StreamingDataflowWorker.java:1086)

at java.base/java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1128)

at java.base/java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:628)

at java.base/java.lang.Thread.run(Thread.java:834)

Caused by: org.apache.beam.sdk.util.UserCodeException: java.lang.OutOfMemoryError: unable to create native thread: possibly out of memory or process/resource limits reached

Hide log summary

Collapse nested fields

Copy to clipboard

Copy link

---

# Failure during pipeline execution

```
class FilterMessagesFn(beam.DoFn):
    BAD_MESSAGE_TAG = 'bad_message'
    GOOD_MESSAGE_TAG = 'good_message'

    def process(self, element, window=beam.DoFn.WindowParam):
        try:
            data = element.decode()
            # tag the elements accordingly
            if 'bad' in data:
                yield pvalue.TaggedOutput(self.BAD_MESSAGE_TAG, element)
            else:
                yield pvalue.TaggedOutput(self.GOOD_MESSAGE_TAG, element)

            # handle any exceptions in the processing
        except Exception as exp:
            logging.getLogger().warning(exp)
            yield pvalue.TaggedOutput(self.BAD_MESSAGE_TAG, element)
```

---

# Failure during pipeline execution

To track failing elements:

- Log the failing elements and check the output using Cloud Logging.





---

# Failure during pipeline execution

To track failing elements:

- Log the failing elements and check the output using Cloud Logging.
- Check the Dataflow worker and worker startup logs for warnings or errors related to work item failures.



---

# Failure during pipeline execution

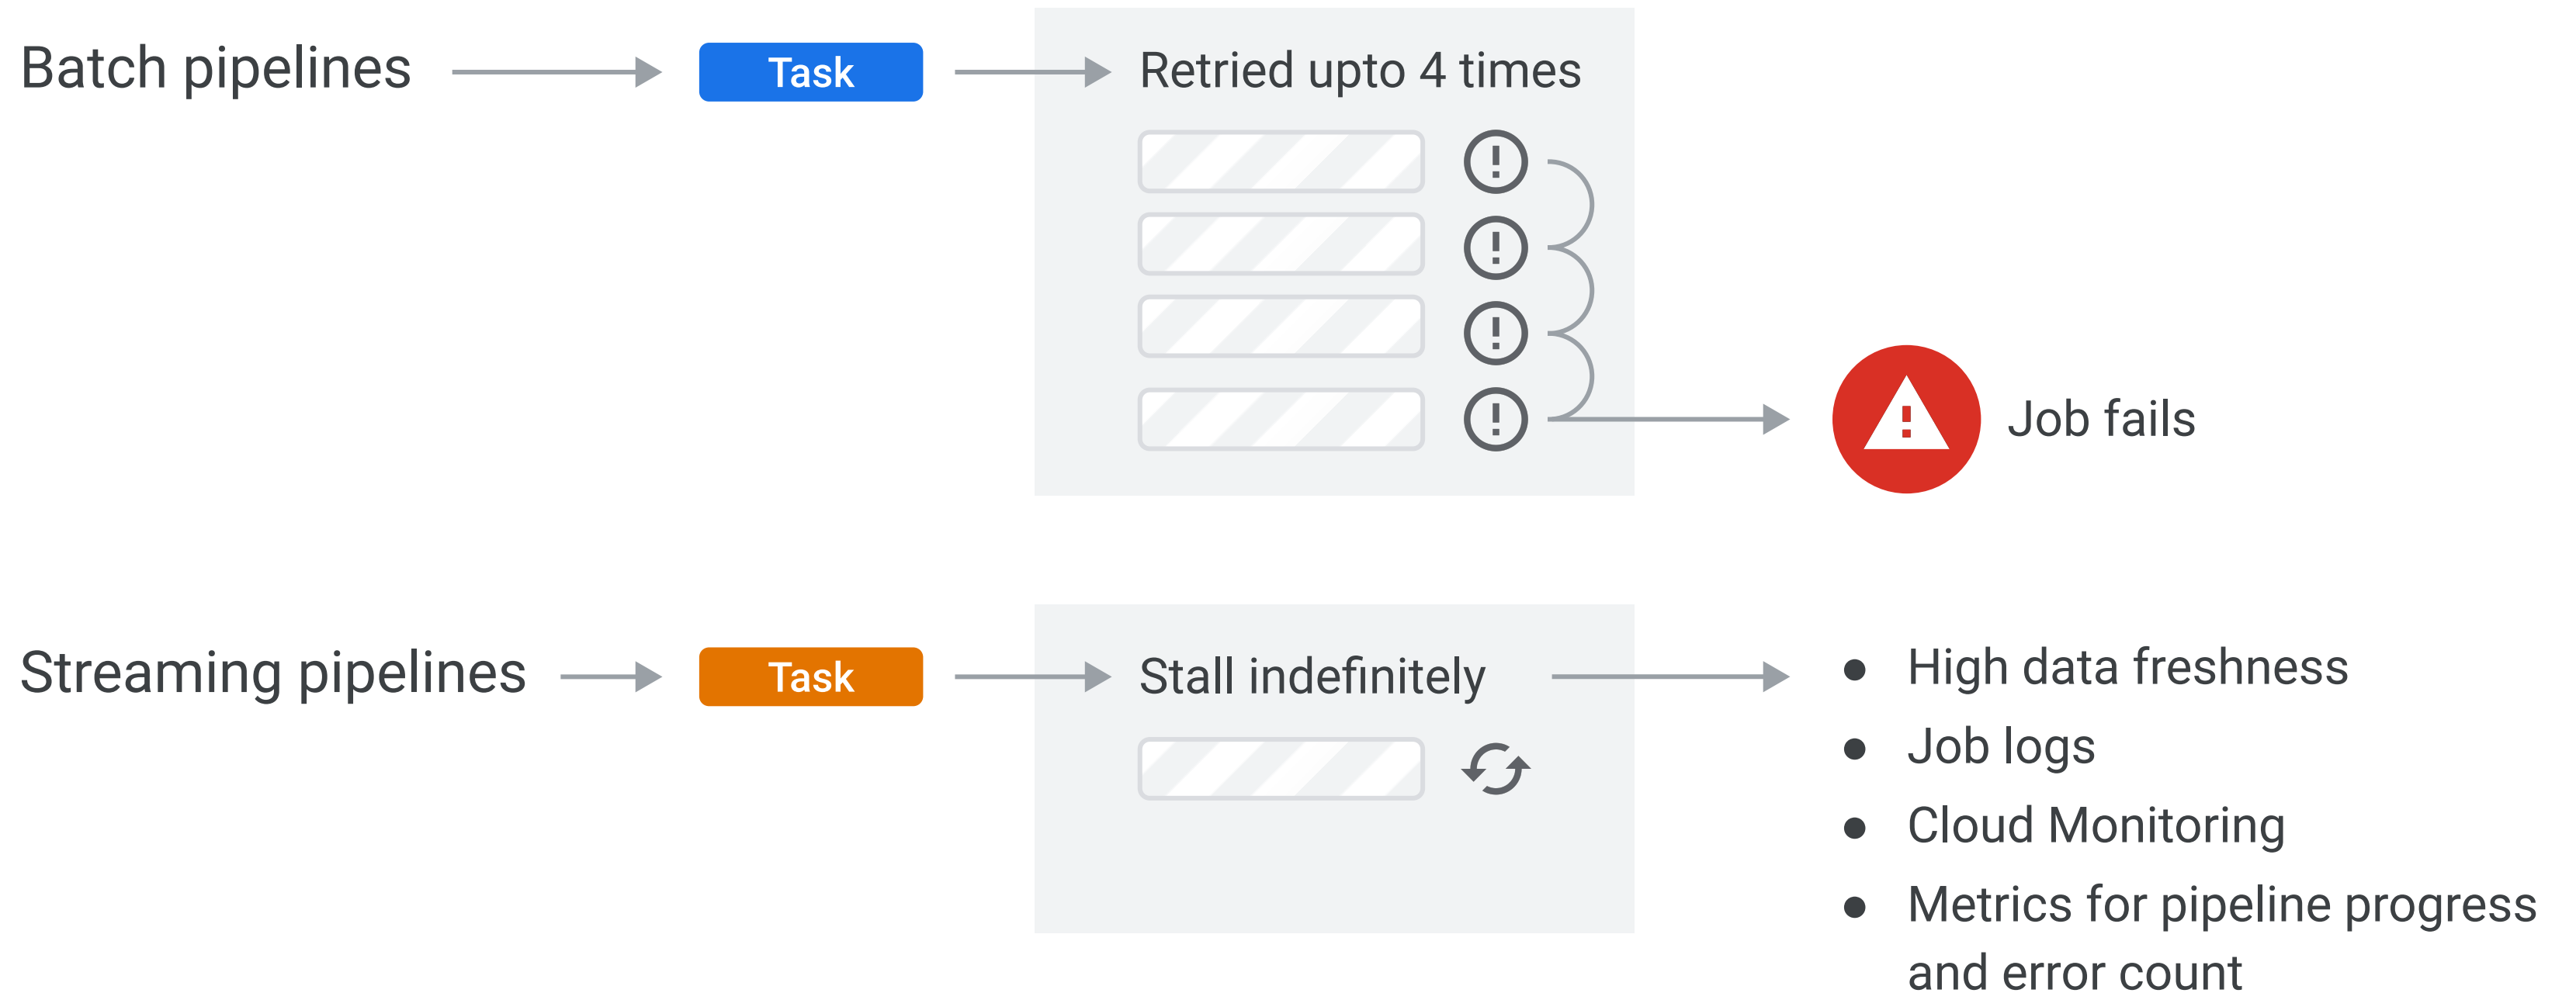
To track failing elements:

- Log the failing elements and check the output using Cloud Logging.
- Check the Dataflow worker and worker startup logs for warnings or errors related to work item failures.
- Set your ParDo to write the failing elements to an additional output for later inspection.




# Failure during pipeline execution

## Batch vs streaming



---

# Types of troubles



Types of  
troubles


- Failure building the pipeline
- Failure starting the pipeline on Dataflow
- Failure during pipeline execution
- Performance problems

---

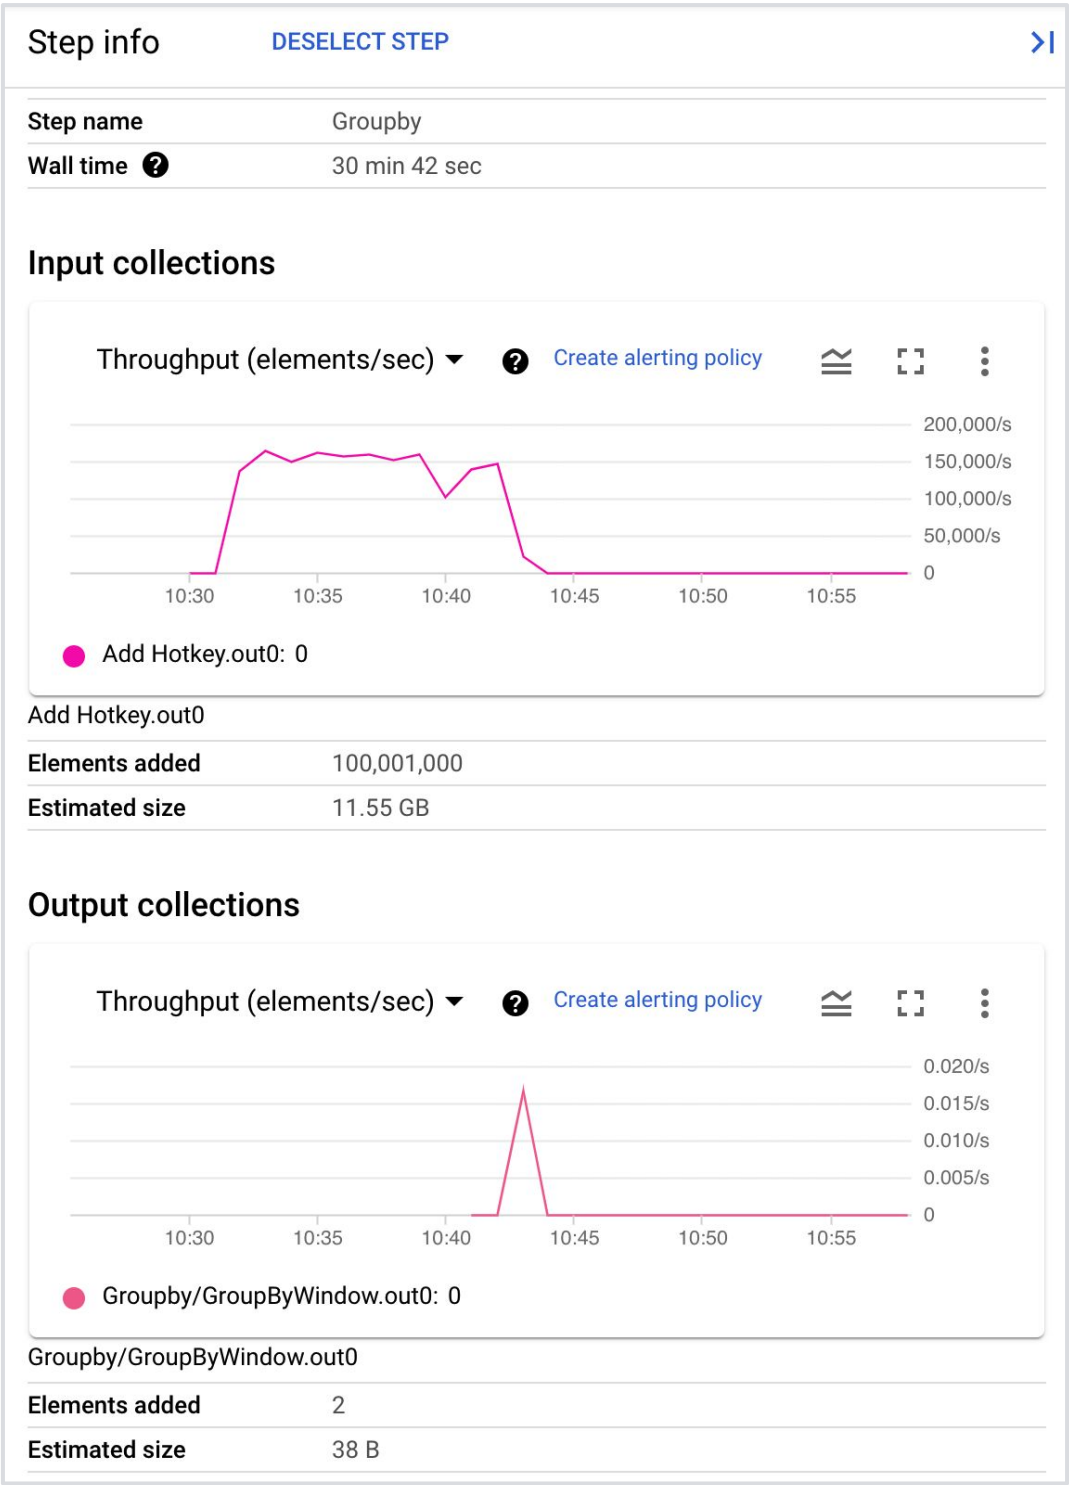
# Performance problems

Factors that can influence the performance of a pipeline:

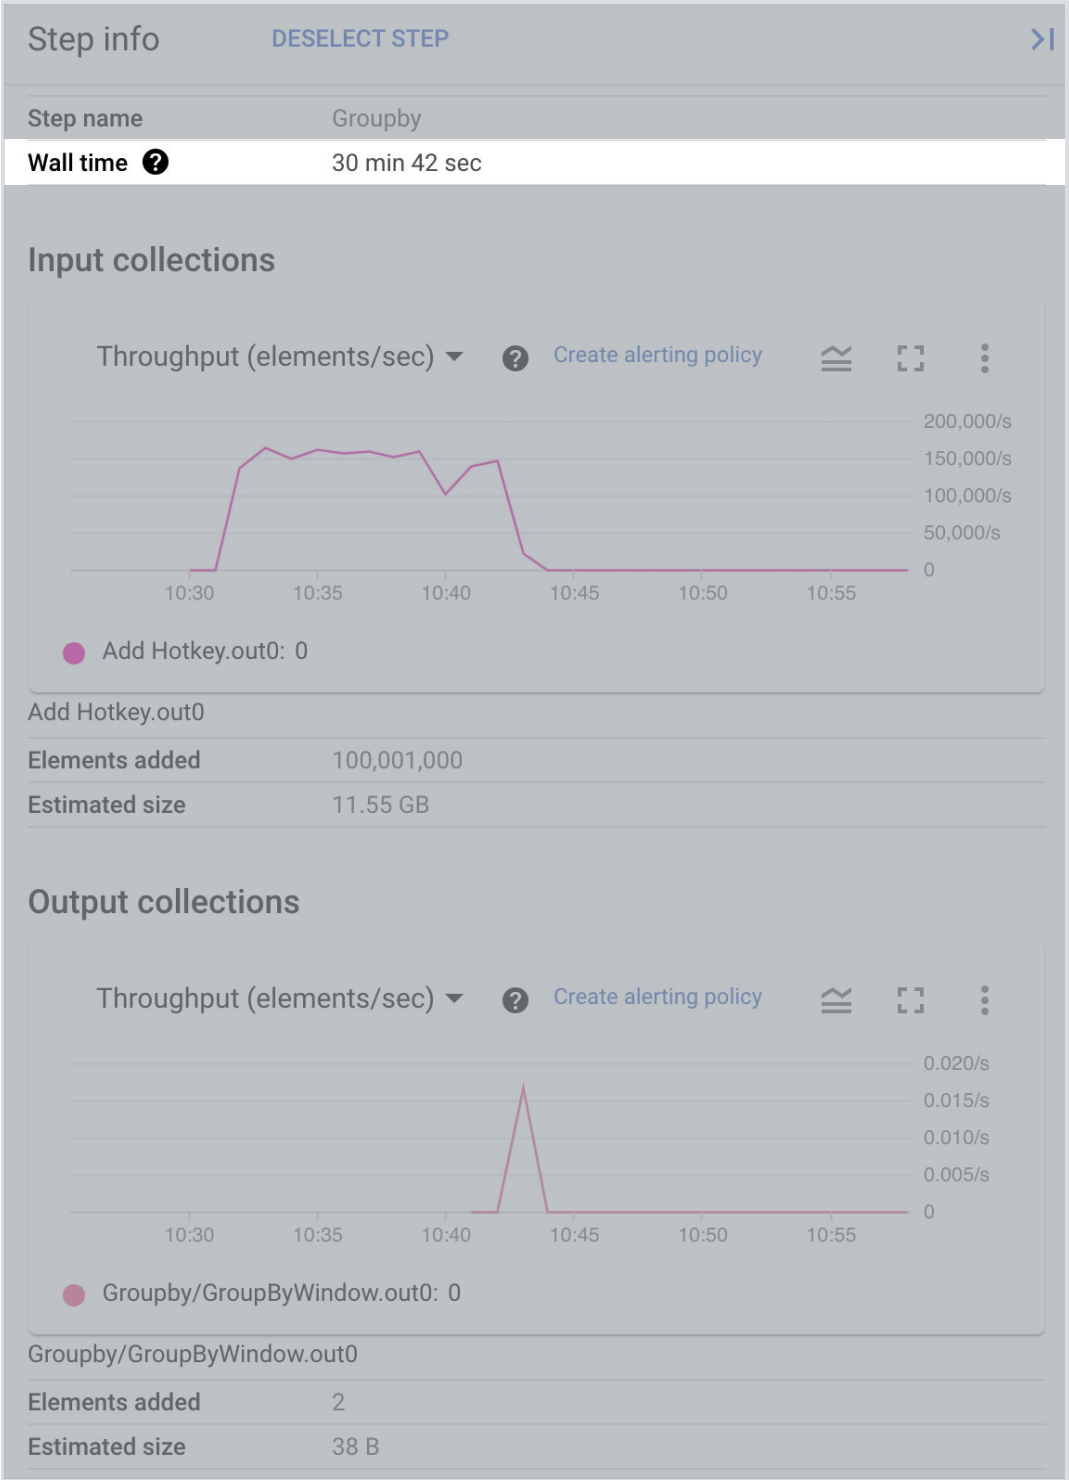
- Pipeline design
- Data shape
- Interactions with sources, sinks and external systems



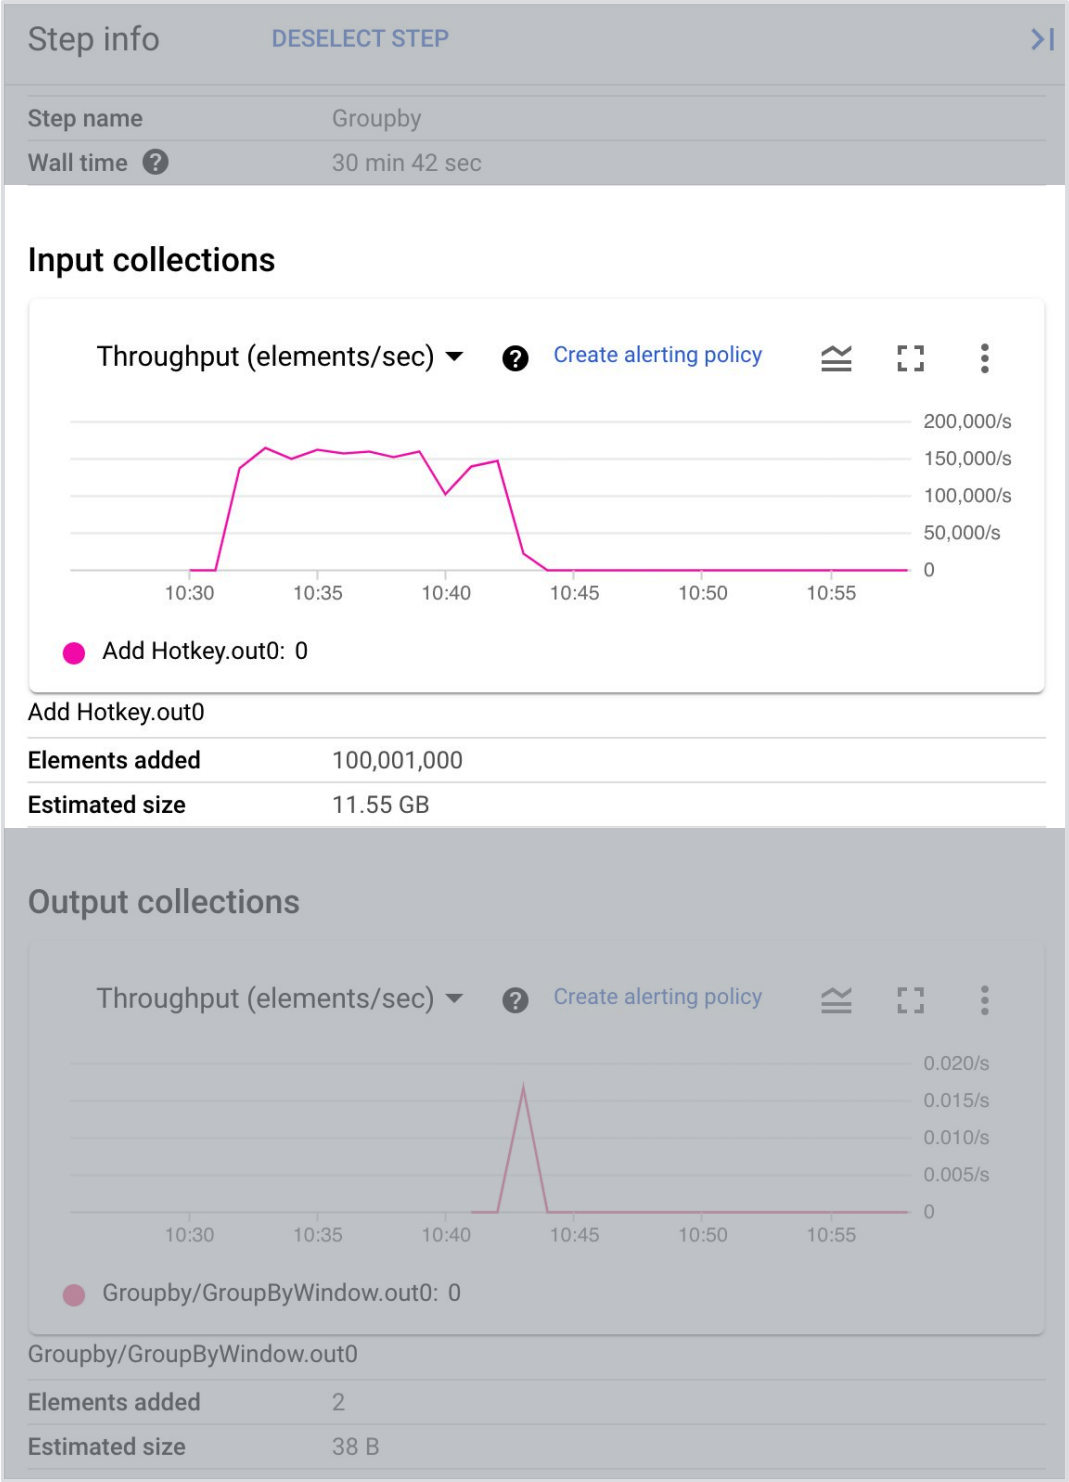
# Performance problems



# Performance problems



# Performance problems





# Performance problems

