



Performance

Ajay Kumar Yadav

Strategic Cloud Engineers,
Google Cloud





Agenda

Course Intro

Monitoring

Logging and Error Reporting

Troubleshooting and Debugging

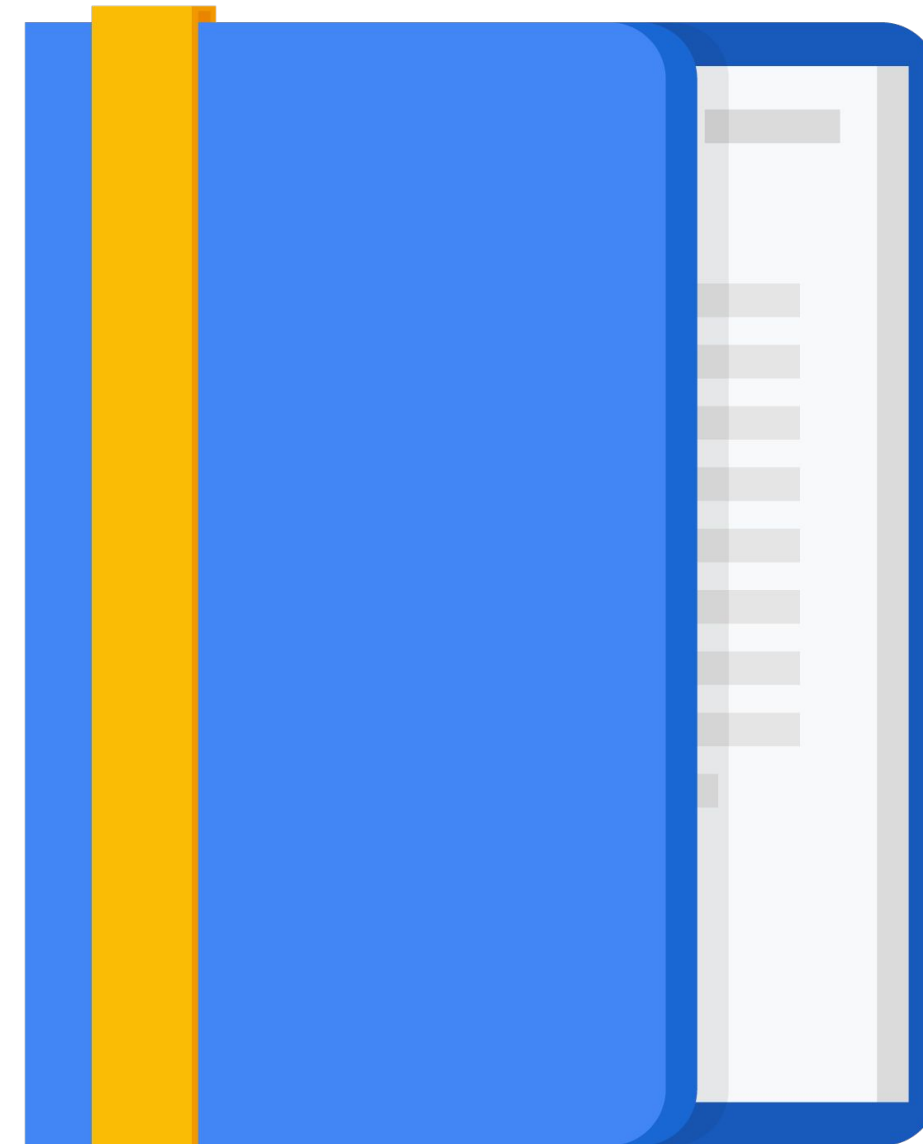
Performance

Testing and CI/CD

Reliability

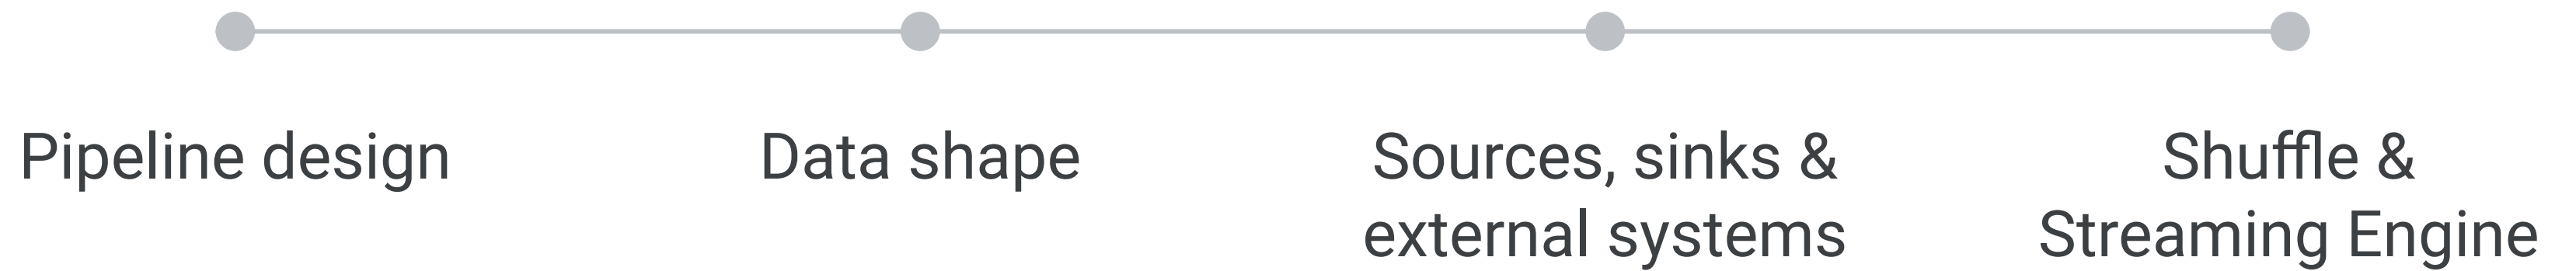
Flex Templates

Course Summary



Performance

Agenda



Performance

Agenda



Pipeline design

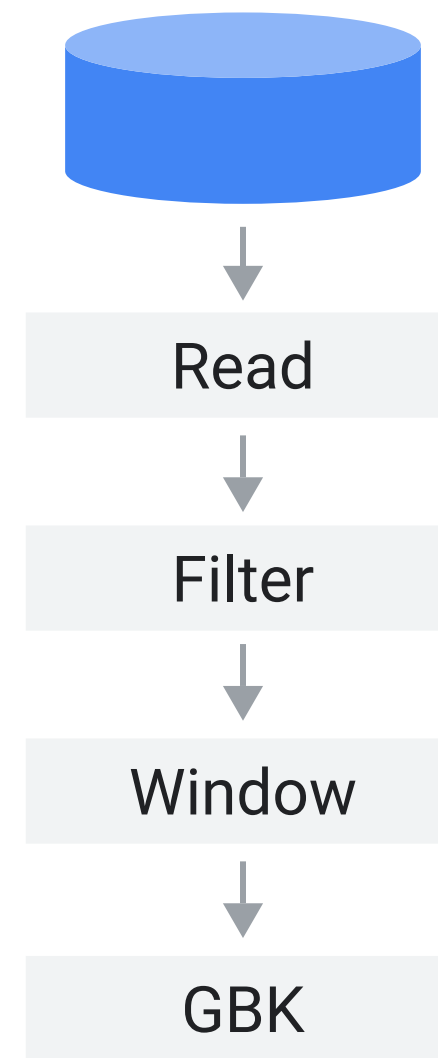
Data shape

Sources, sinks &
external systems

Shuffle &
Streaming Engine

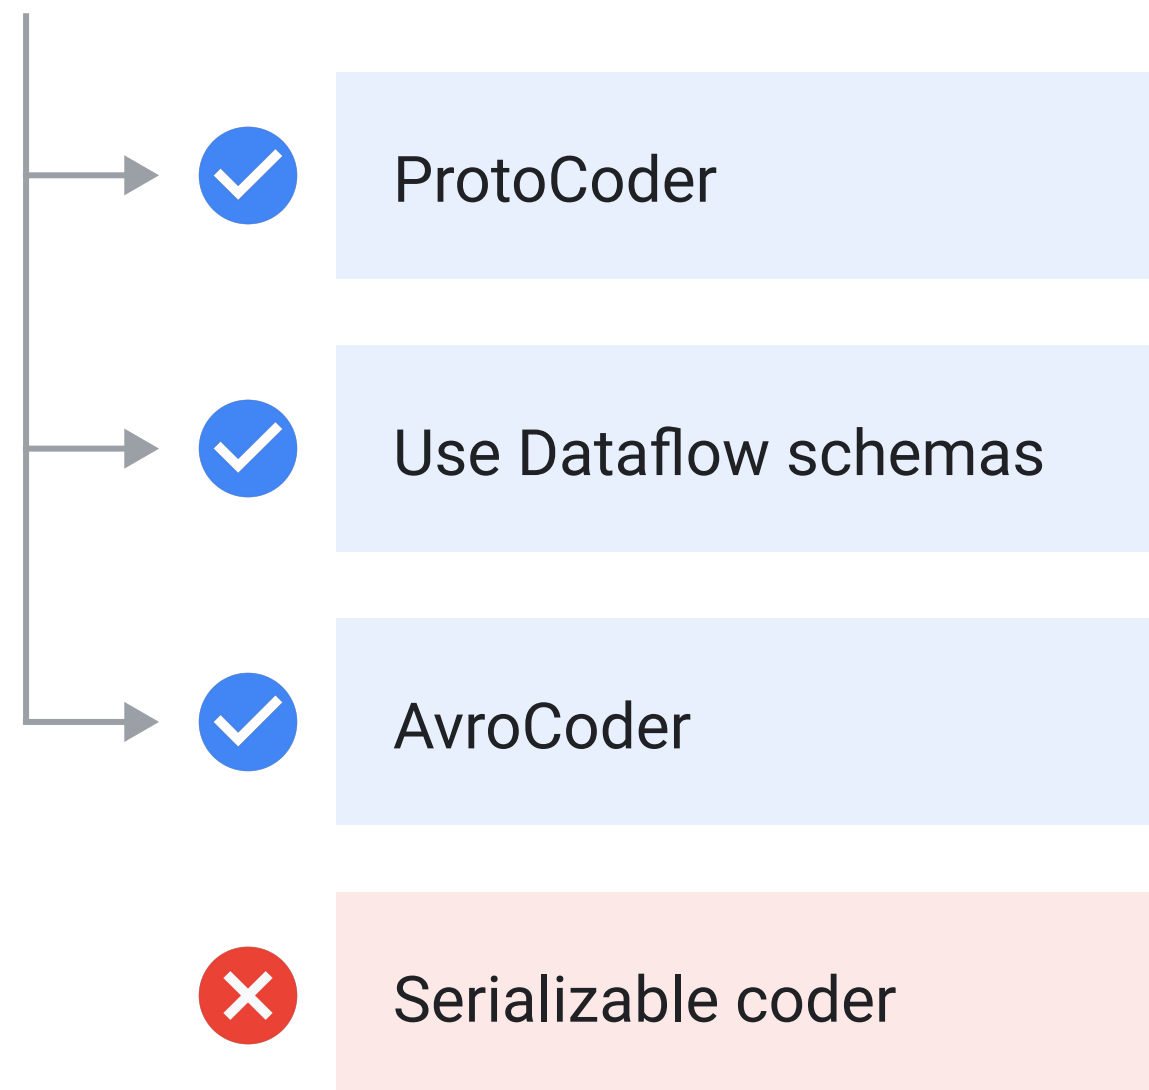
Design: Topology

Filter early in pipeline and place transformations that reduce the volume of data as high up on the graph as possible.



Design: Coders

Utilize efficient coders



Annotating a custom data type [Java]

```
@DefaultCoder(AvroCoder.class)
public class MyCustomDataType {
    ...
}
```

Set a default coder [Java]

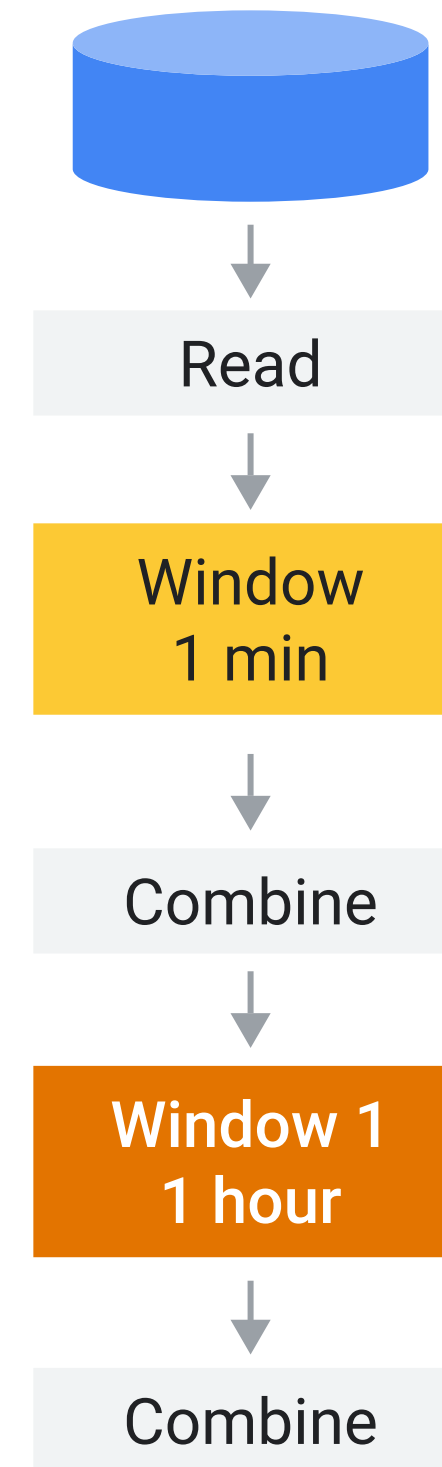
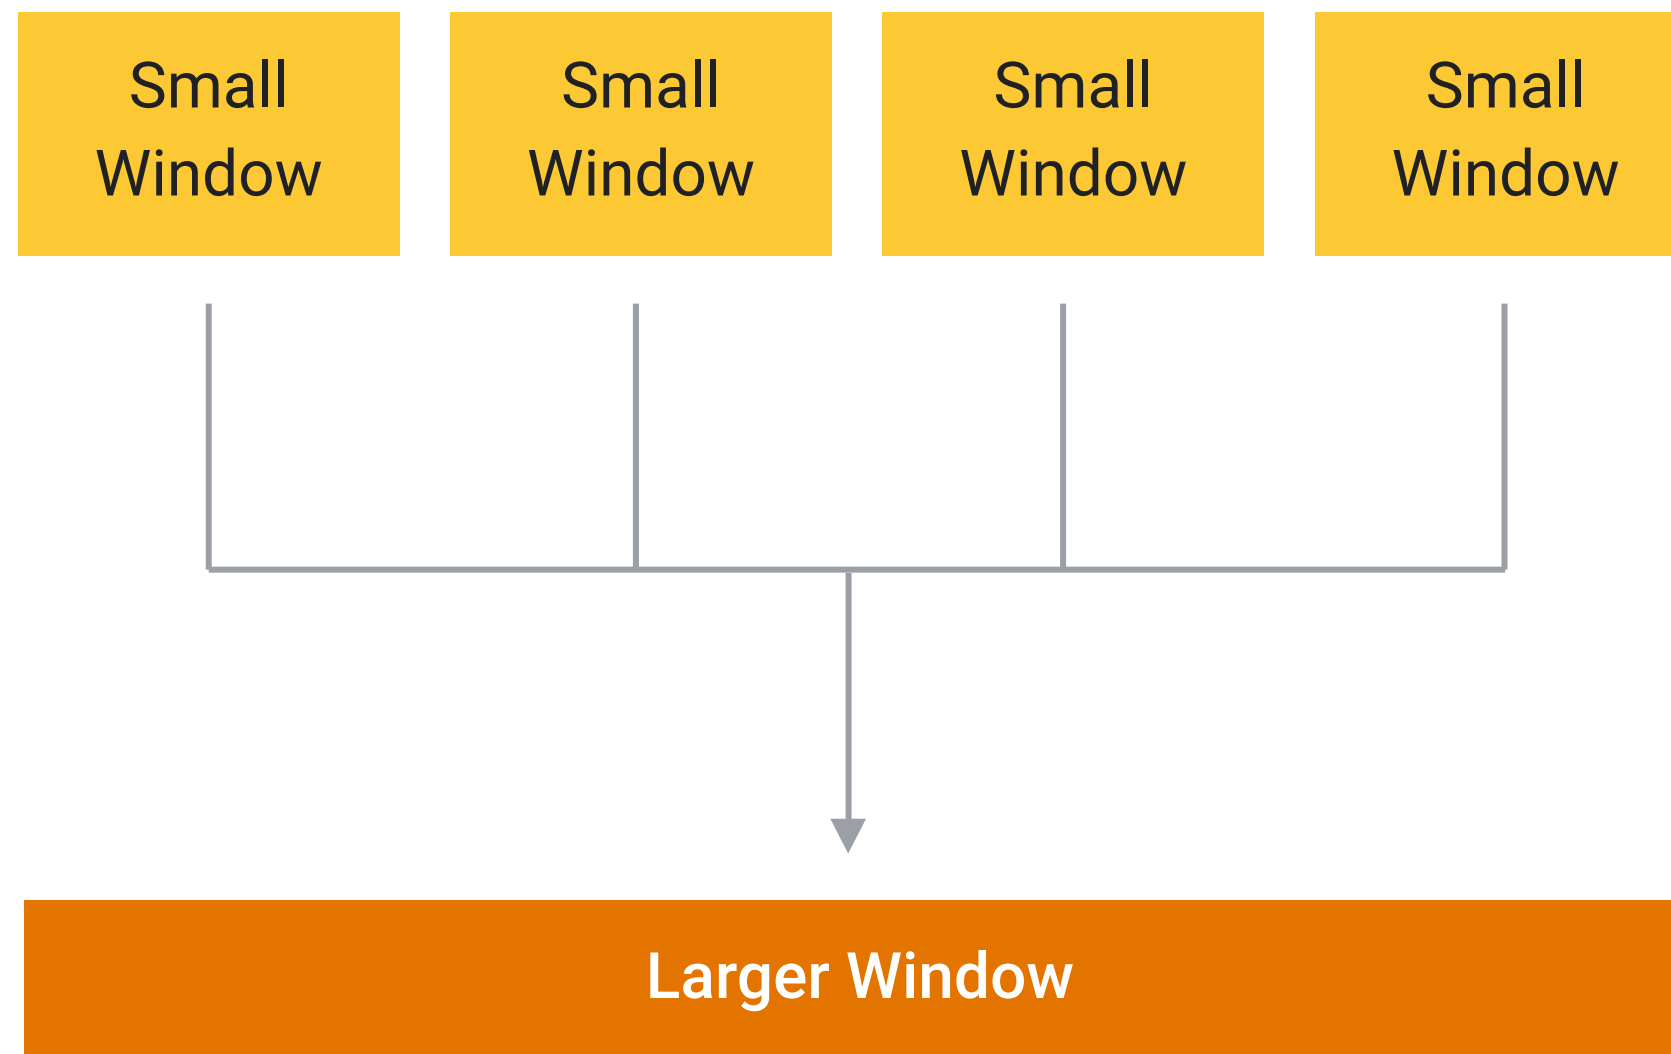
```
PipelineOptions options =
PipelineOptionsFactory.create();
Pipeline p = Pipeline.create(options);

CoderRegistry cr = p.getCoderRegistry();
cr.registerCoder(Integer.class,
BigEndianIntegerCoder.class);
```

Set a default coder [Python]

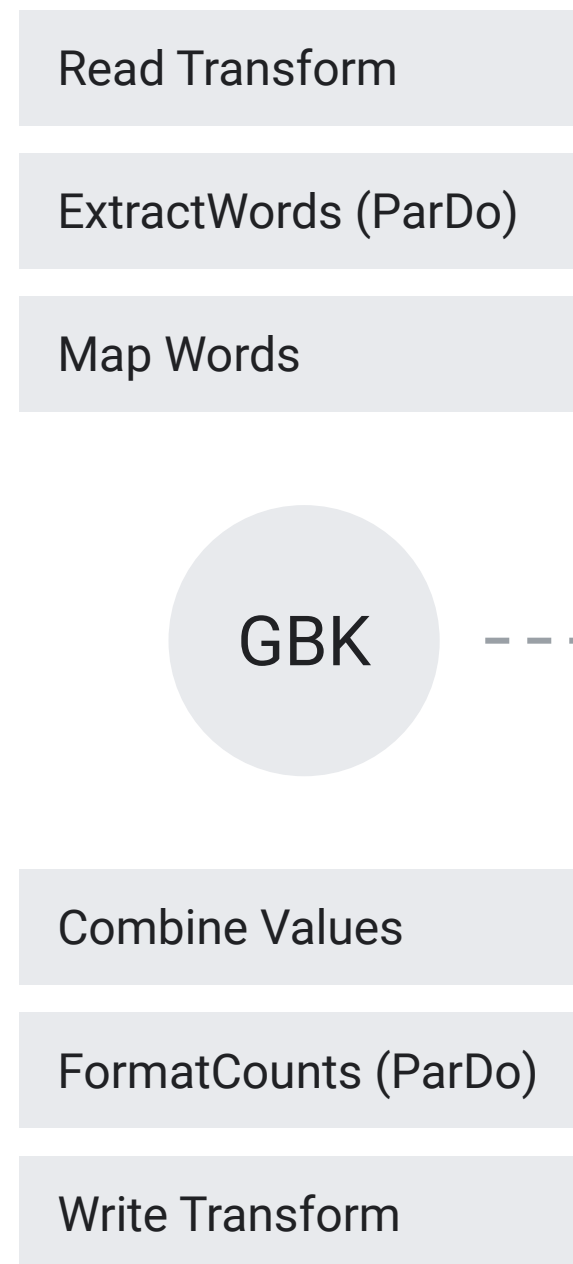
```
apache_beam.coders.registry.register_coder(i
nt, BigEndianIntegerCoder)
```

Design: Window considerations

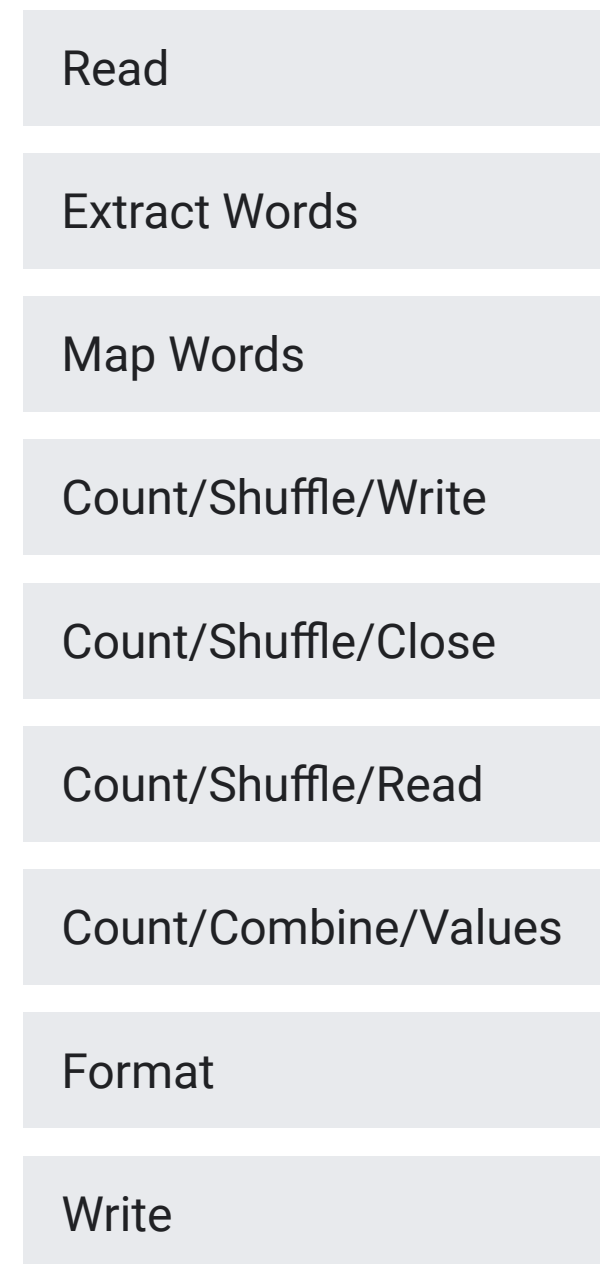


Design: Graph optimization

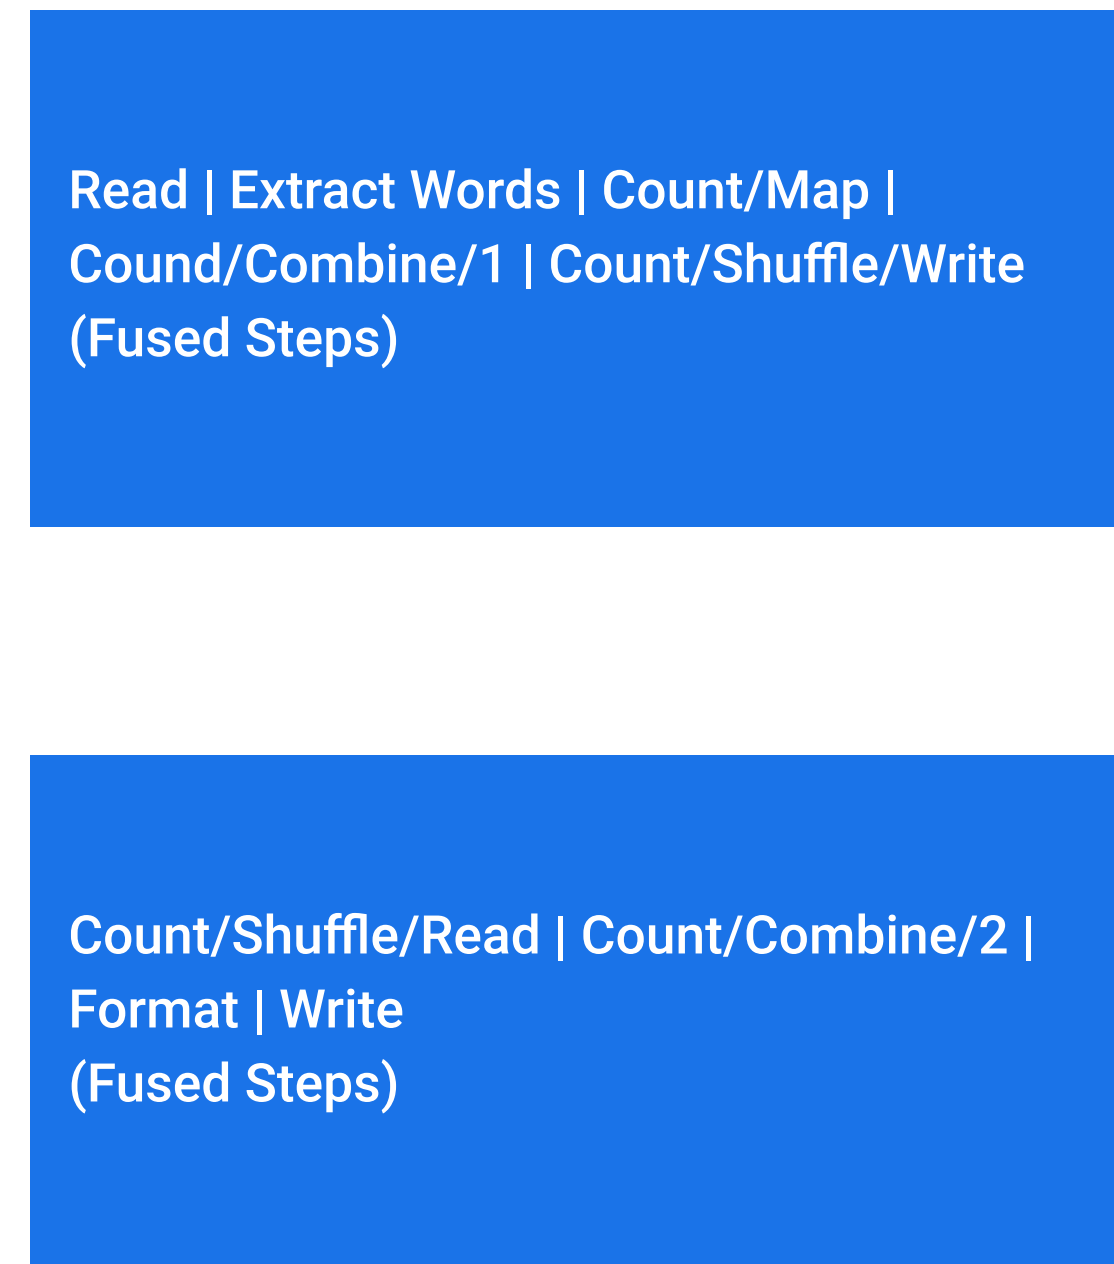
Inlined transformations



Expanded execution graph

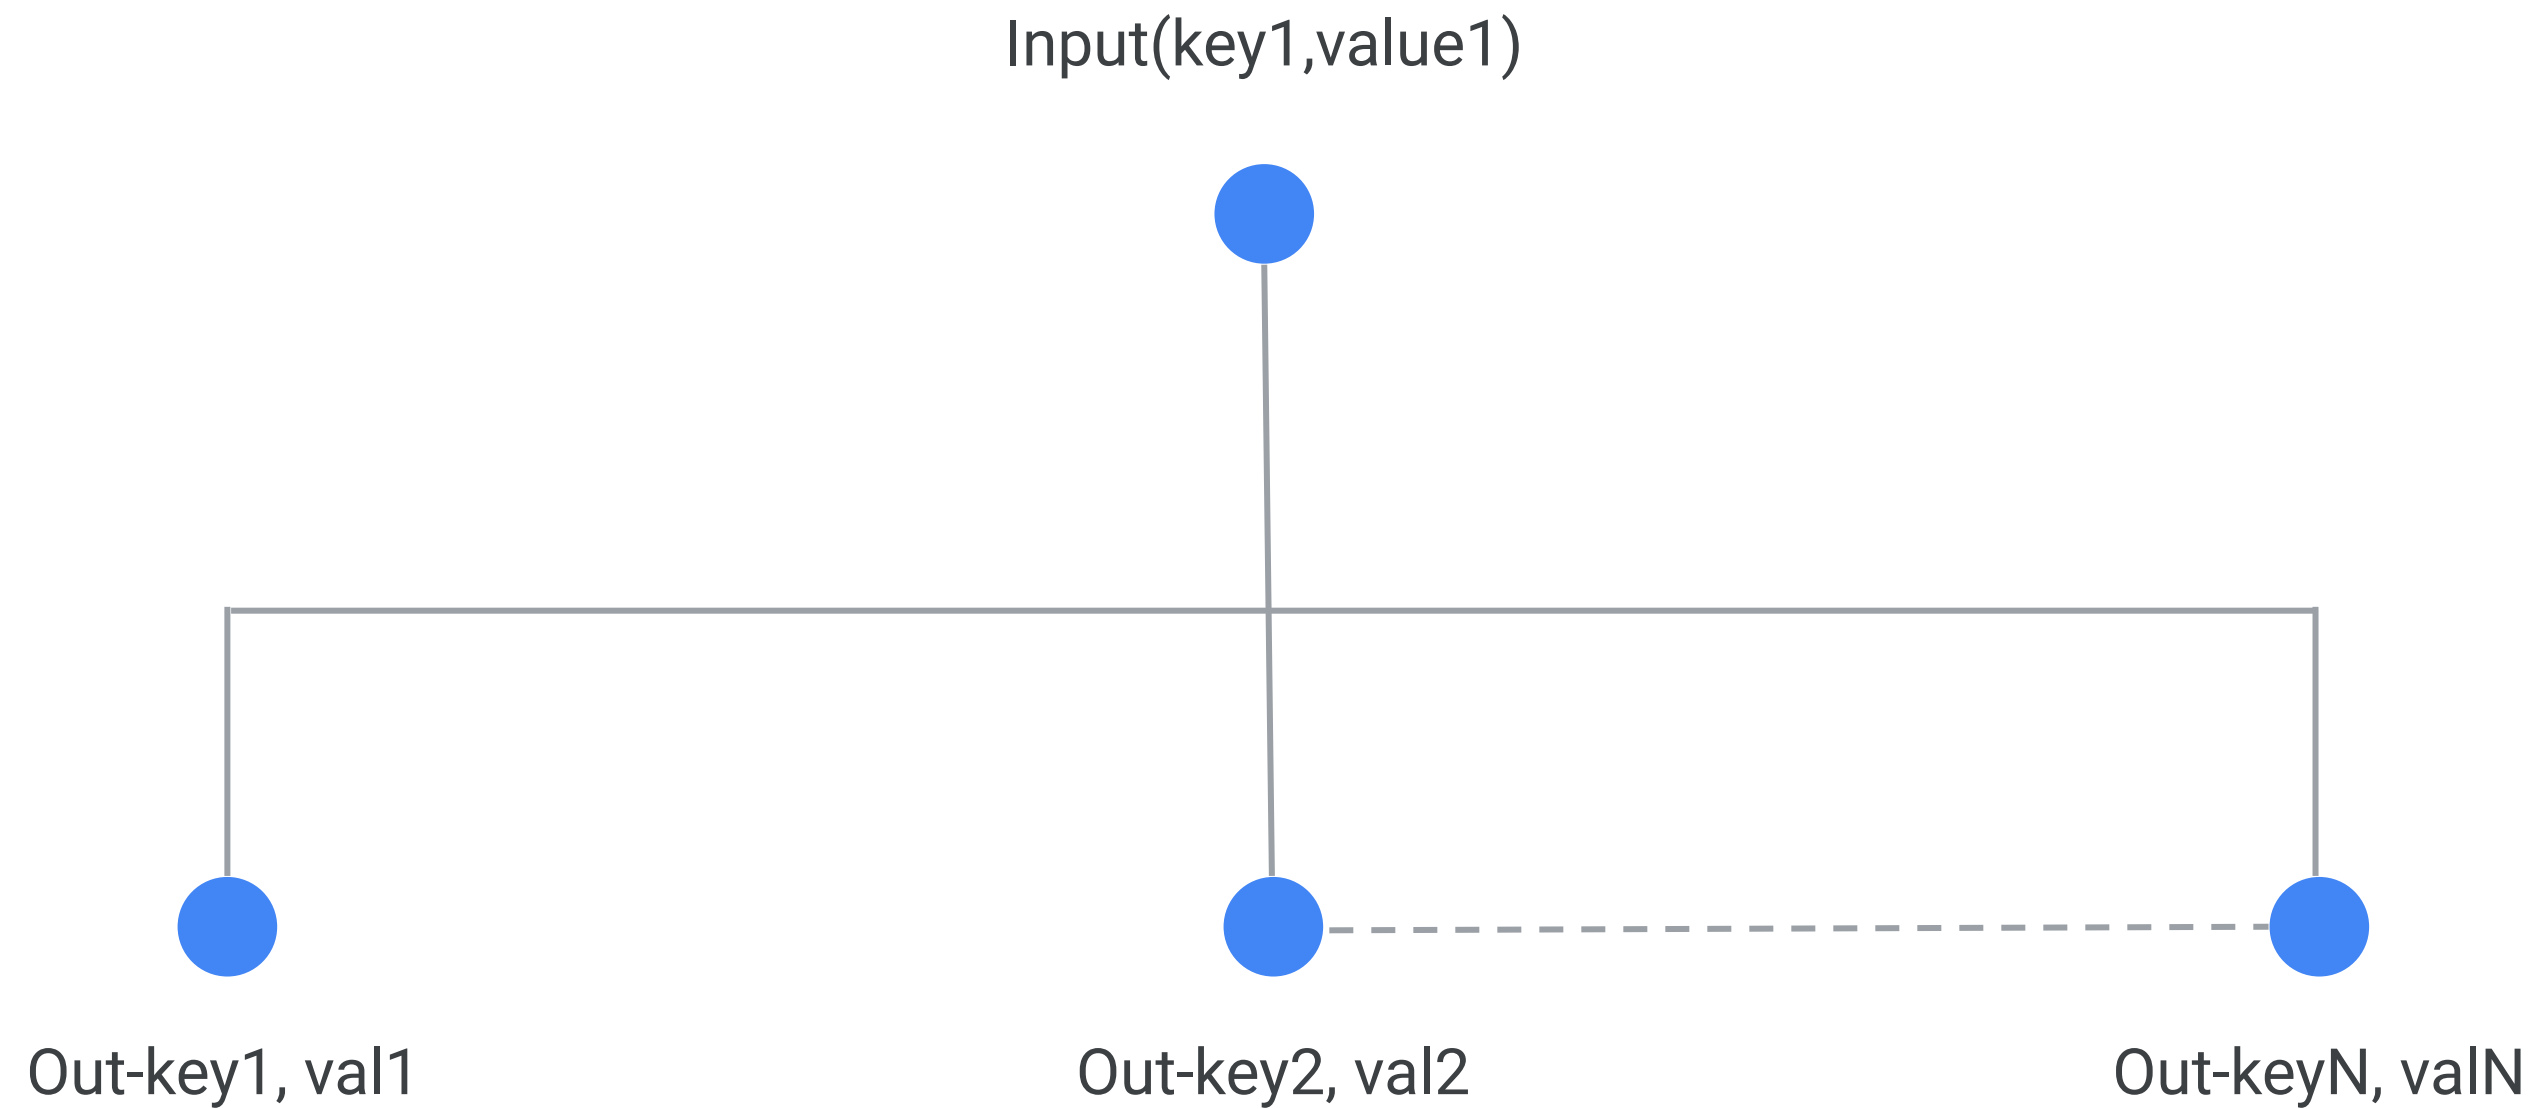


Optimized execution graph



Design: Graph optimization

Fanout transformation

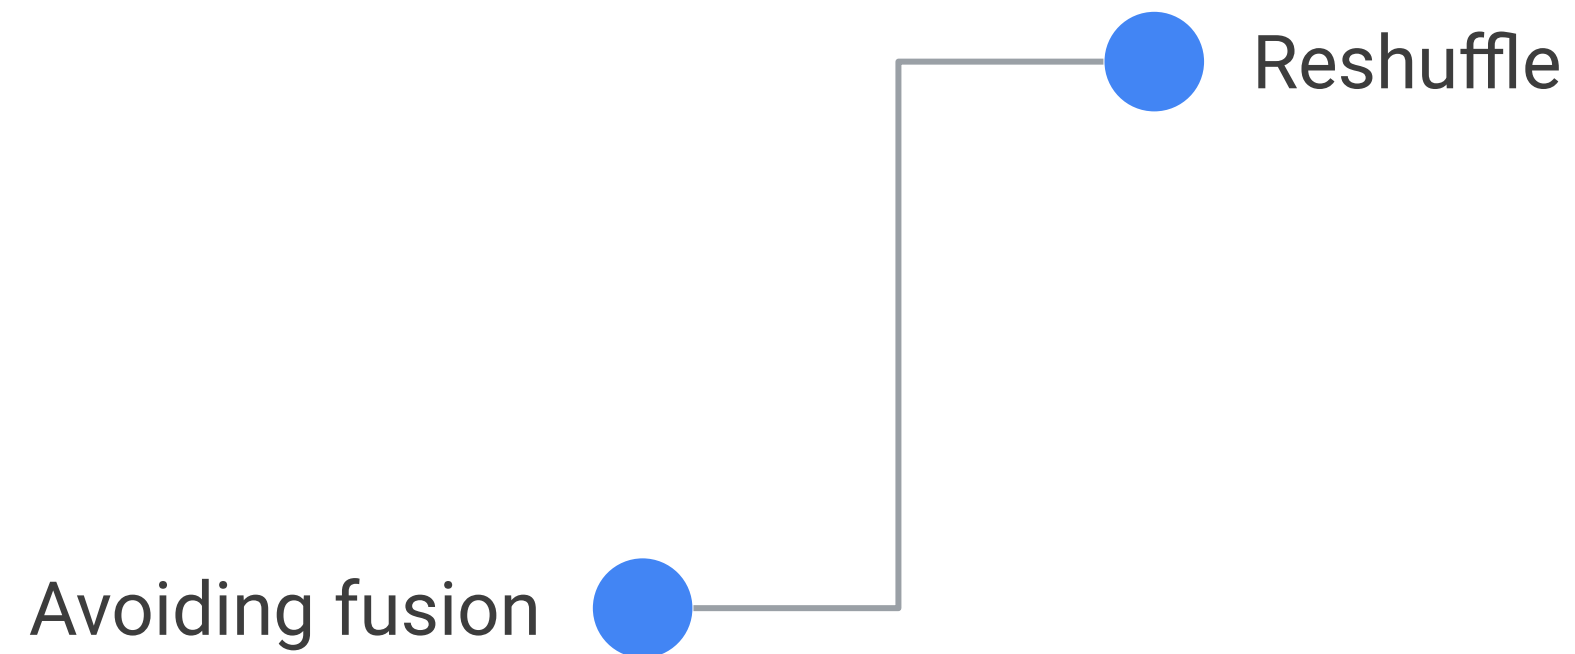


Design: Graph optimization

Avoiding fusion



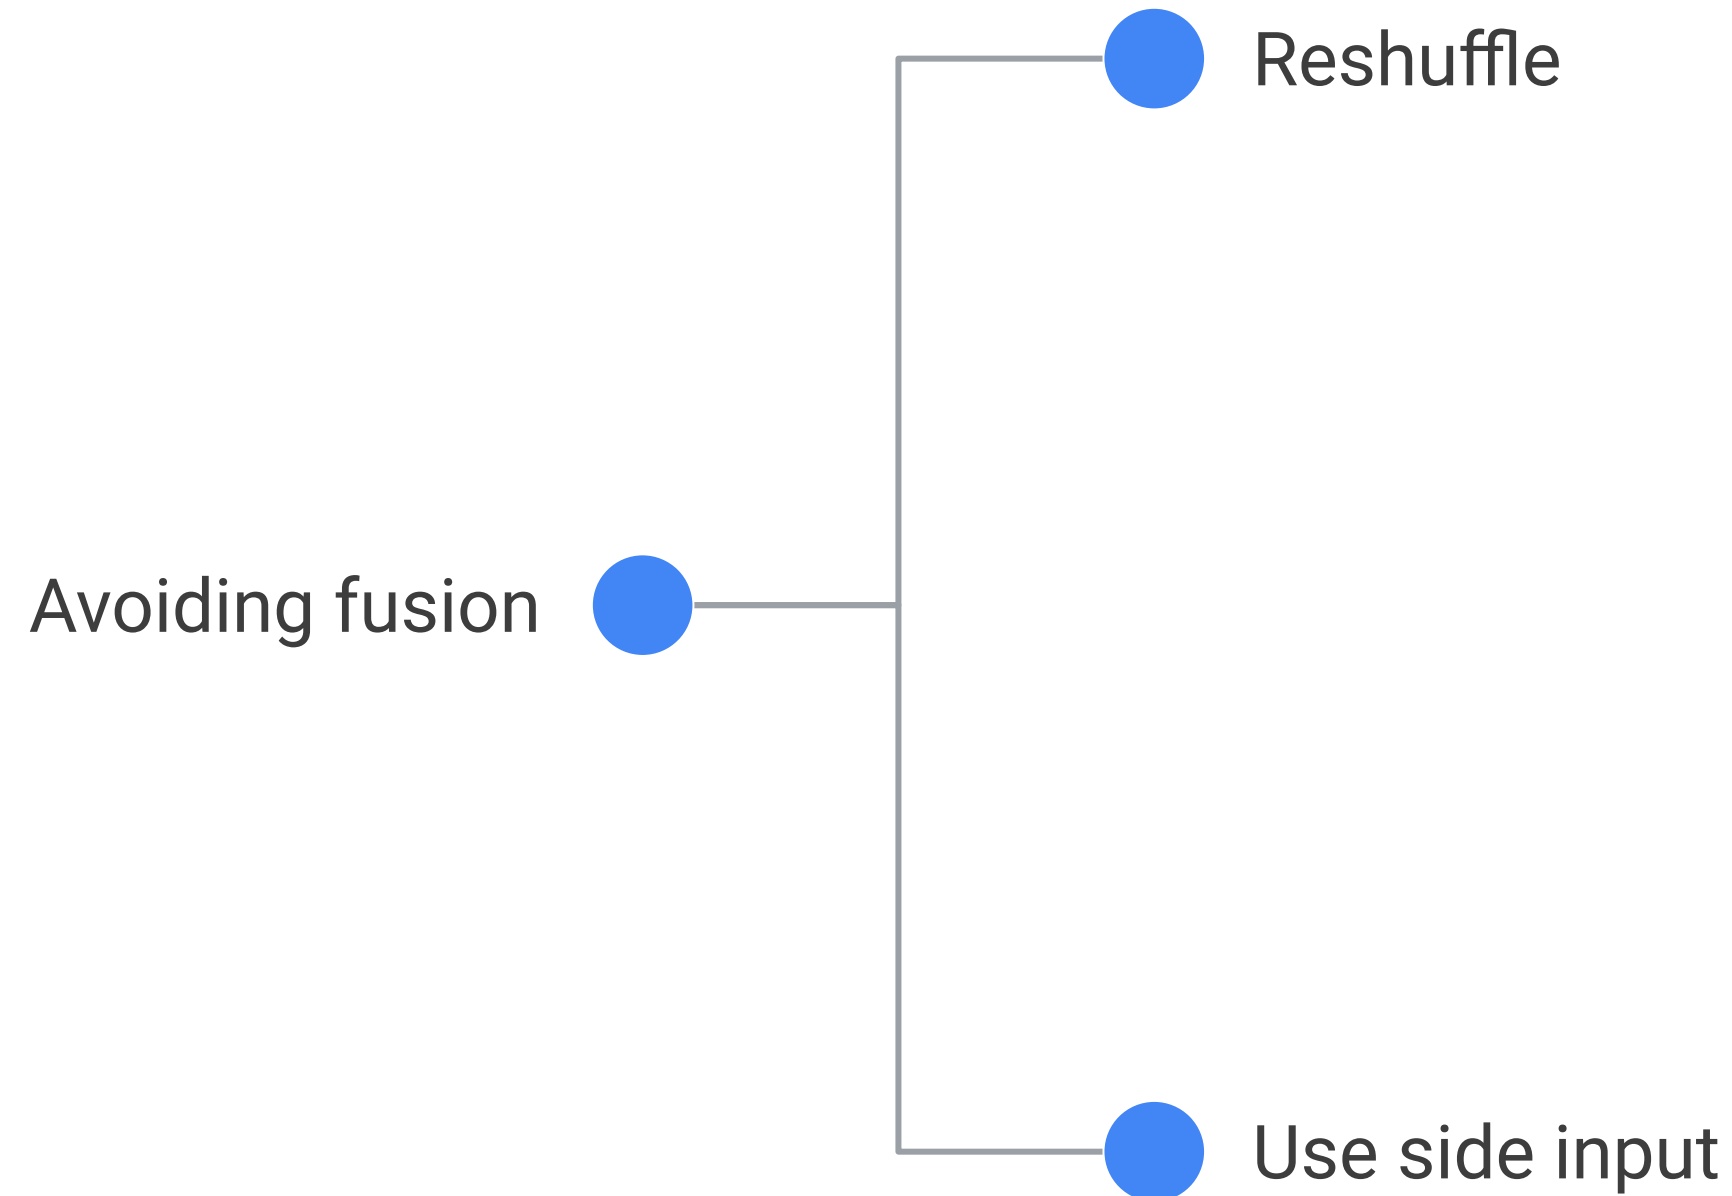
Design: Graph optimization



```
output_collection = output_collection |  
'Reshuffle' >> beam.Reshuffle()
```

```
input.apply("Reshuffle", Reshuffle.viaRandomKey())
```

Design: Graph optimization



```
output_collection = output_collection |  
'Reshuffle' >> beam.Reshuffle()
```

```
input.apply("Reshuffle", Reshuffle.viaRandomKey())
```

```
messages.apply(ParDo.of(new DoFn<PubsubMessage,  
TableRow>() {  
  @ProcessElement  
  public void processElement(ProcessContext c) {  
    TableRow sideInputData =  
      c.sideInput(mapData).get(key);  
    // Do something with side input  
  }).withSideInputs(mapData));
```

```
result = ( main_input | 'ApplyCrossJoin' >>  
beam.FlatMap(cross_join, rights=beam.pvalue.AsIter  
(side_input)))
```

Design: Logging

1

Logging considerations

Design: Logging

1

Logging considerations

- Strike a balance b/w excessive logging vs no/little logging at all.

Design: Logging

1

Logging considerations

- Strike a balance b/w excessive logging vs no/little logging at all.
- Avoid logging at info level against PCollection element granularity.

Design: Logging

1

Logging considerations

- Strike a balance b/w excessive logging vs no/little logging at all.
- Avoid logging at info level against PCollection element granularity.

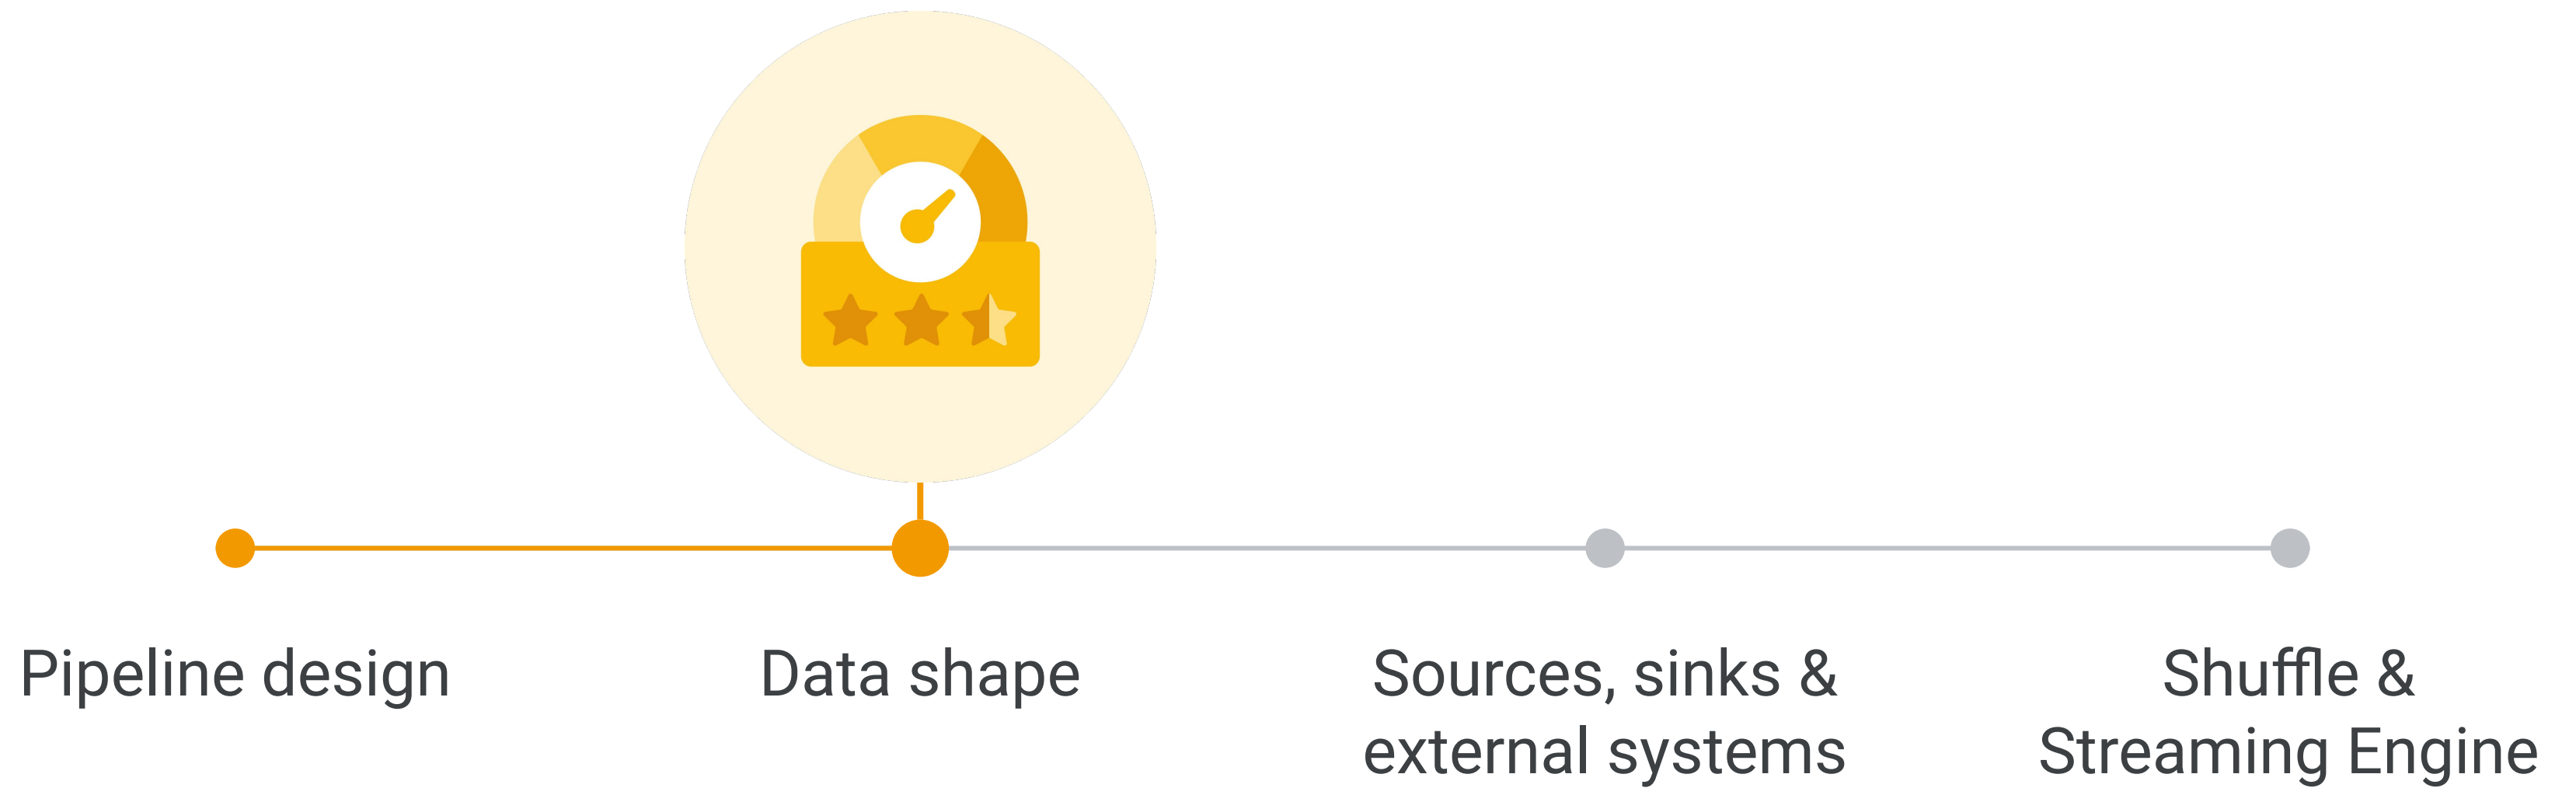
2

Use a dead letter pattern

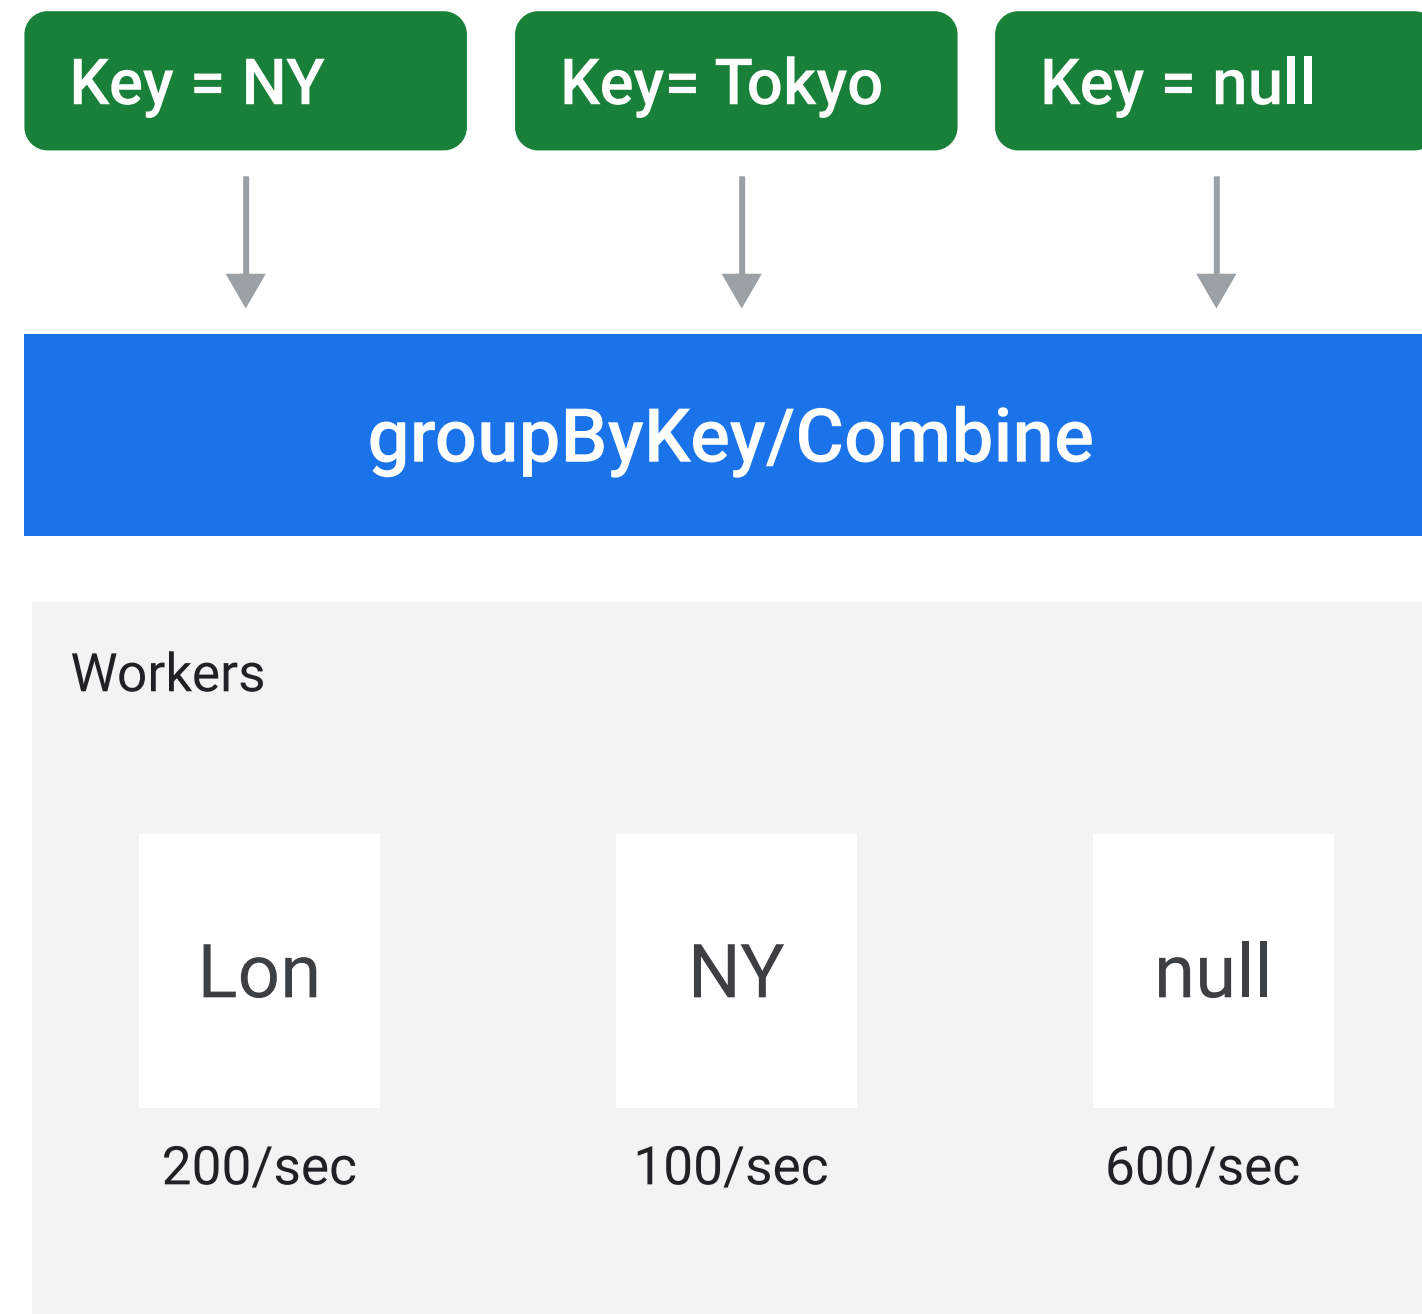
- Use a dead letter pattern followed by a count per window (ex: 5 mins) for reporting data errors.

Performance

Agenda

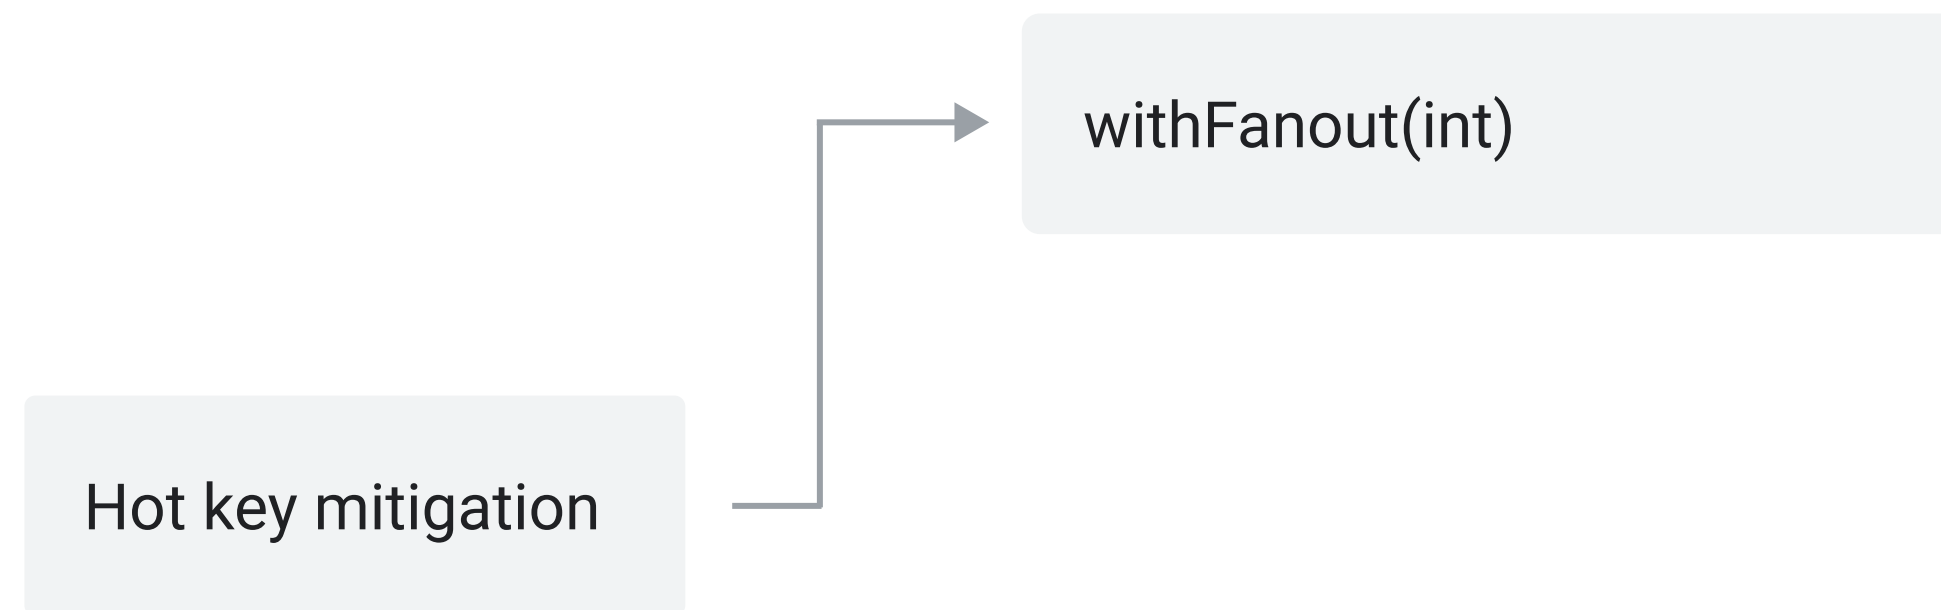


Design: Data skew



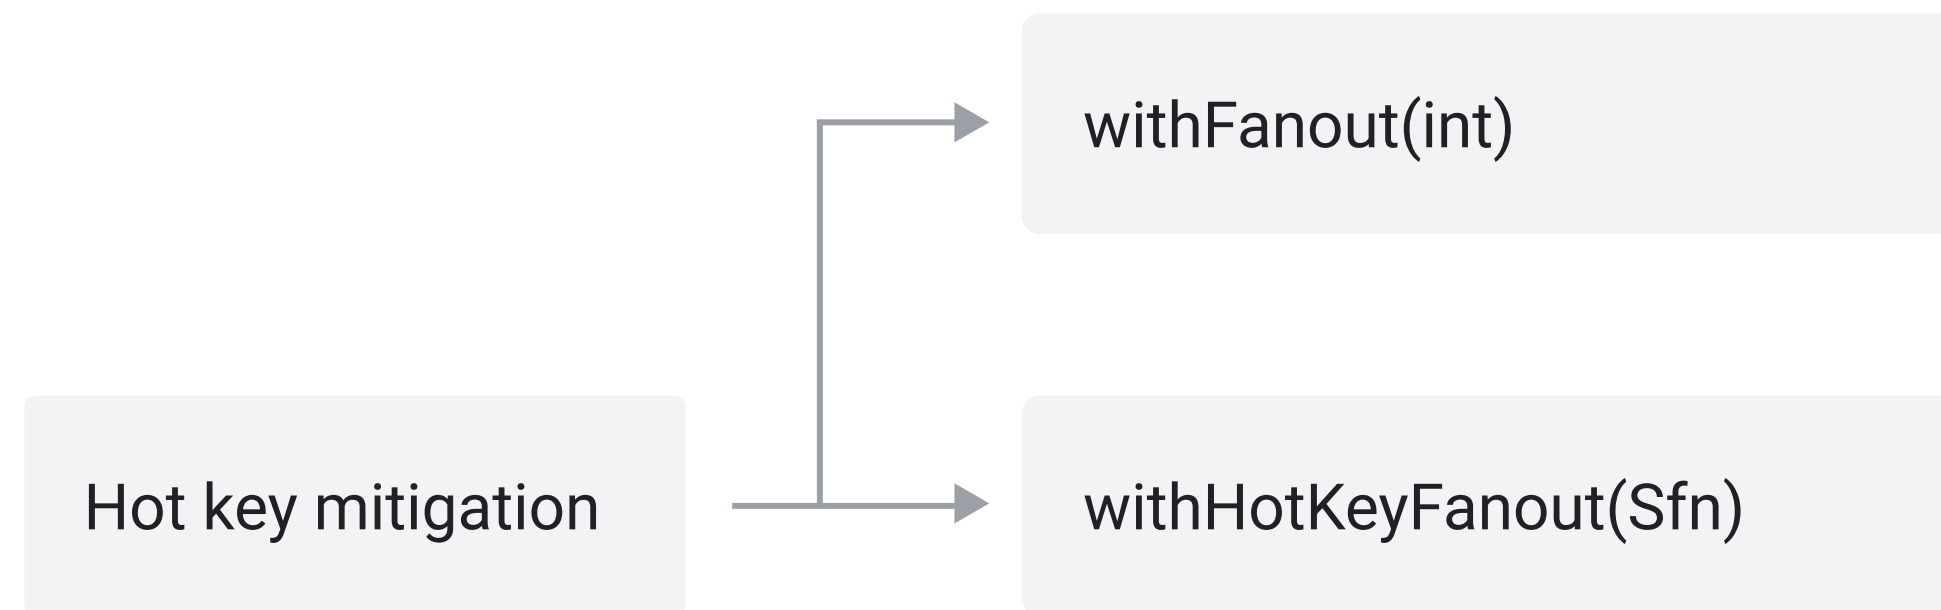
Design: Data skew

Log hotkeys by setting `hotKeyLoggingEnabled` to true



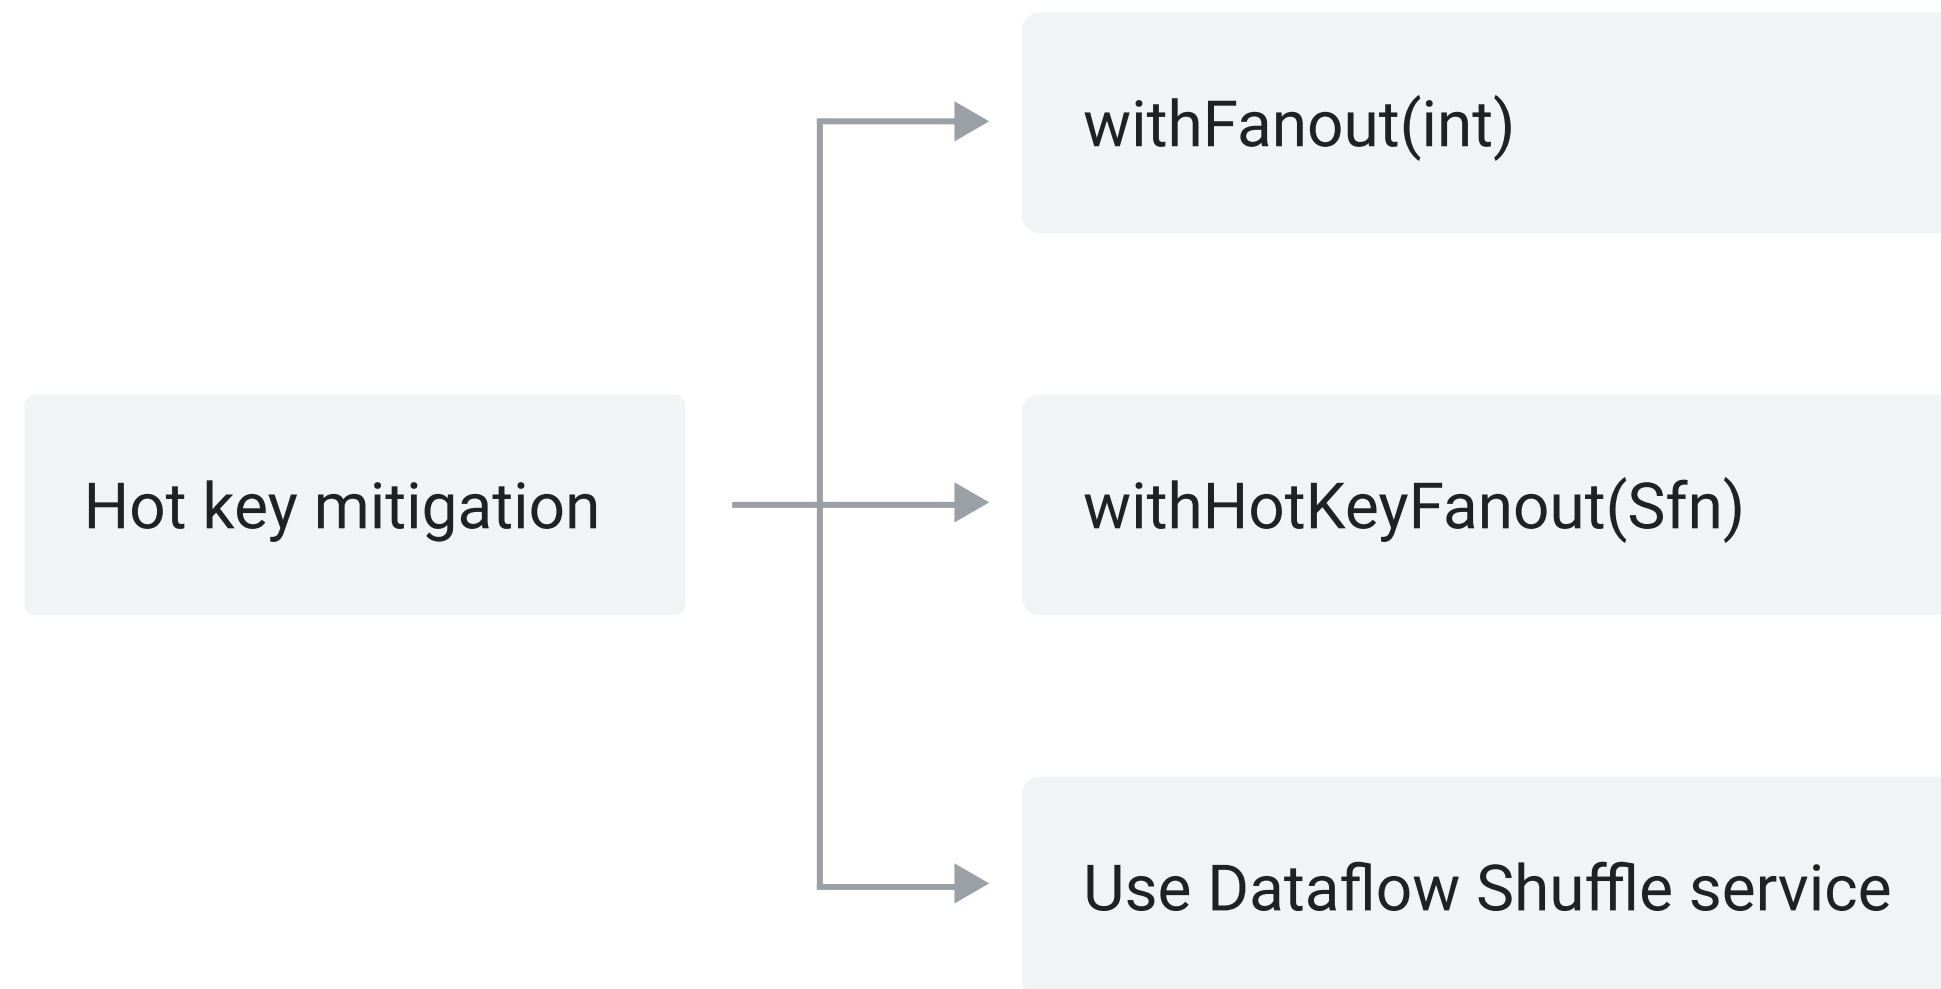
Design: Data skew

Log hot keys by setting `hotKeyLoggingEnabled` to `true`



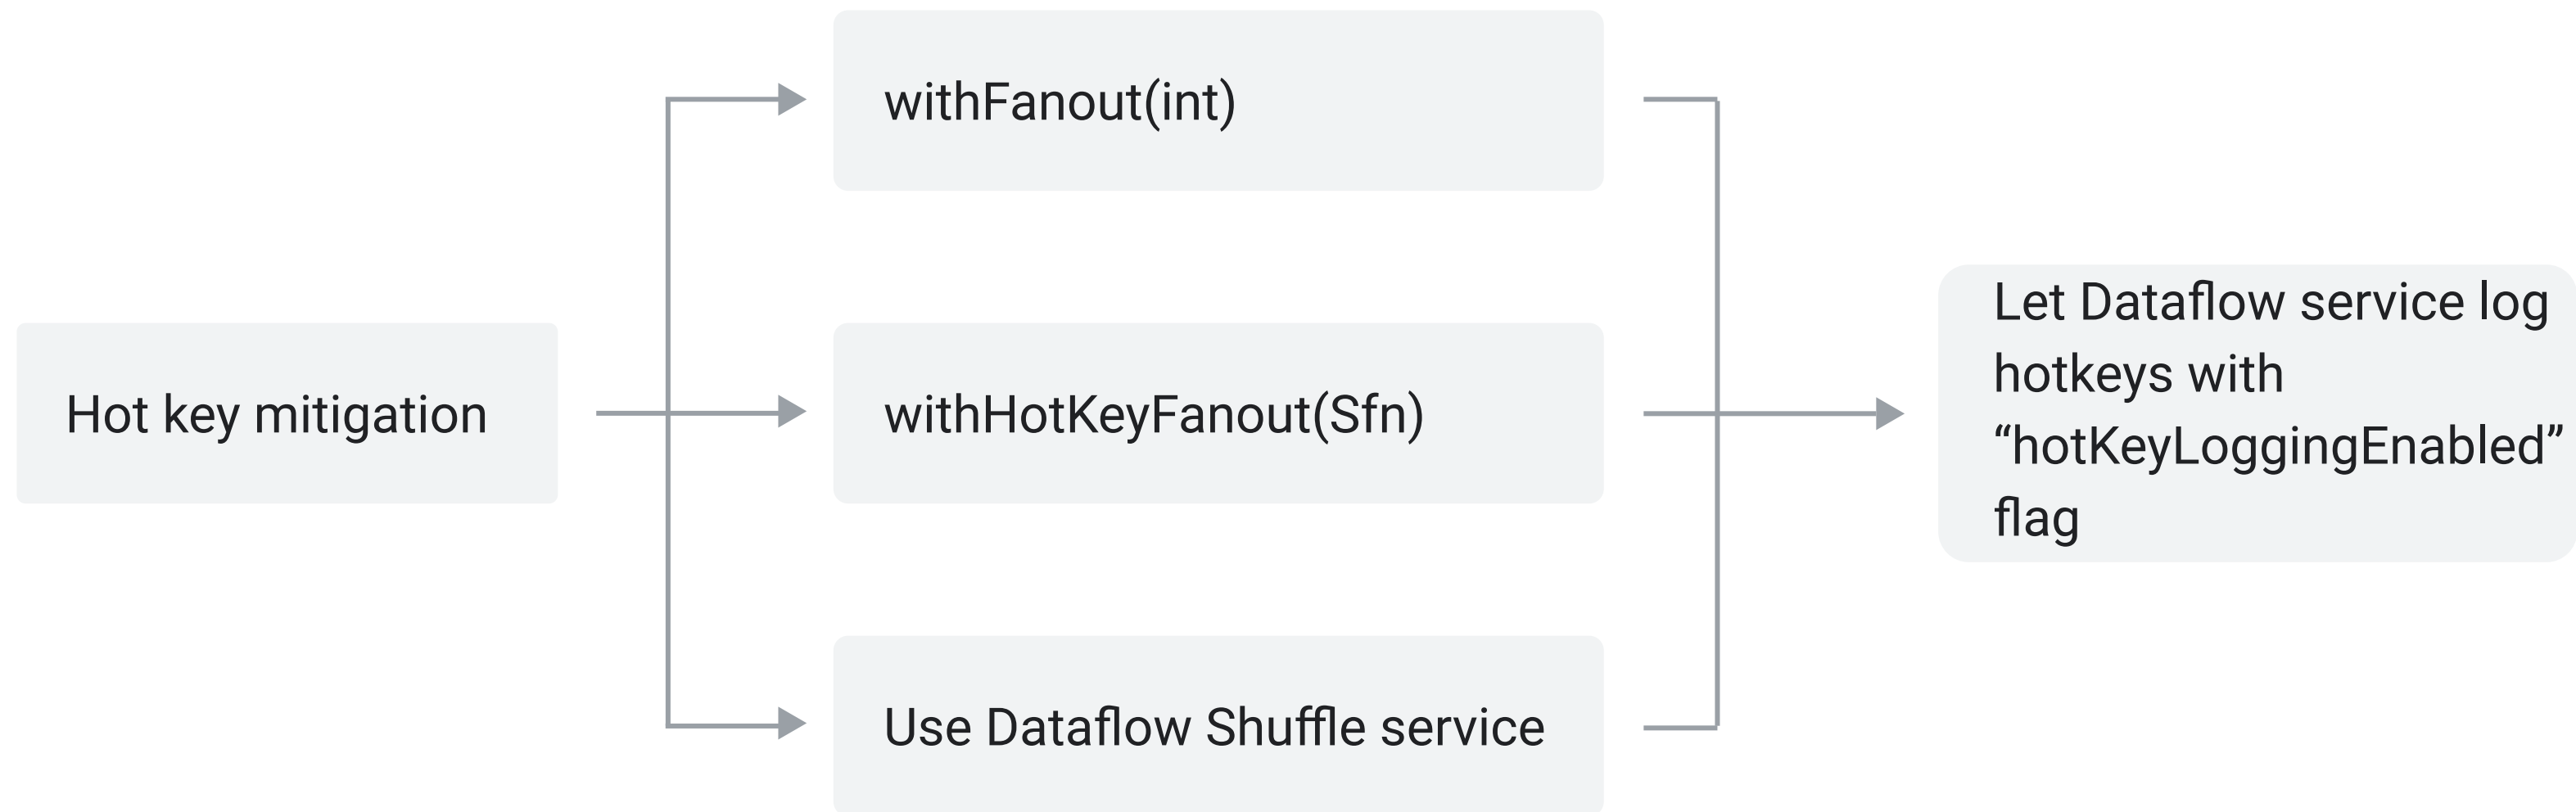
Design: Data skew

Log hot keys by setting `hotKeyLoggingEnabled` to `true`



Design: Data skew

Log hot keys by setting `hotKeyLoggingEnabled` to true



Design: Key space and parallelism

Low parallelism (too few keys)

1. Increase number of keys
Example: If windows are distinct, use composite (window + key) keys.
2. If reading from files, prefer reading from splittable compression formats like Avro

parallelism = no. of keys

Design: Key space and parallelism

Low parallelism (too few keys)

1. Increase number of keys
Example: If windows are distinct, use composite (window + key) keys.
2. If reading from files, prefer reading from splittable compression formats like Avro

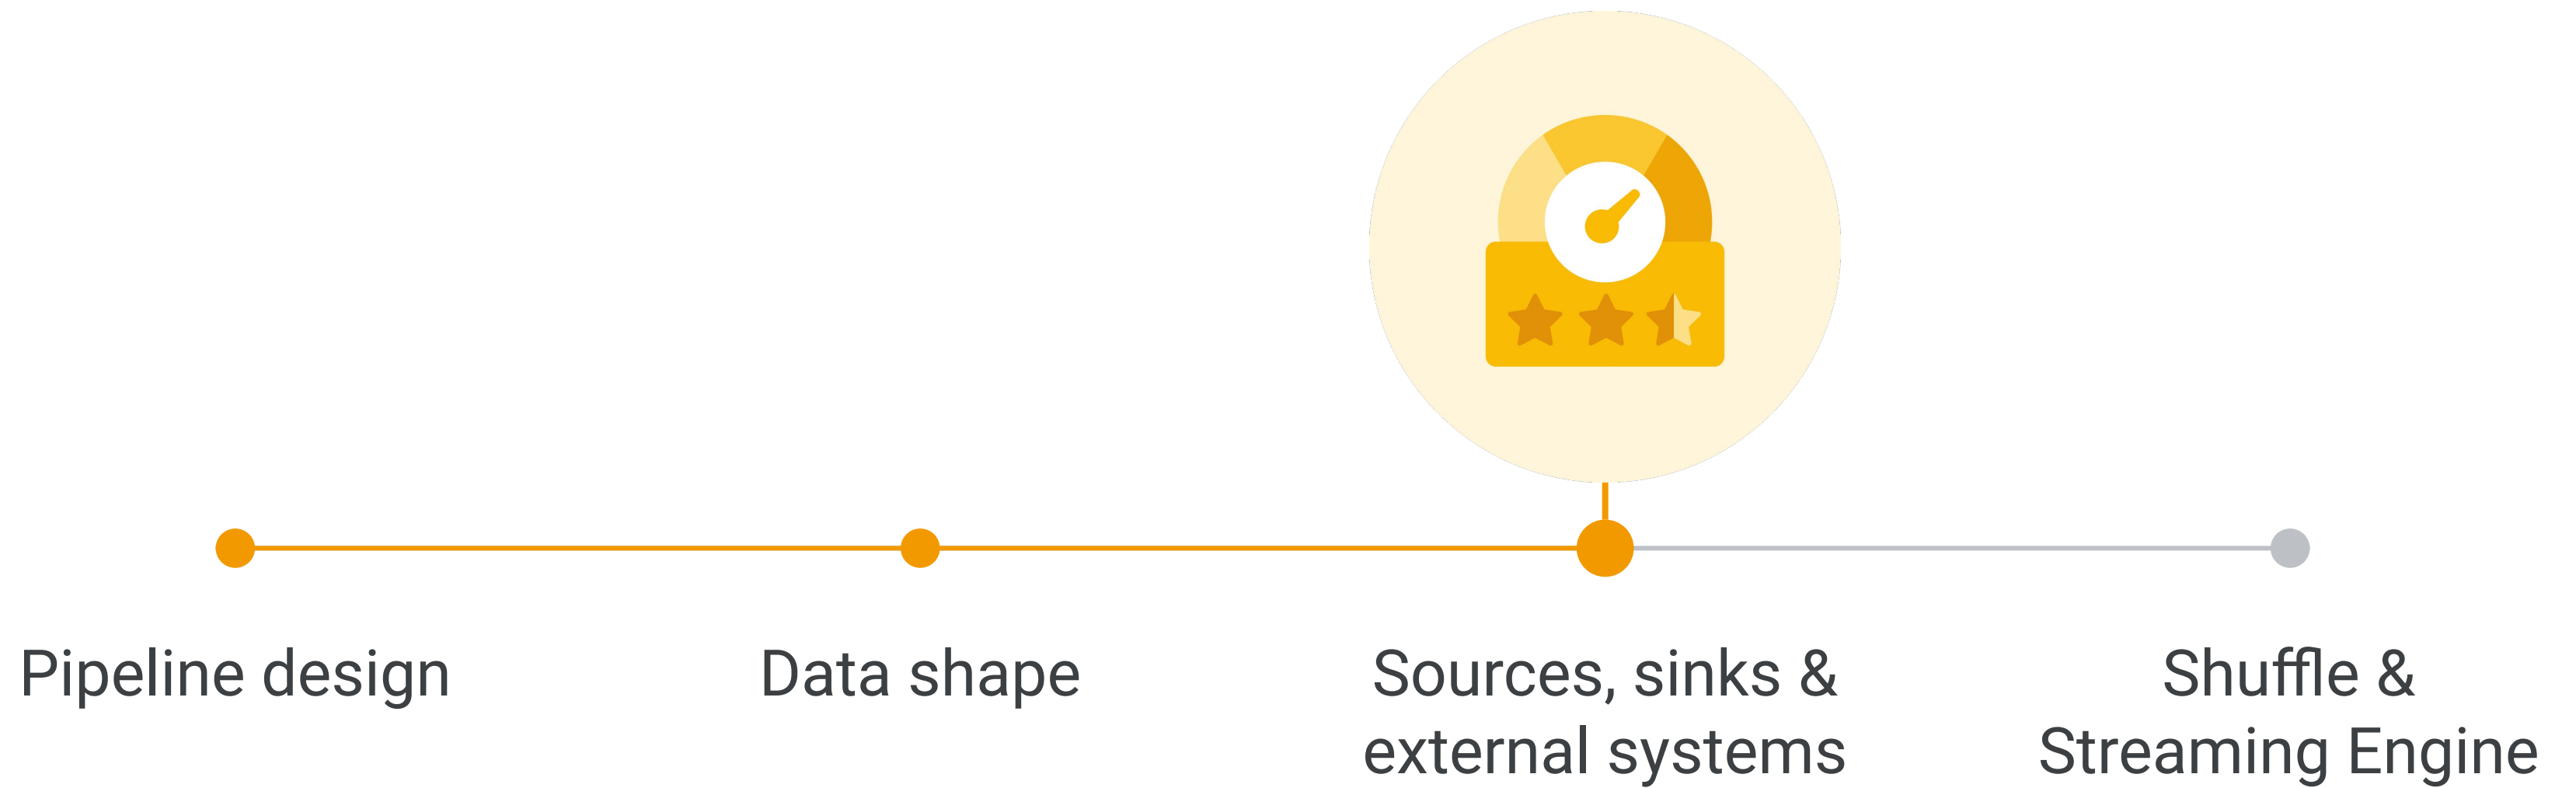
Too-high parallelism (too many keys)

1. If key space is very large, consider using hashes separating keys out internally.
2. "Re-use" processing keys from the past that are not active

parallelism = no. of keys

Performance

Agenda



Source, Sinks & External Systems

TextIO + Compressed file

Source, sinks, and external systems

TextIO + Compressed file

- ✗ Only one machine can read the compressed file.

Source, sinks, and external systems

TextIO + Compressed file

- ✗ Only one machine can read the compressed file.
- ✗ Fused stages will need to run on the same worker that read the data.

Source, Sinks & External System

TextIO + Compressed file

- ✗ Only one machine can read the compressed file.
- ✗ Fused stages will need to run on the same worker that read the data.
- ✗ A single machine will need to push all the data from the file to all other machines.

Source, Sinks & External System

TextIO + Compressed file

- ✗ Only one machine can read the compressed file.
- ✗ Fused stages will need to run on the same worker that read the data.
- ✗ A single machine will need to push all the data from the file to all other machines.

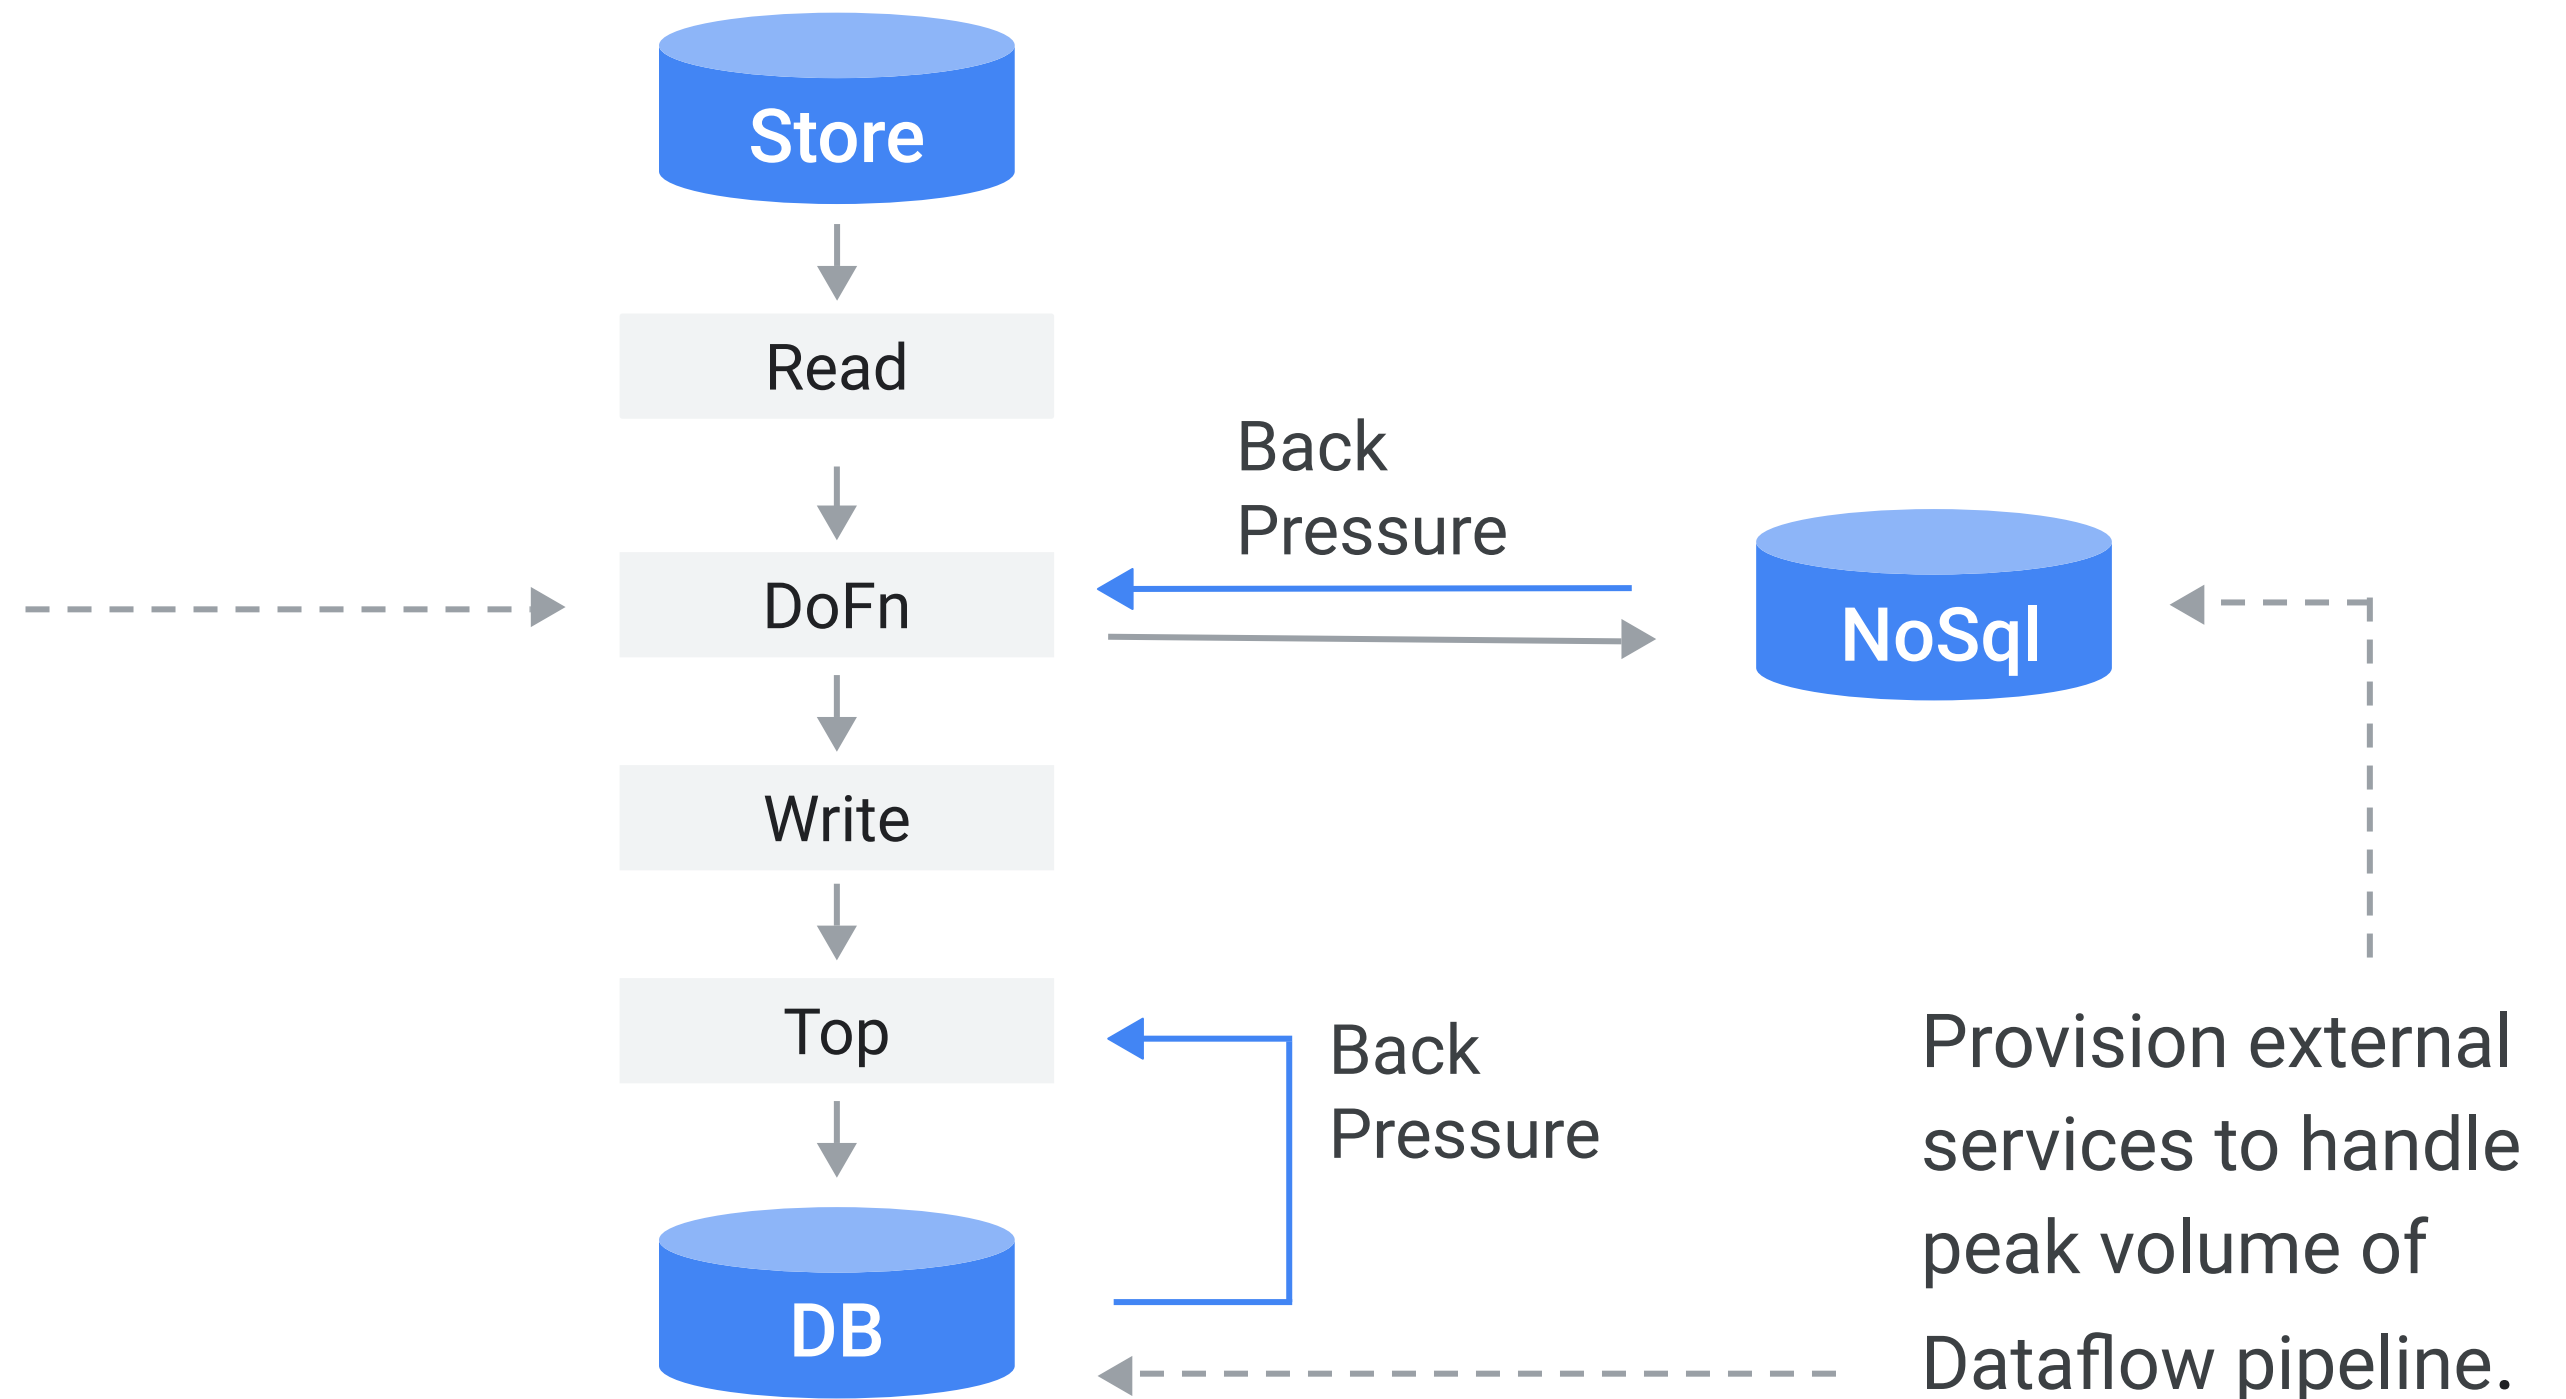
Switch to

AvroIO

TextIO +
Uncompressed

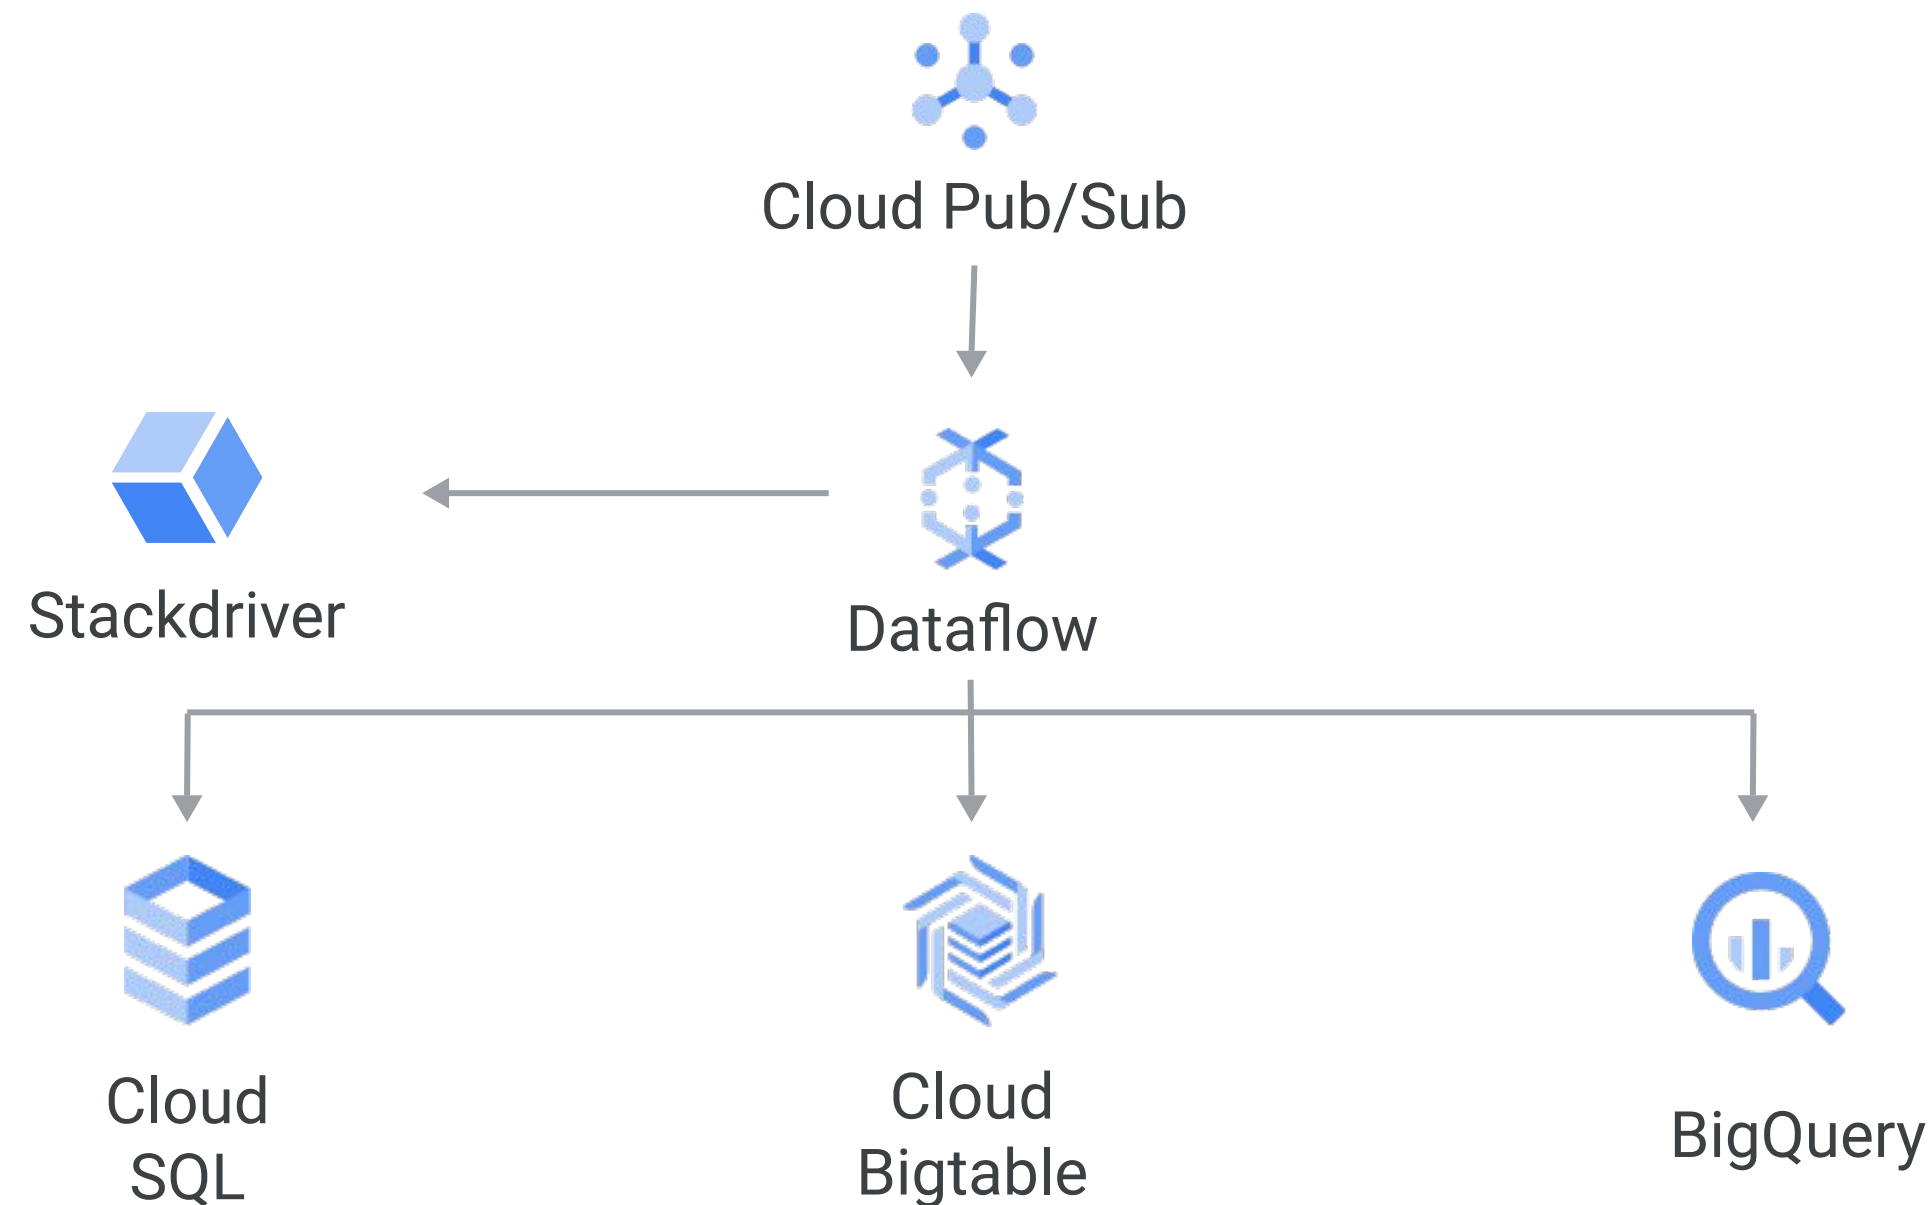
Source, Sinks & External Systems

Make use of
@StartBundle &
@FinishBundle for
micro-batching.



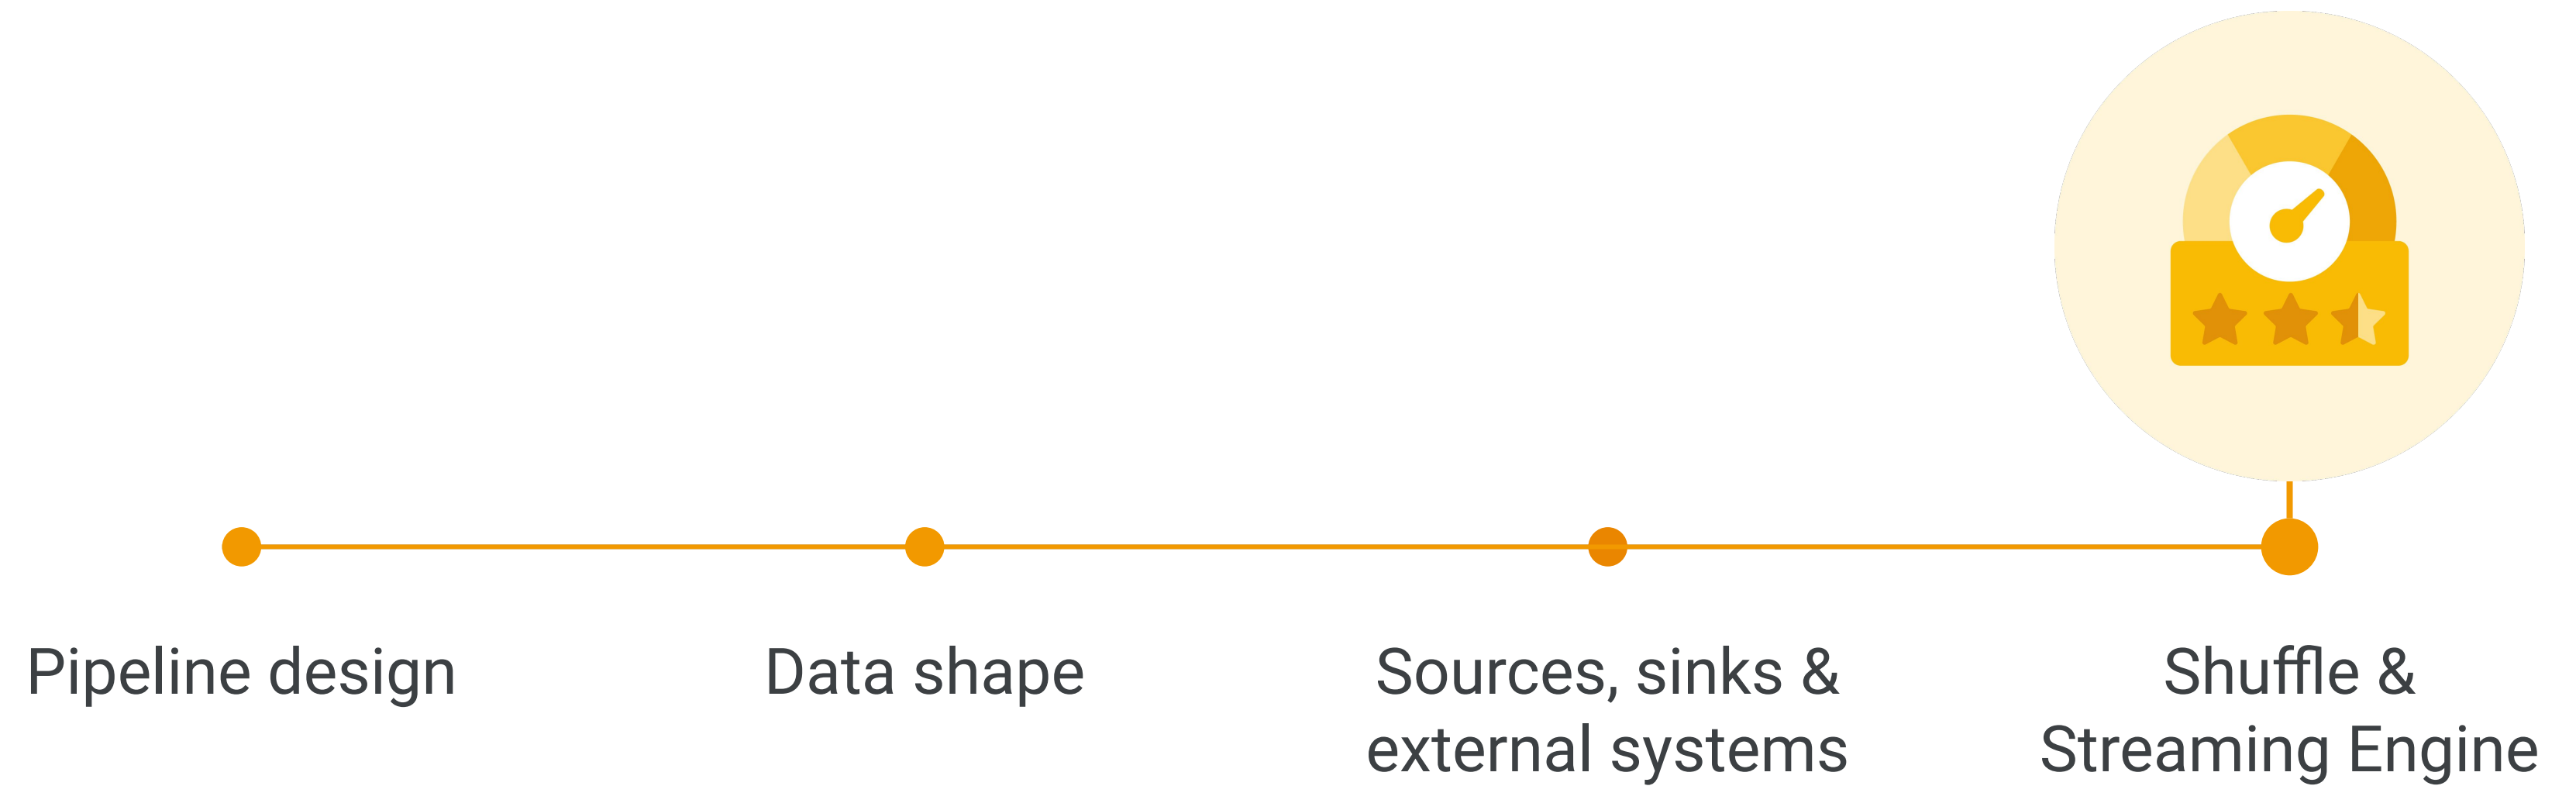
Source, Sinks & External Systems

Colocation (i.e same region/zone) helps in performance

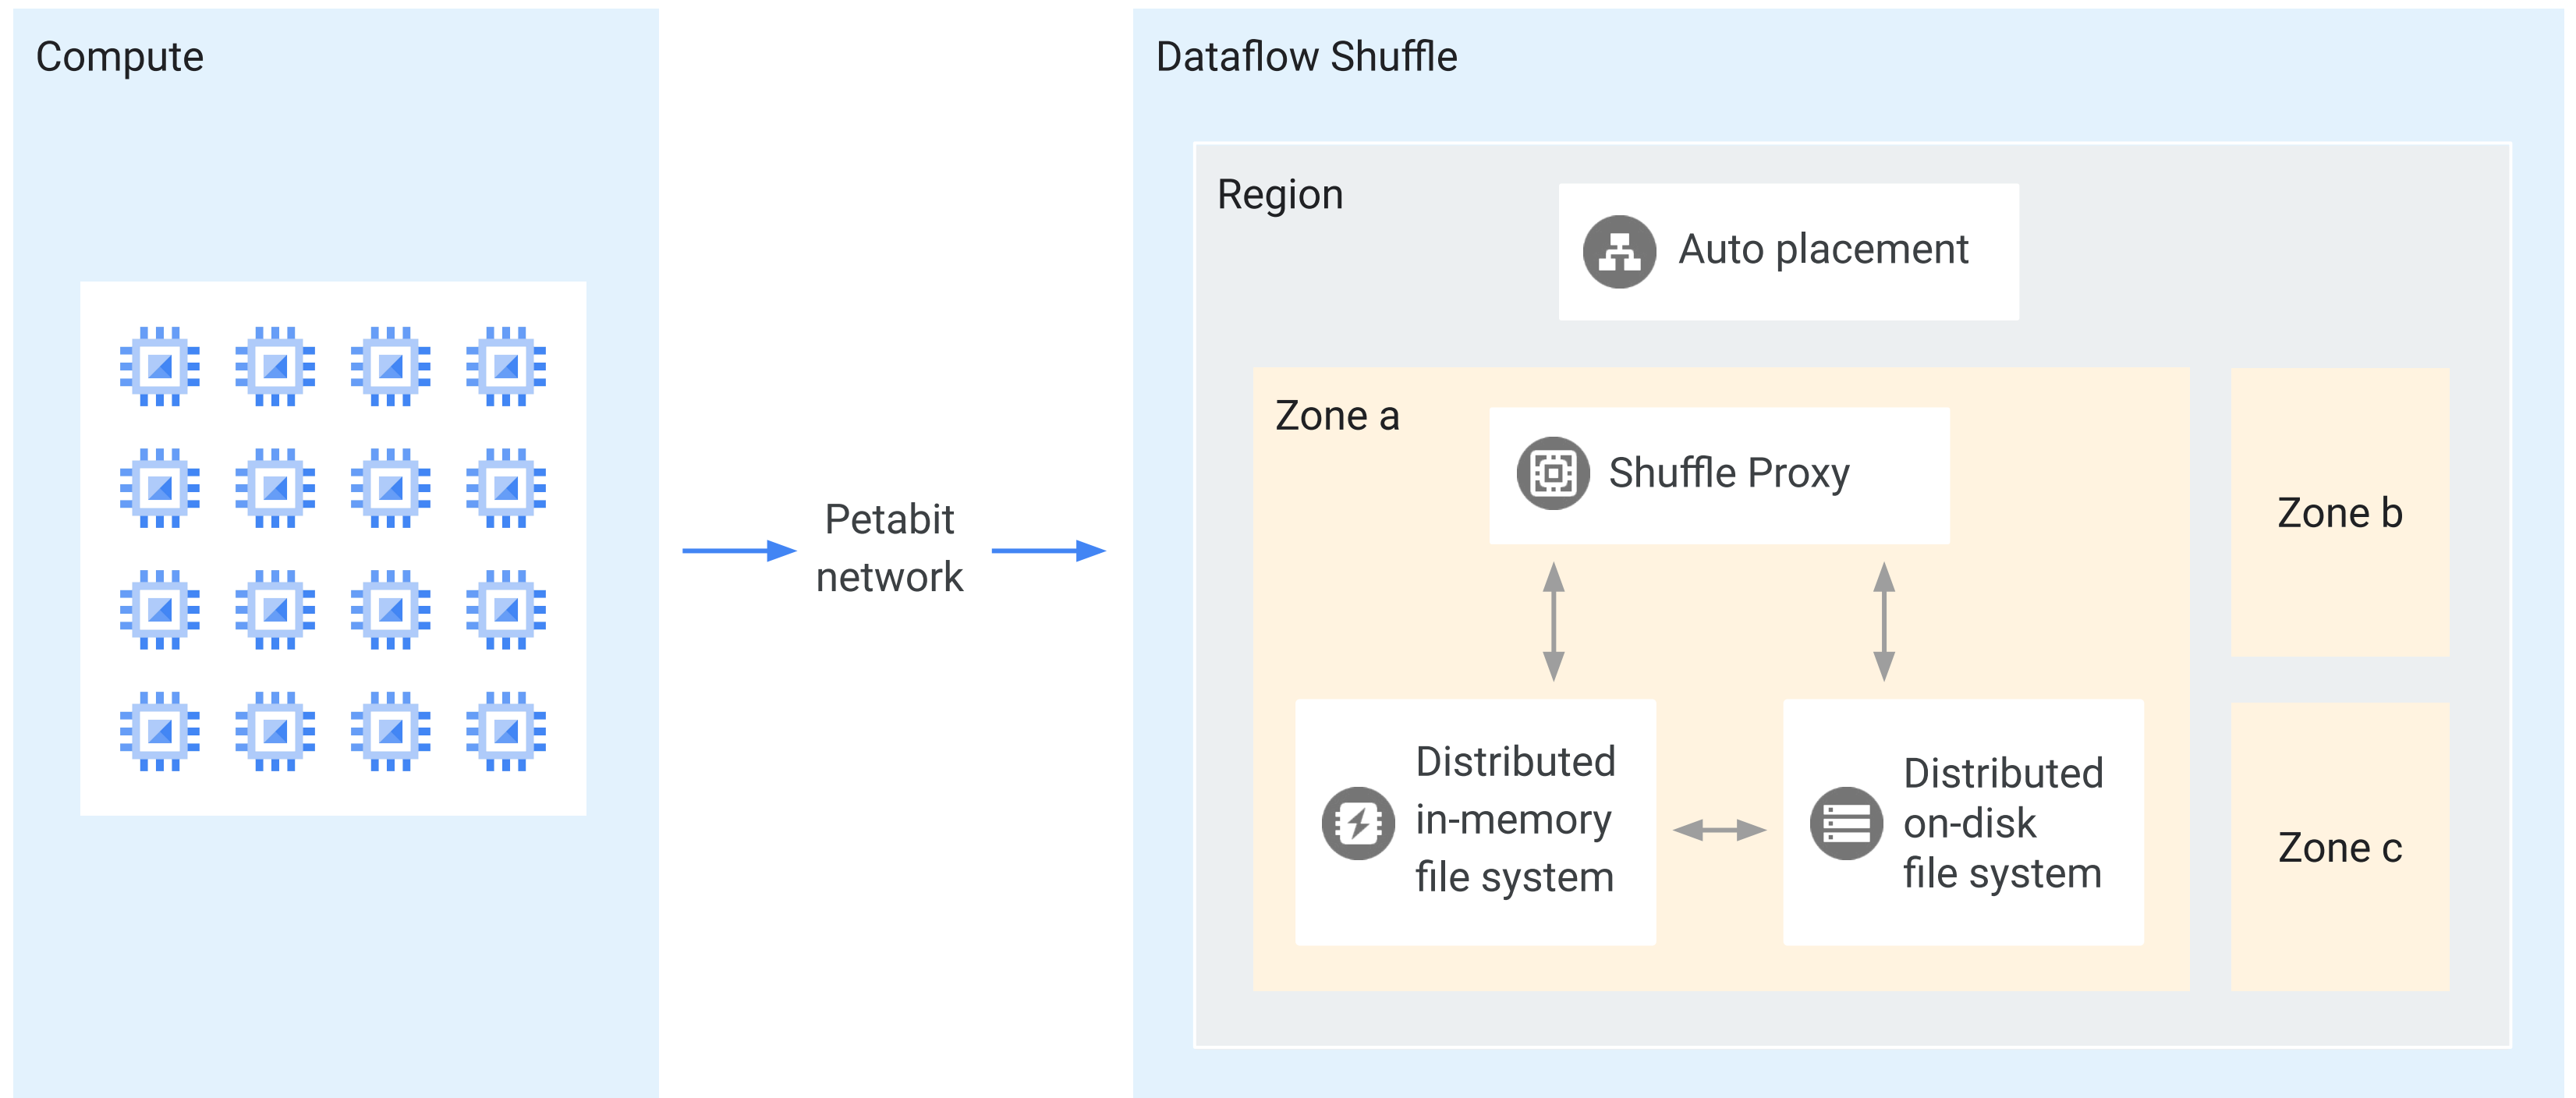


Performance

Agenda



Dataflow Shuffle service



Benefits of Cloud Dataflow Shuffle

- Faster execution time of batch pipelines.



Benefits of Cloud Dataflow Shuffle

- Faster execution time of batch pipelines.
- A reduction in consumed CPU, memory, and persistent disk storage resources on the worker VMs.



Benefits of Cloud Dataflow Shuffle

- Faster execution time of batch pipelines.
- A reduction in consumed CPU, memory, and persistent disk storage resources on the worker VMs.
- Better autoscaling.

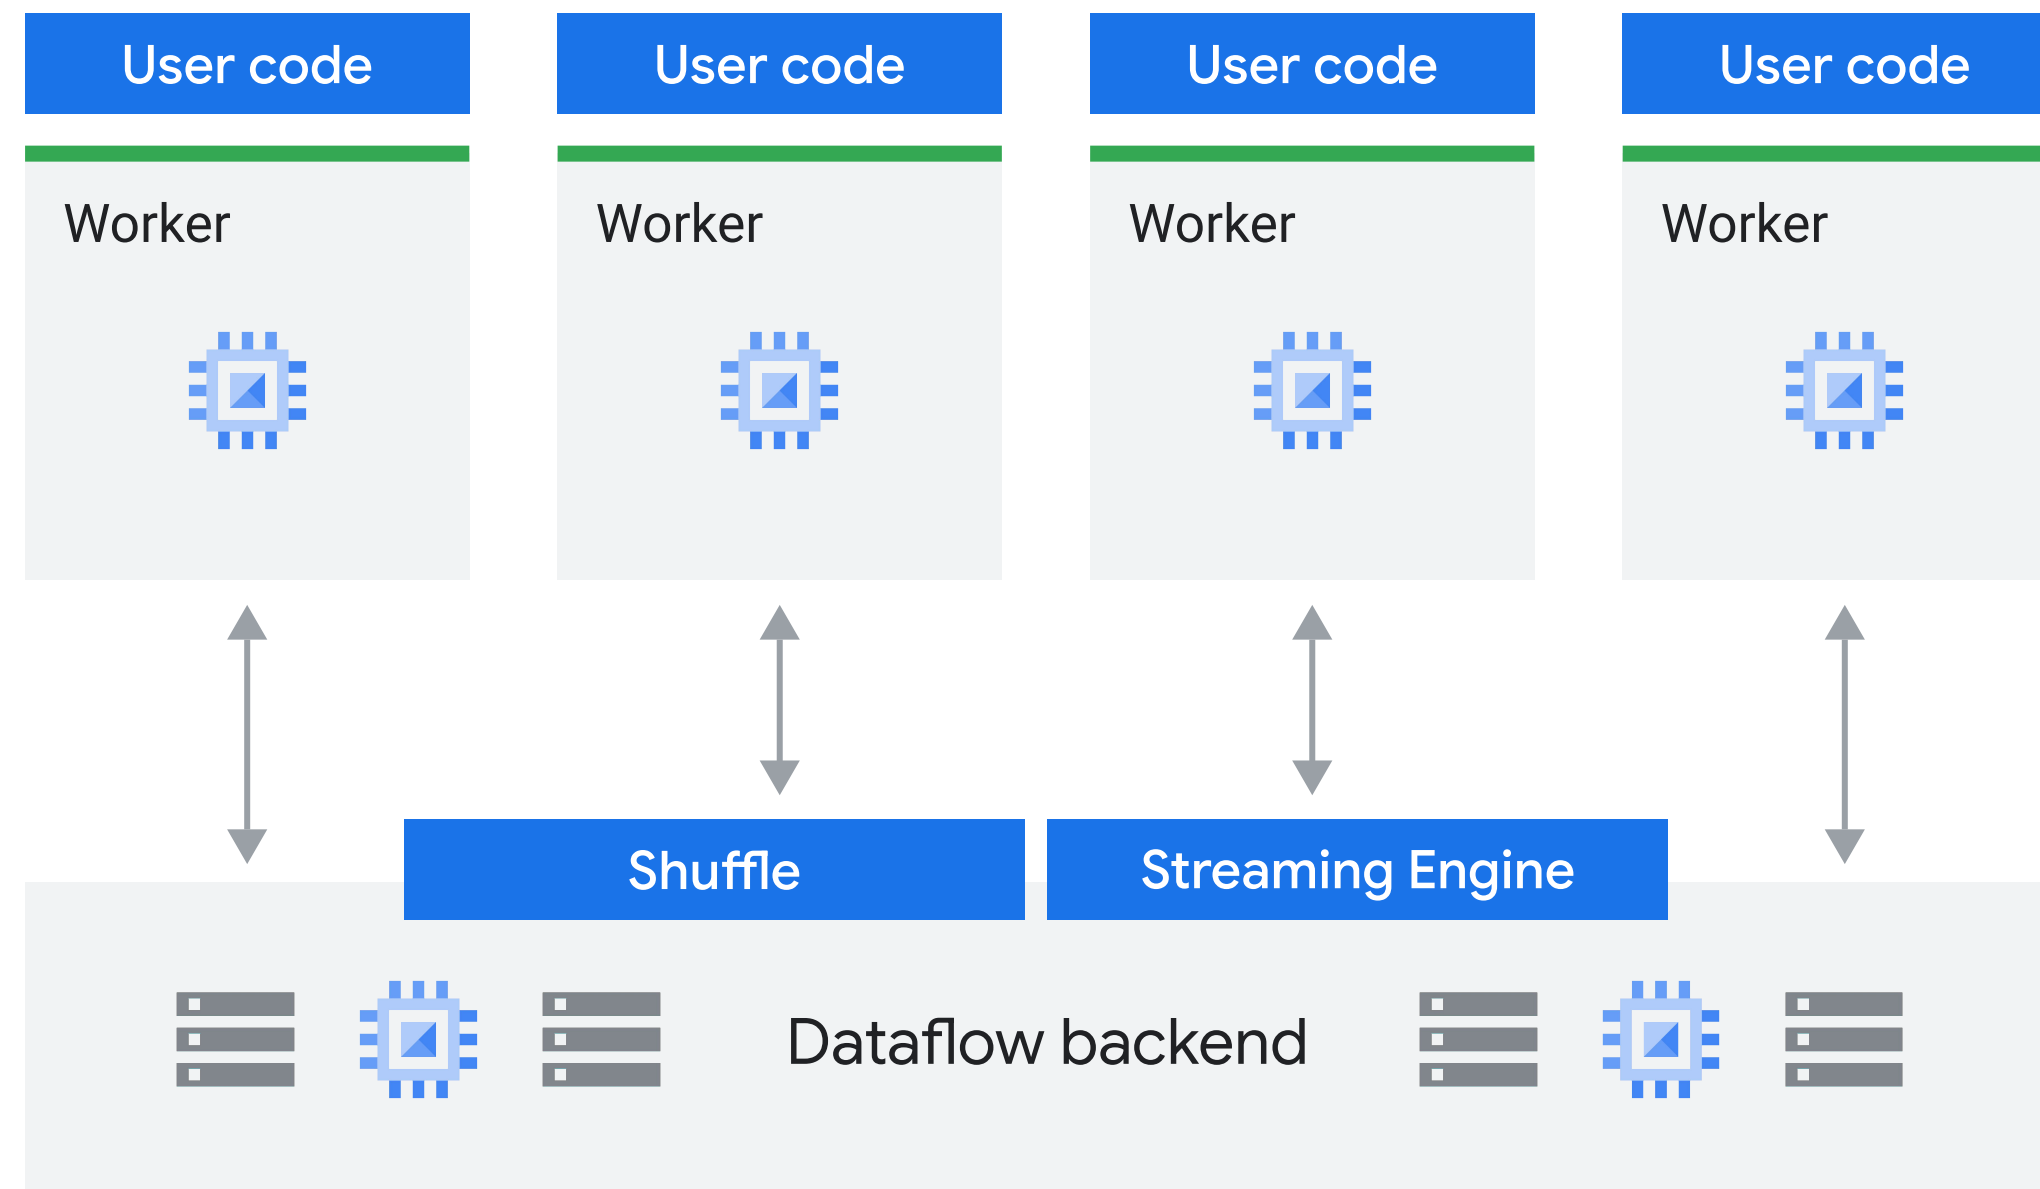


Benefits of Cloud Dataflow Shuffle

- Faster execution time of batch pipelines.
- A reduction in consumed CPU, memory, and persistent disk storage resources on the worker VMs.
- Better autoscaling.
- Better fault tolerance.



Dataflow shuffle and streaming service





Summary

