



# Sources and Sinks

Wei Hsia

Customer Engineer, Google Cloud





---

# Agenda

Course Intro

Beam Concepts Review

Windows, Watermarks, and Triggers

**Sources and Sinks**

Schemas

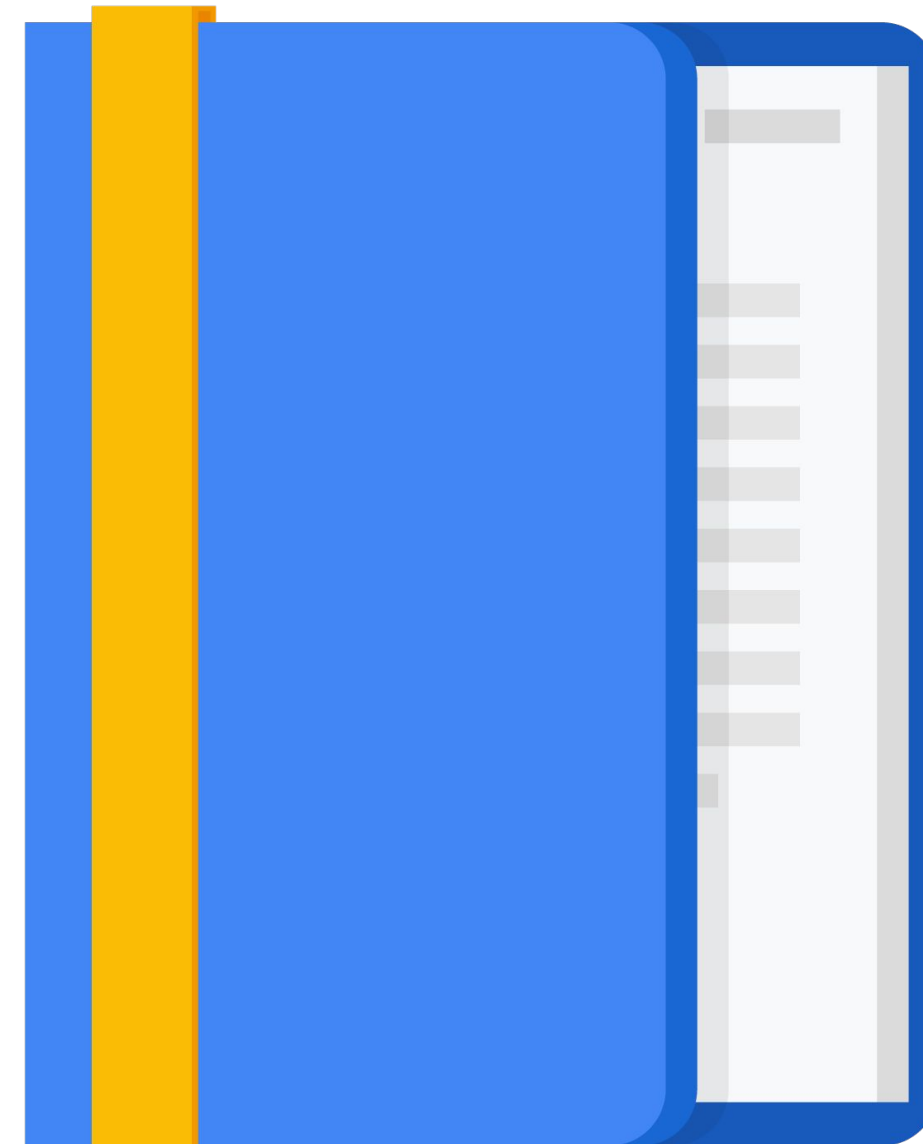
State and Timer

Best Practices

SQL and DataFrames


Beam Notebooks

Summary



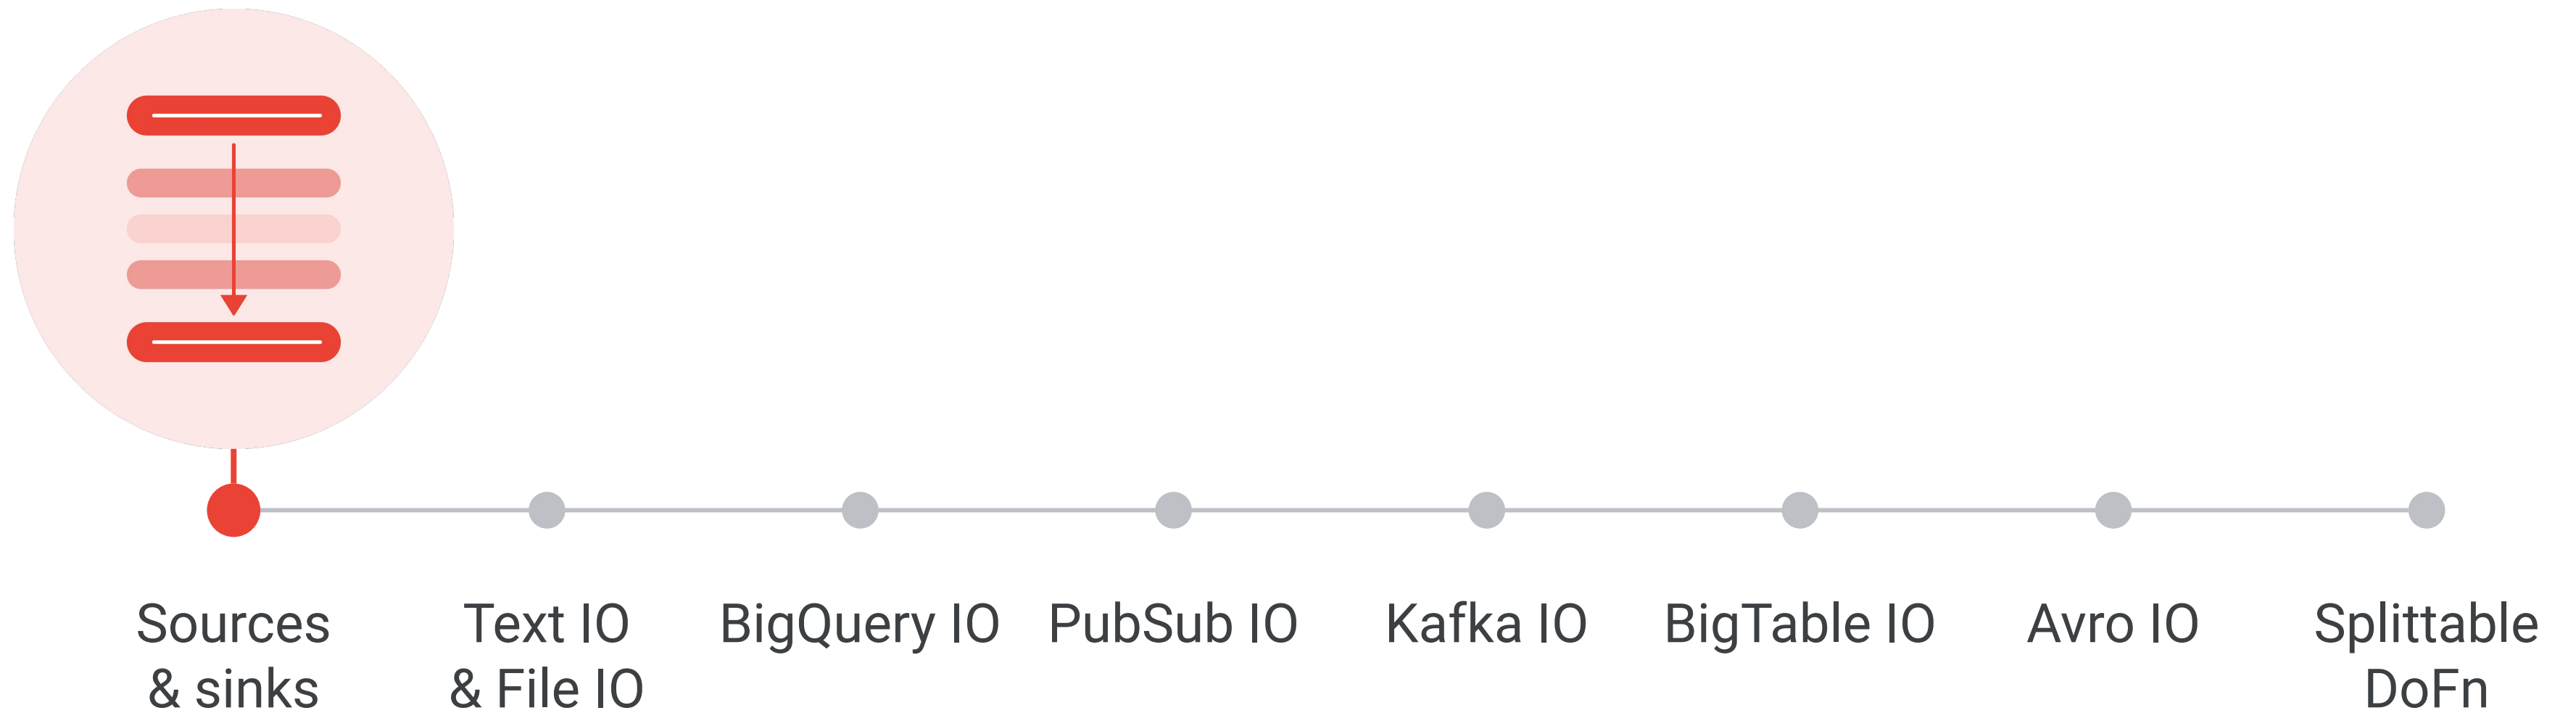
# Sources and Sinks

## Agenda



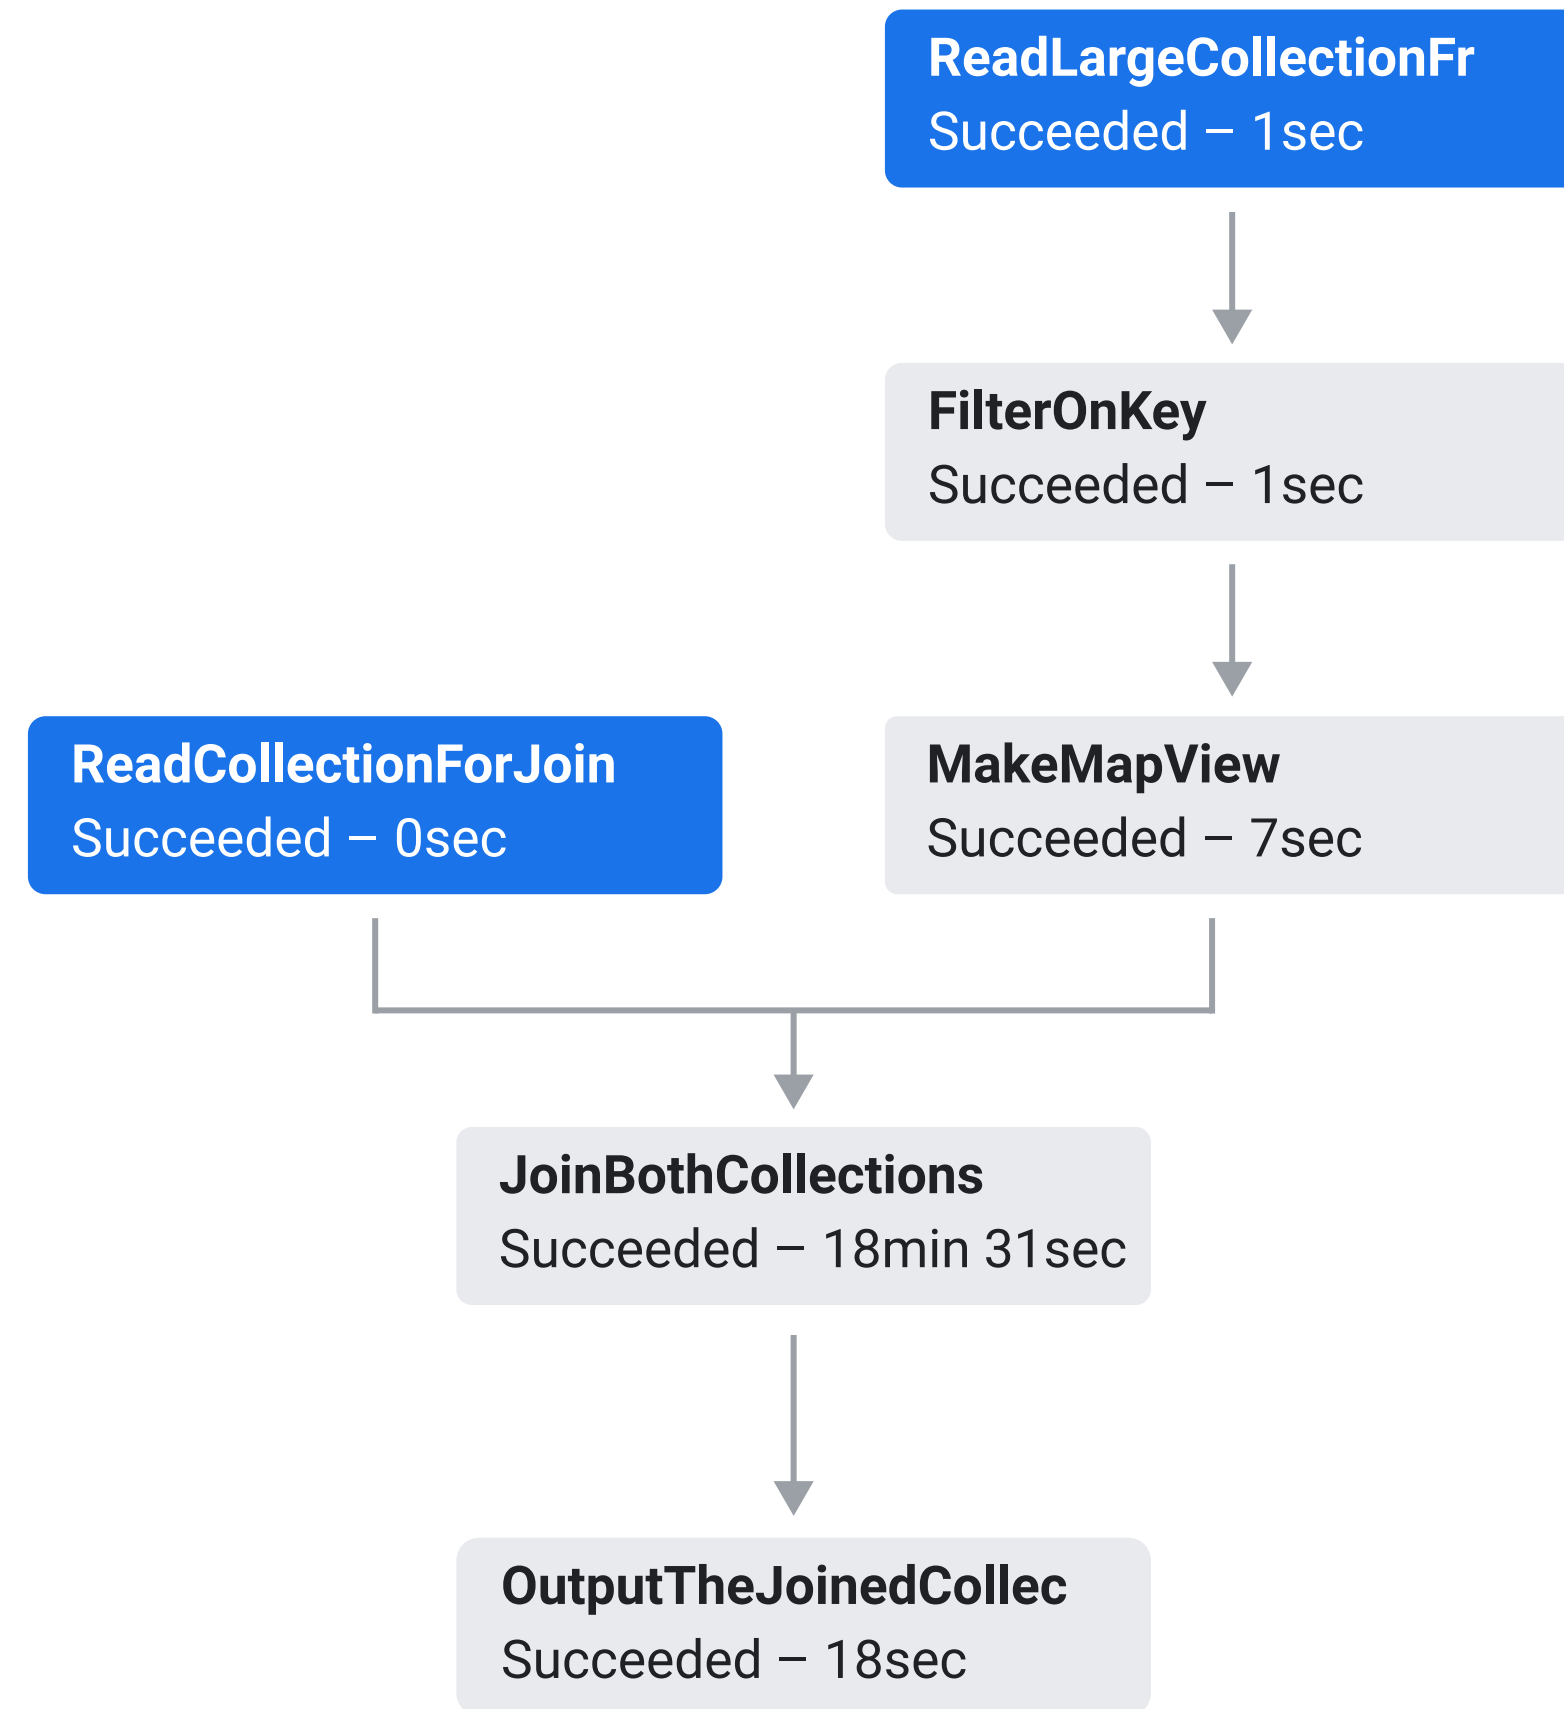
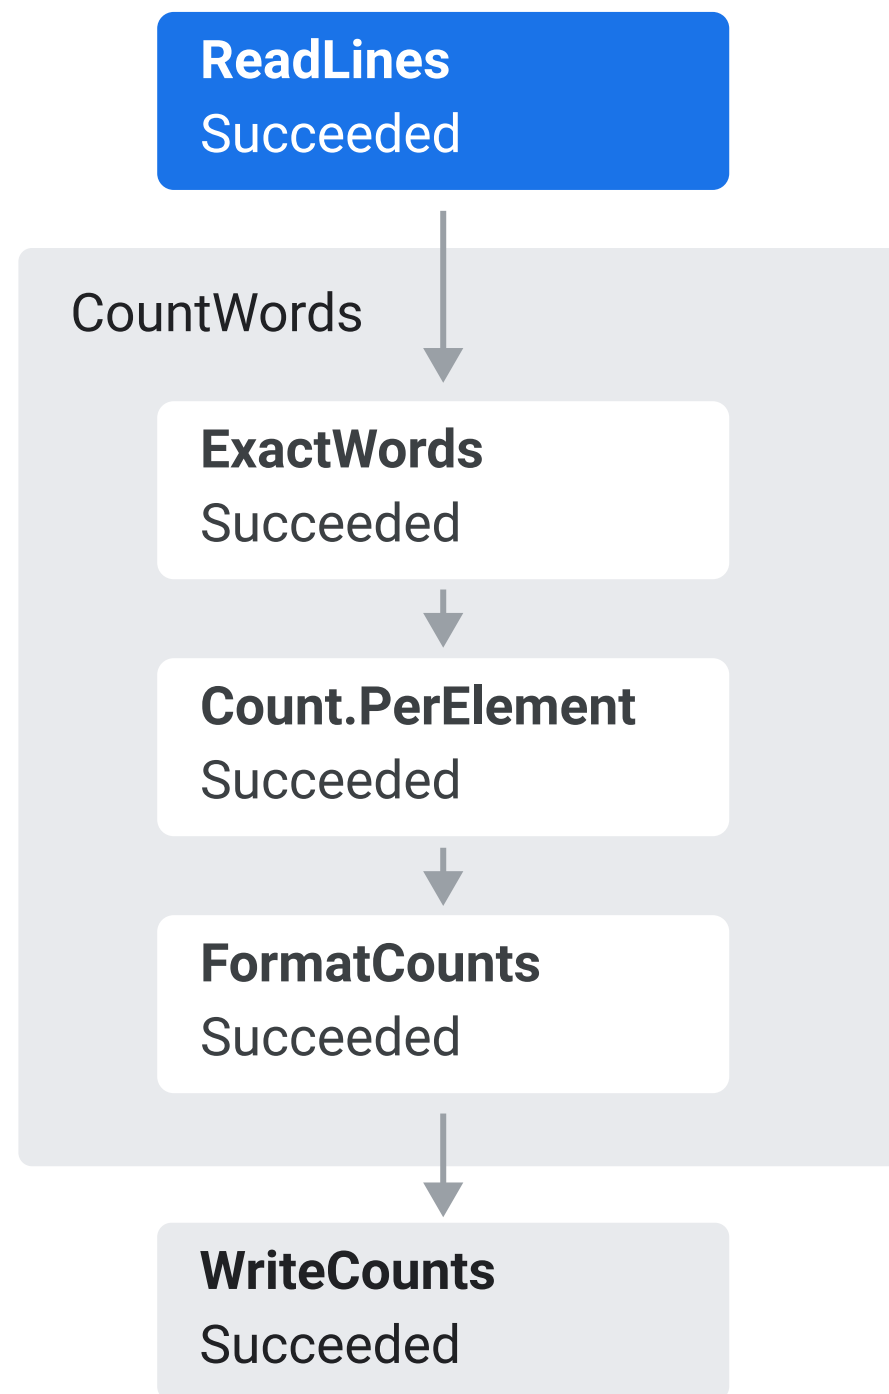
# Sources and Sinks

## Agenda



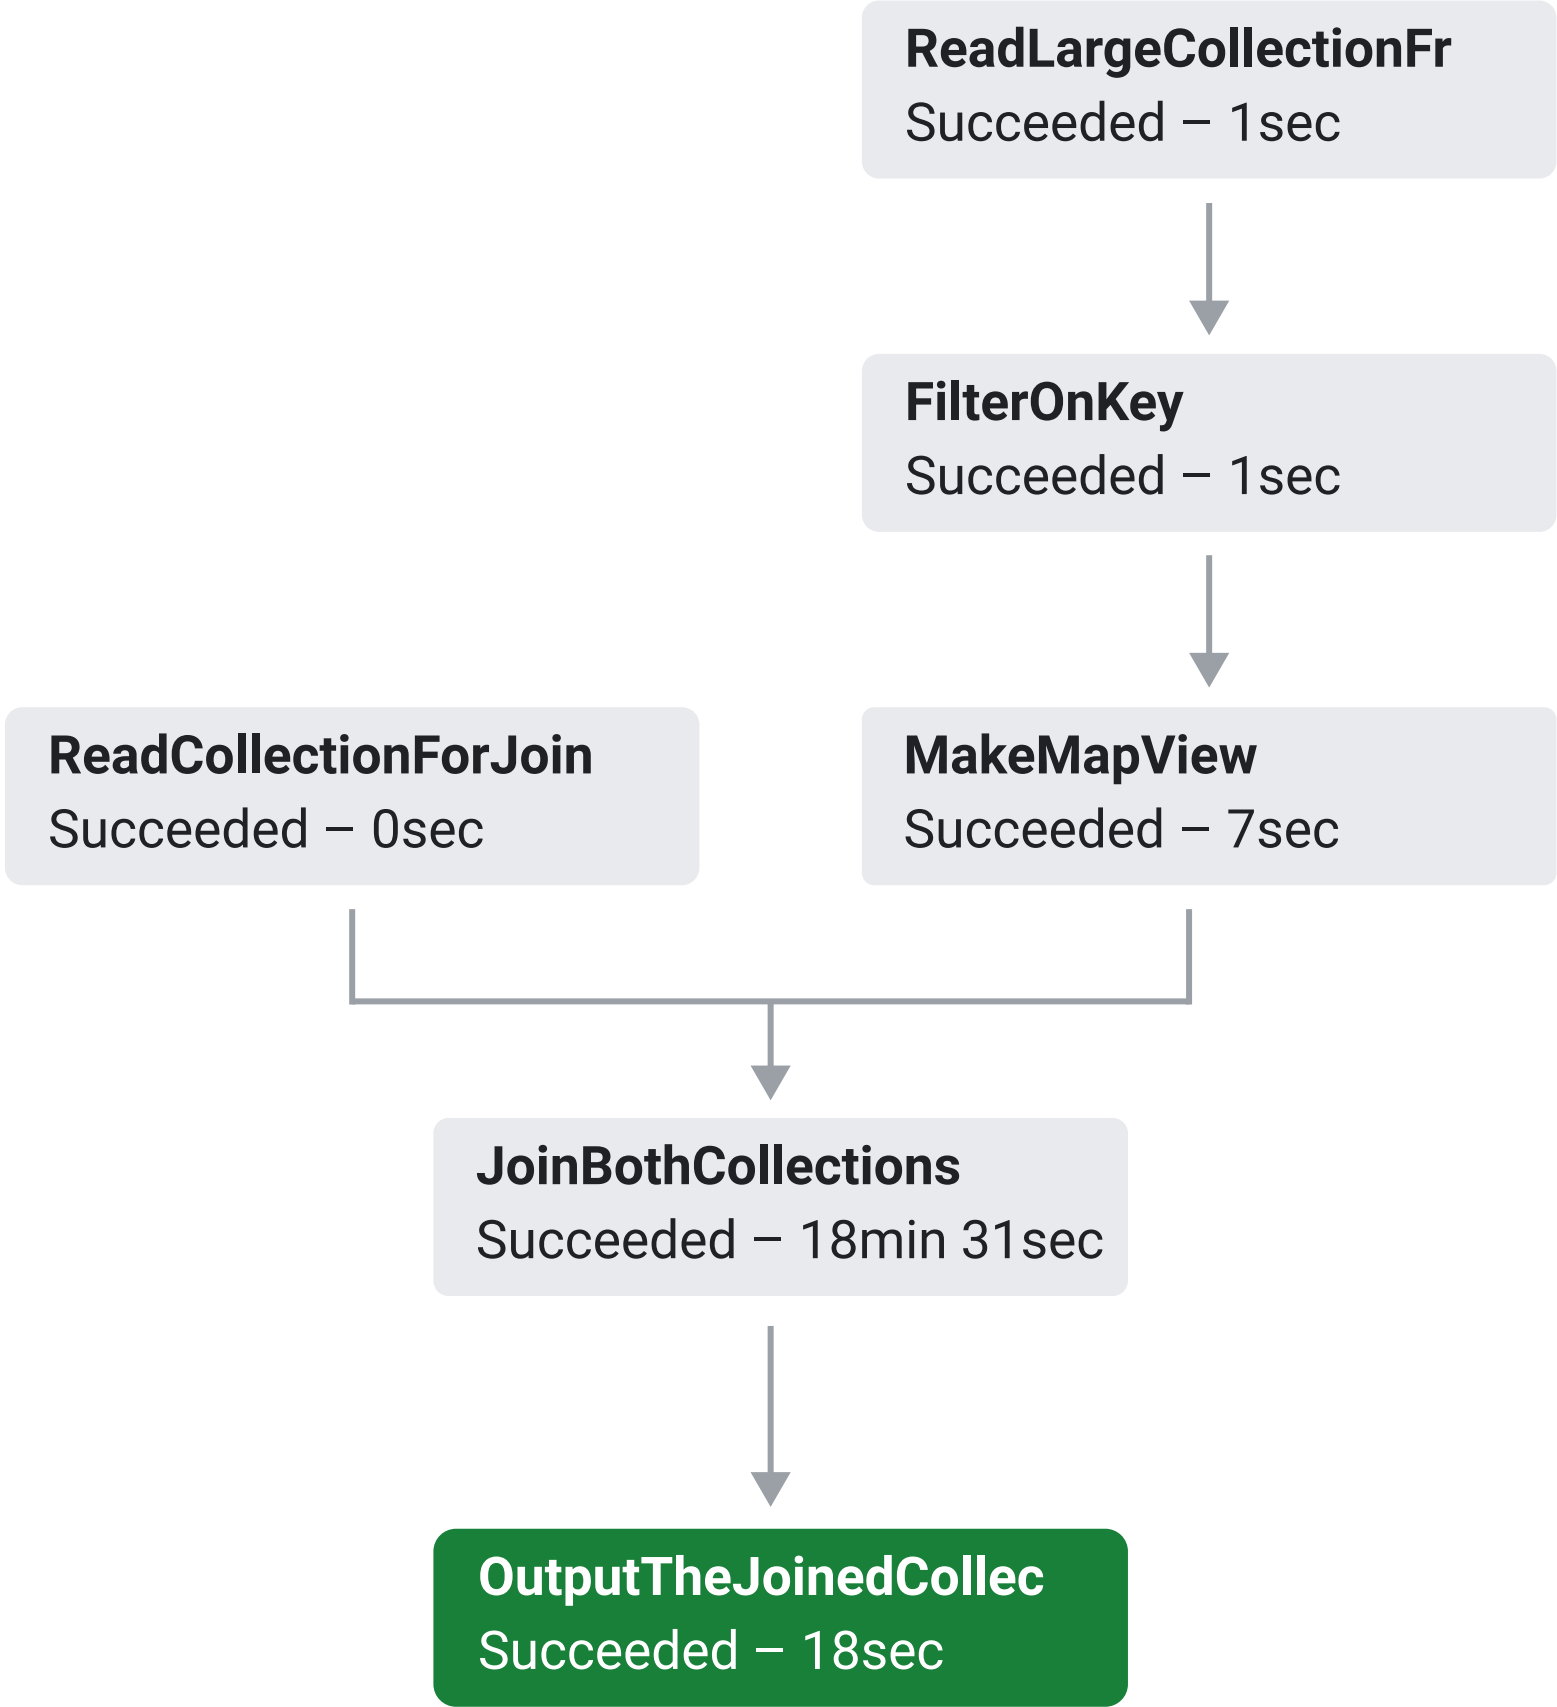
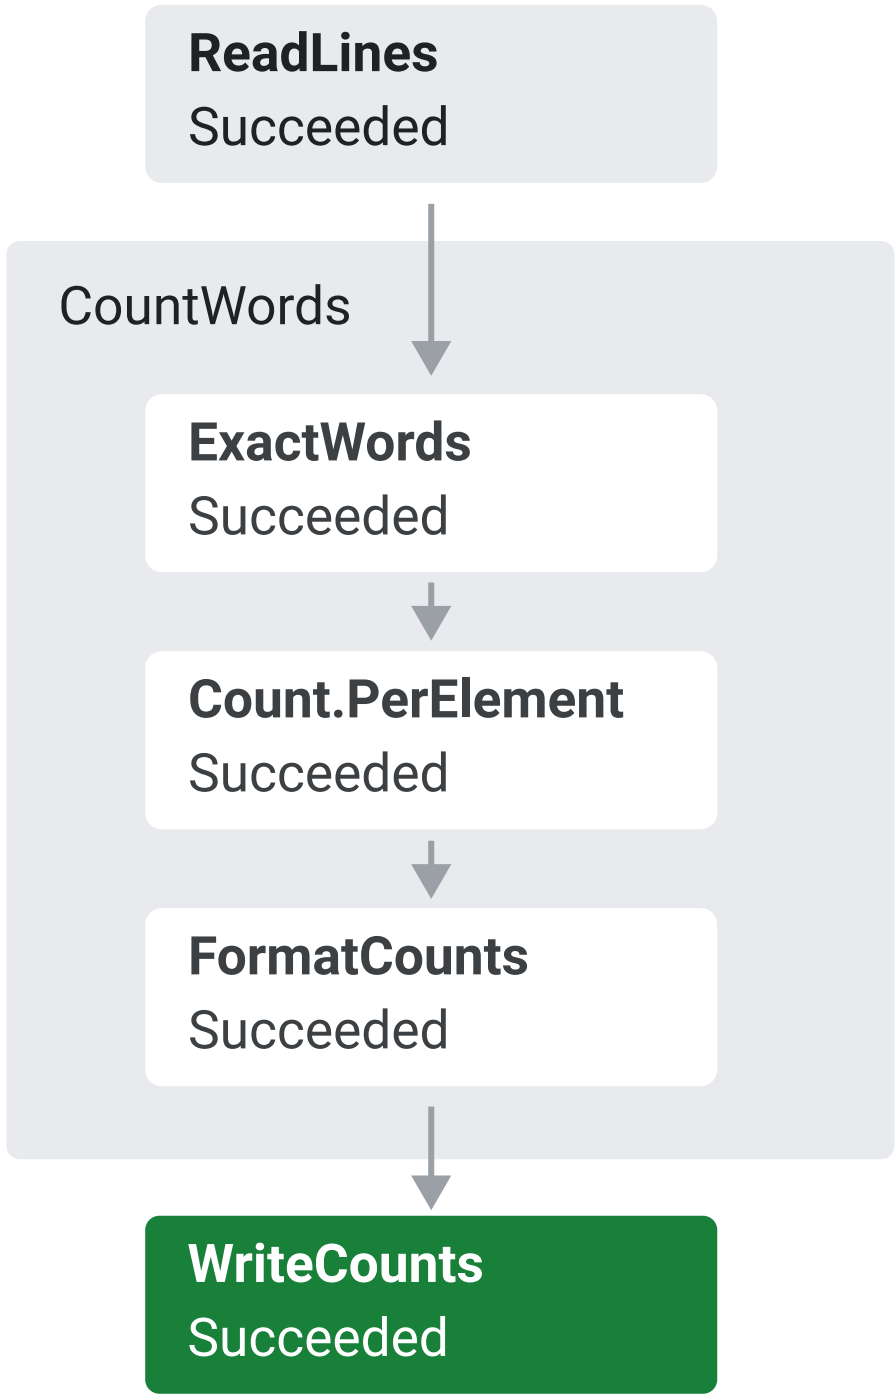
# Reading input data

Source



# Writing output data

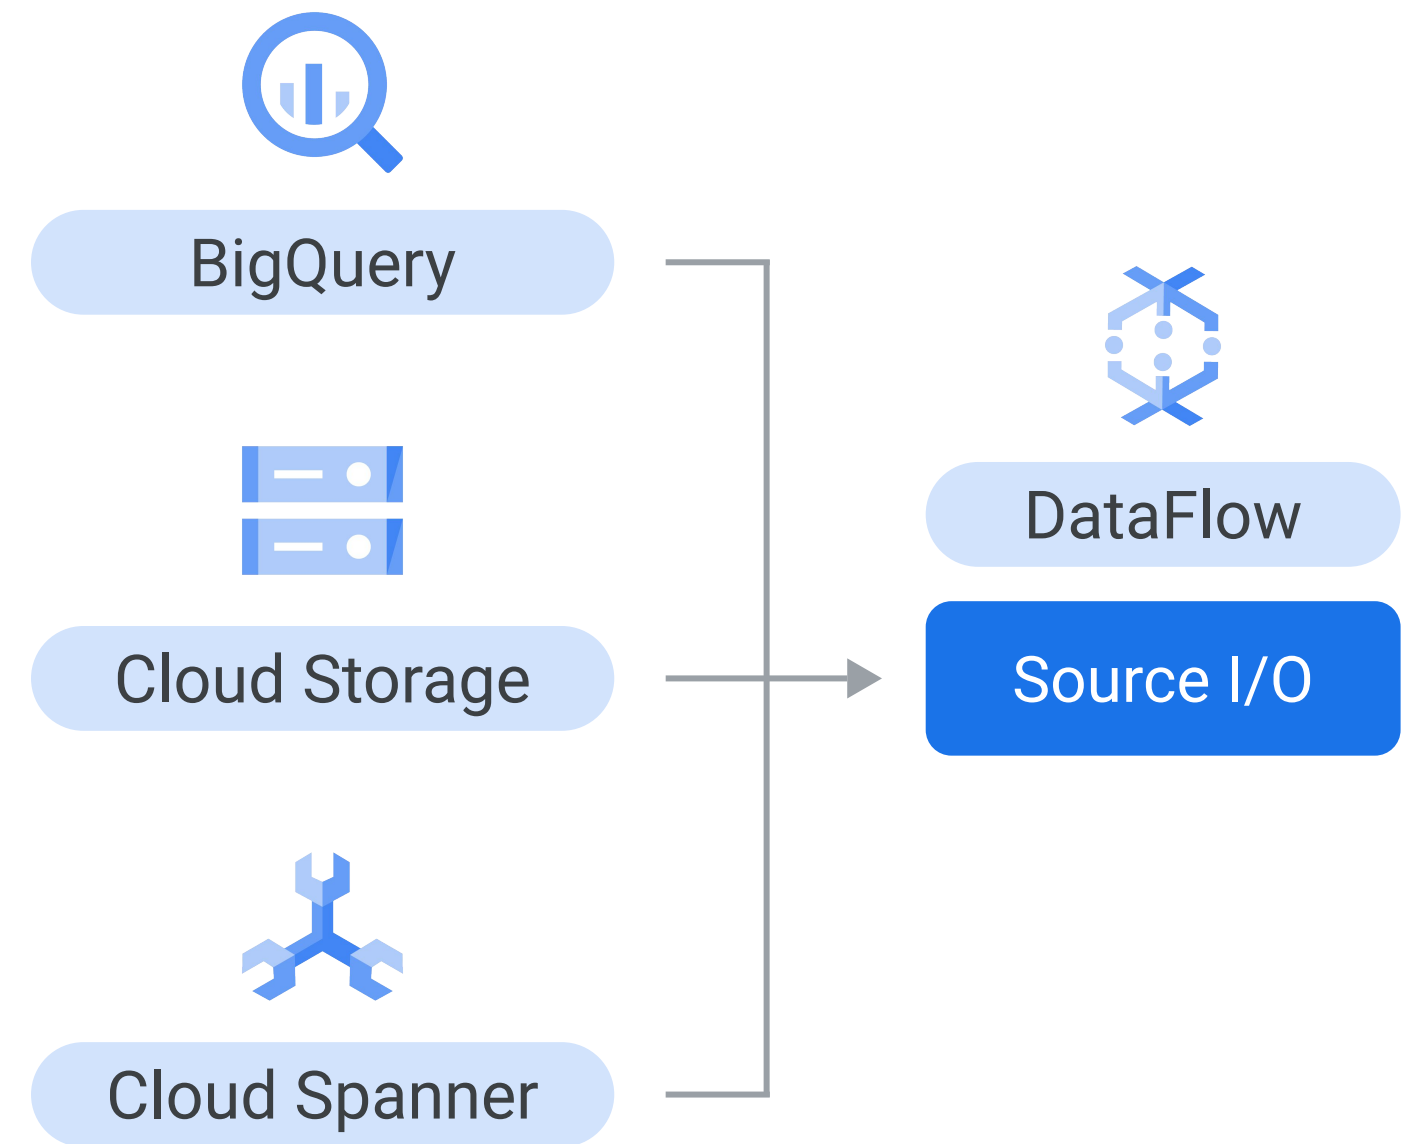
Sink



---

# Bounded sources

Sources that read a finite amount of input.



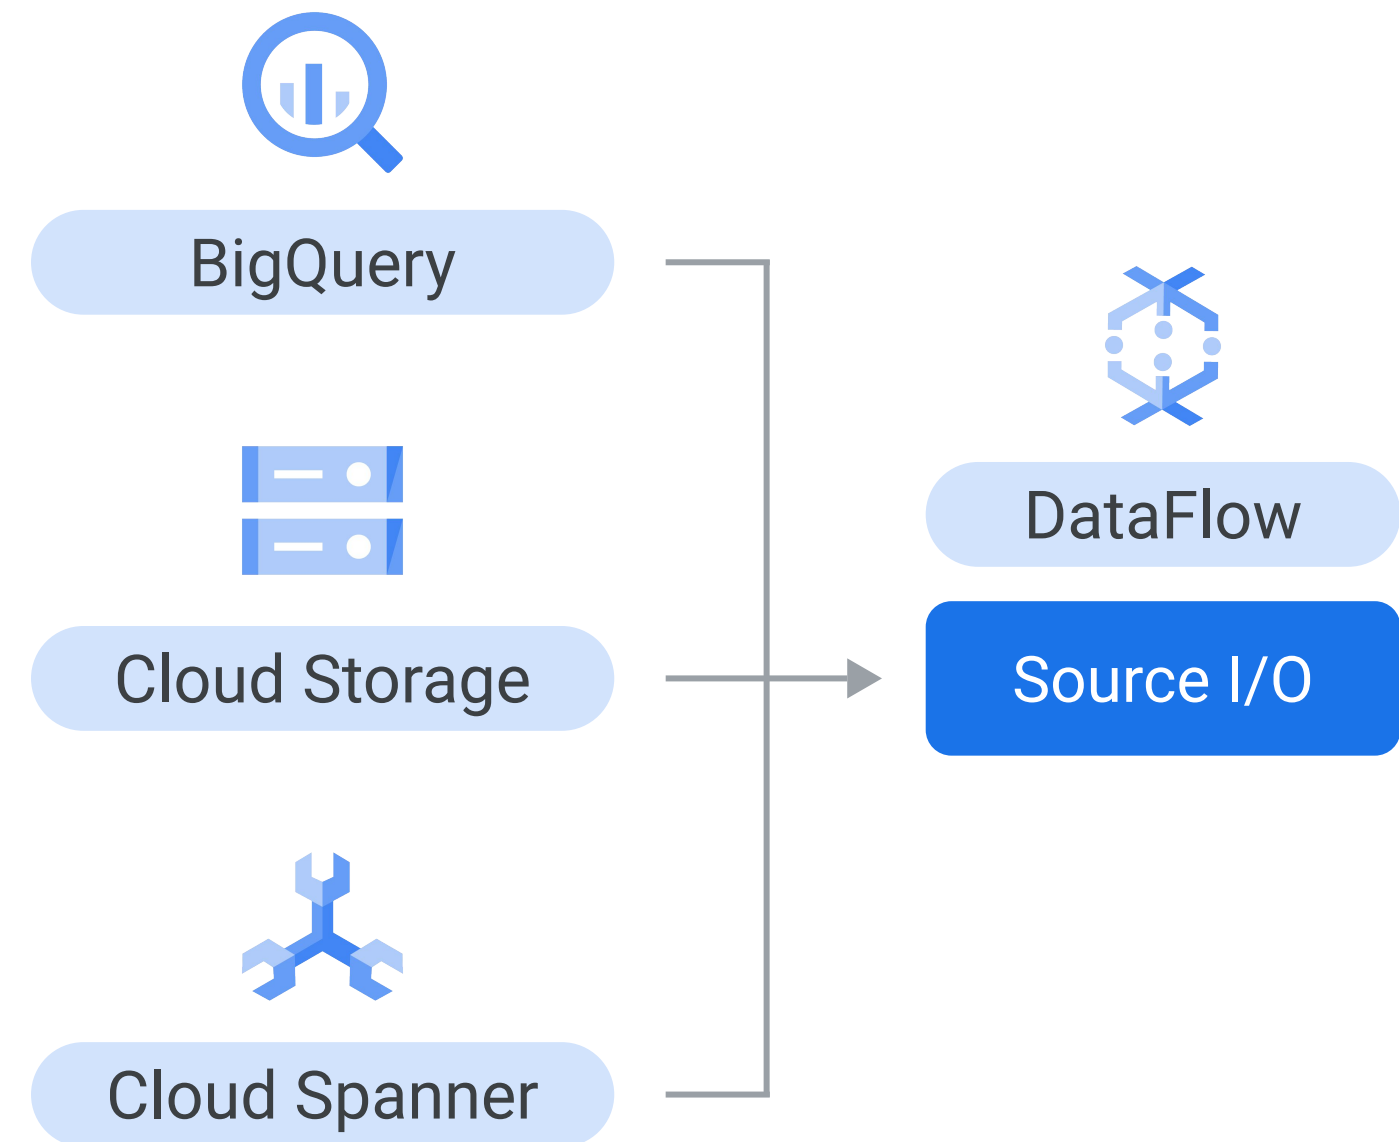


---

# Bounded sources

Sources that read a finite amount of input.

- Split the work of reading into smaller chunks, known as **bundles**.

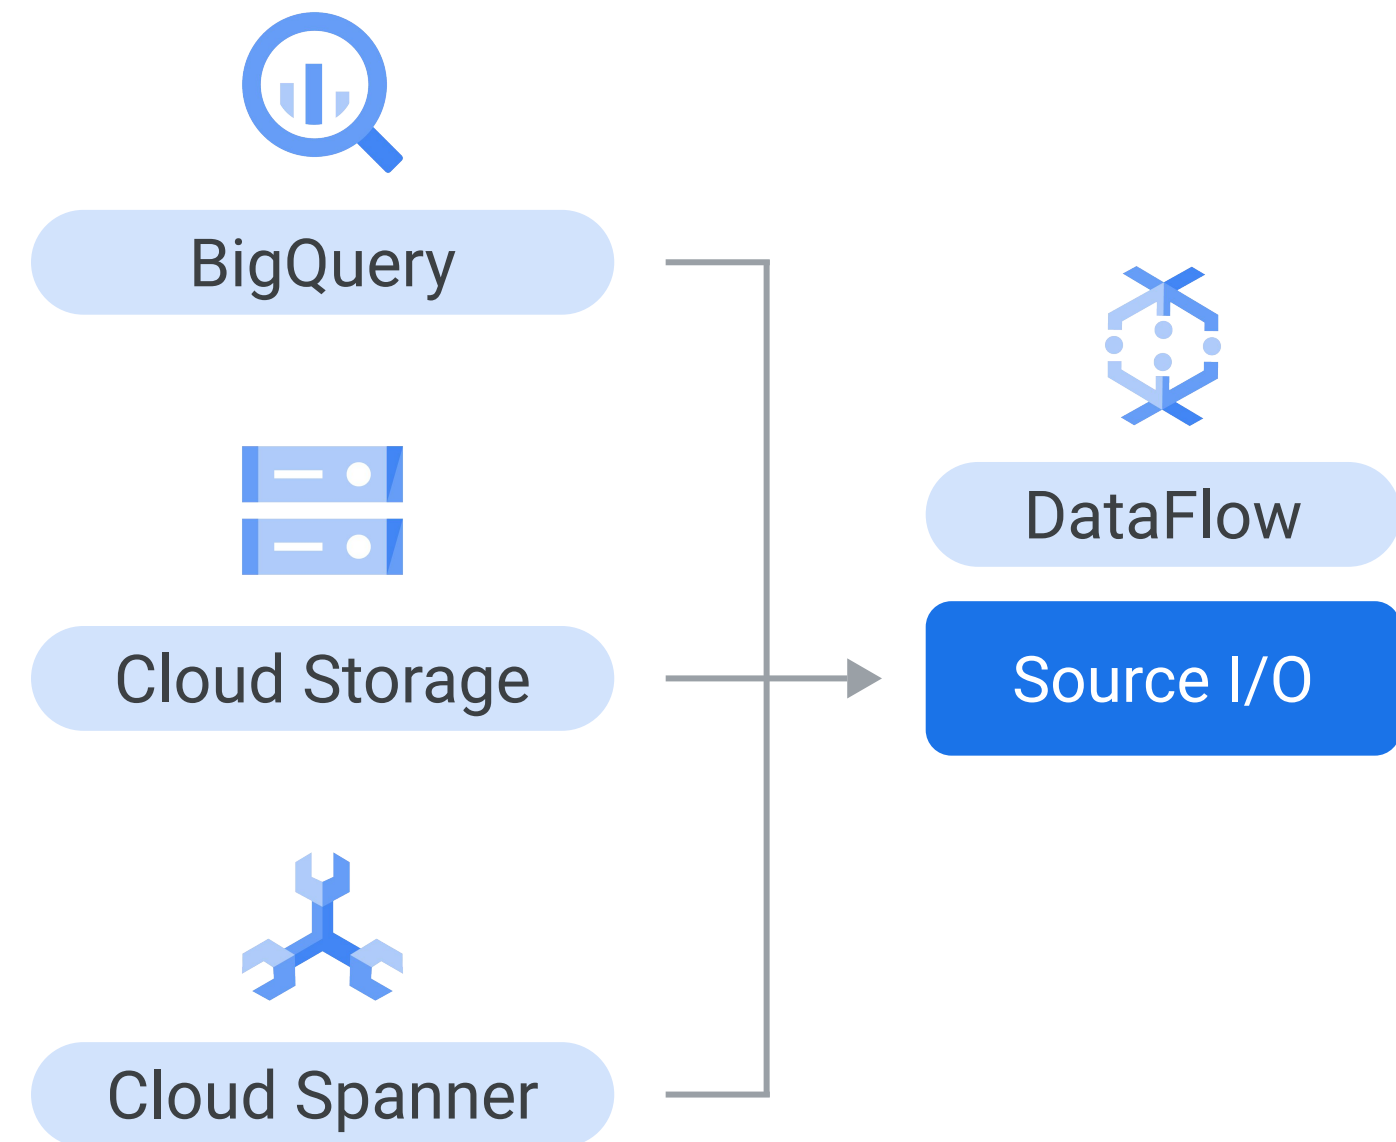


---

# Bounded sources

Sources that read a finite amount of input.

- Split the work of reading into smaller chunks, known as **bundles**.
- Provide **estimates** of progress to the service and number of bytes to be processed.

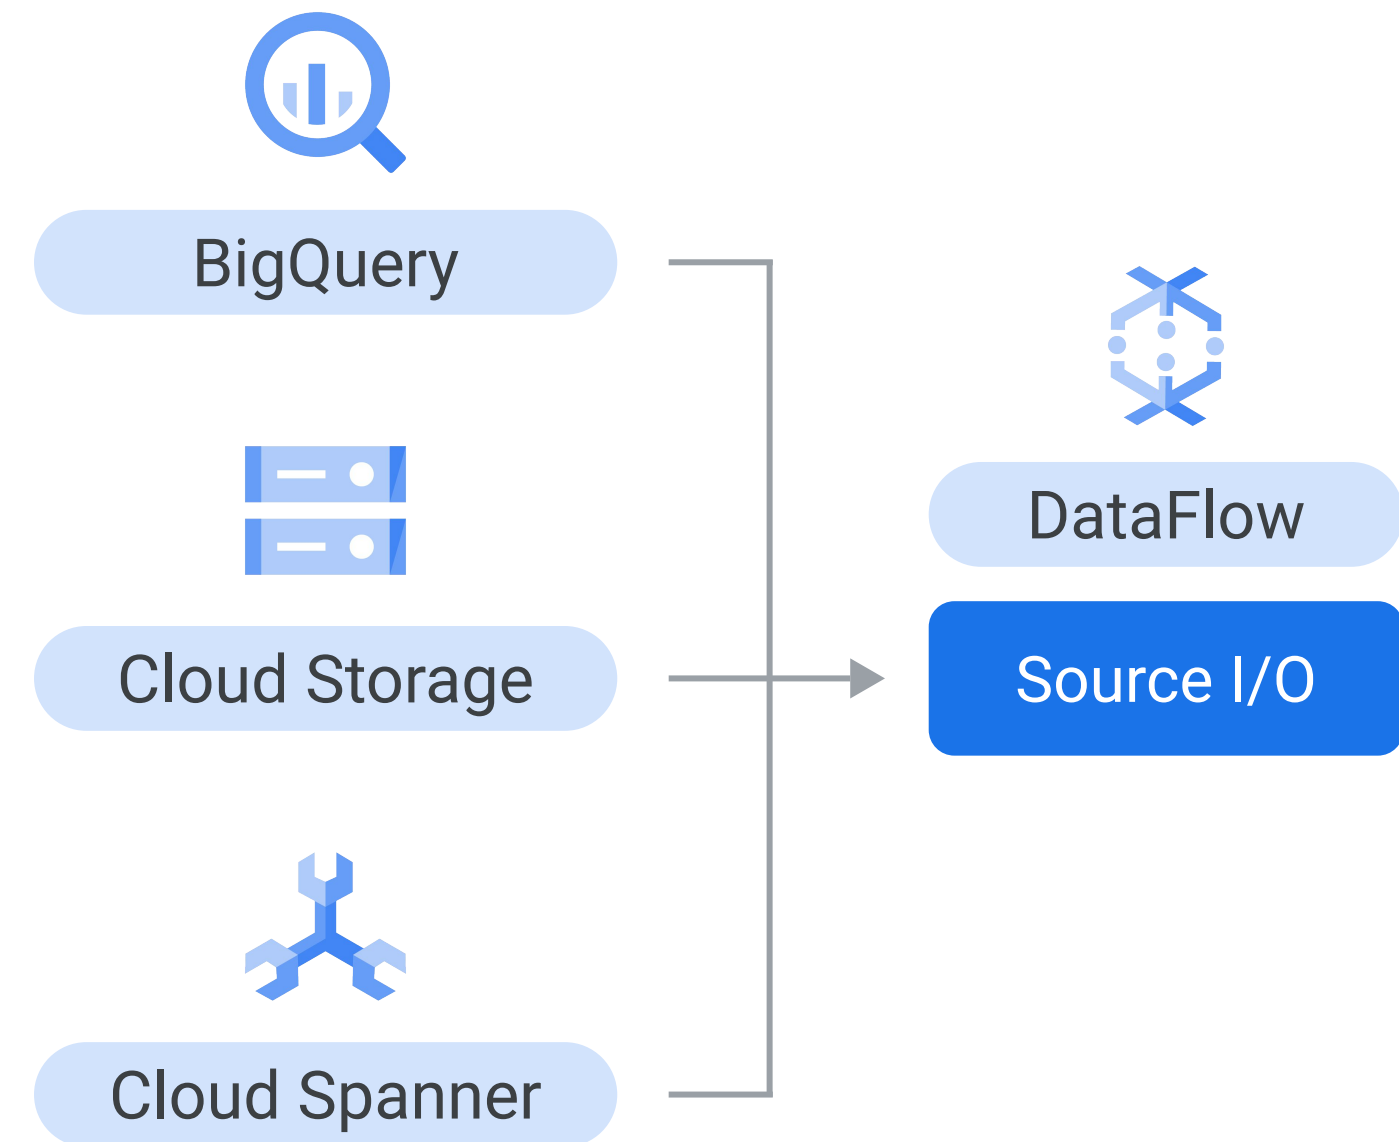


---

# Bounded sources

Sources that read a finite amount of input.

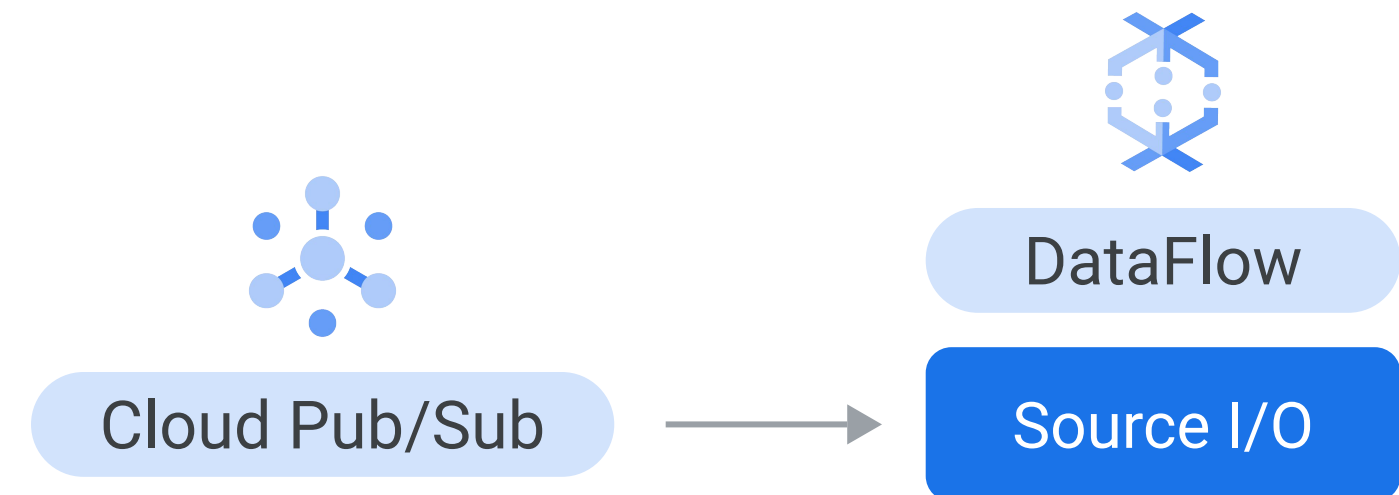
- Split the work of reading into smaller chunks, known as **bundles**.
- Provide **estimates** of progress to the service and number of bytes to be processed.
- Track if the units of work (bundles) can be broken down into smaller chunks for **dynamic work rebalancing** and carry out the **split operation** if needed.



---

# Unbounded sources

Source that reads an unbounded amount of input (e.g., streaming).

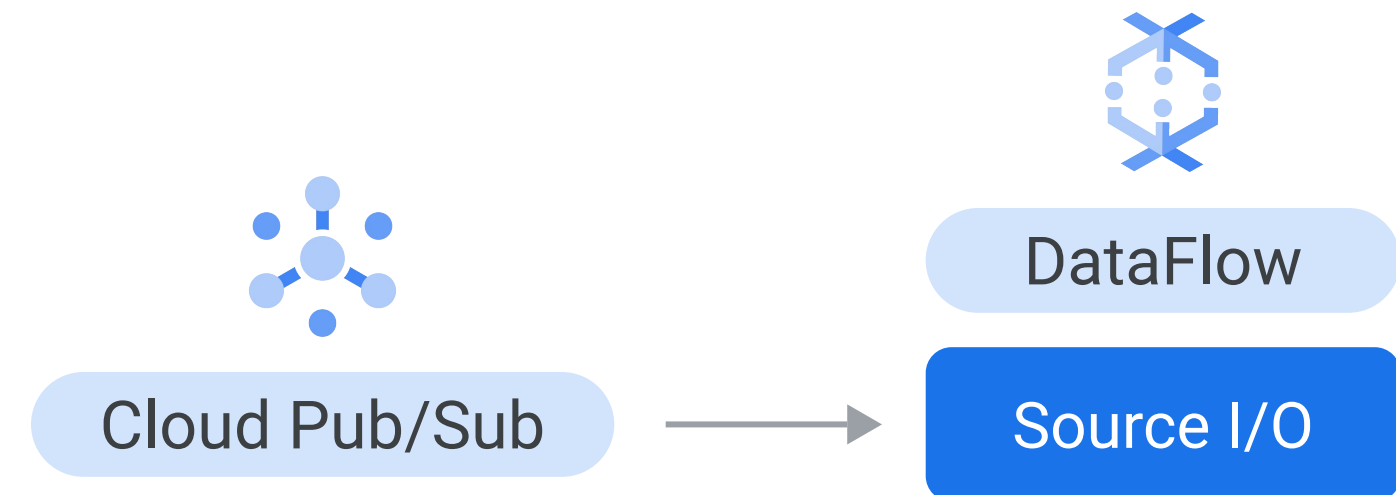


---

# Unbounded sources

Source that reads an unbounded amount of input (e.g., streaming).

- Allowing for the source to not re-read the same data by using **checkpoints**.

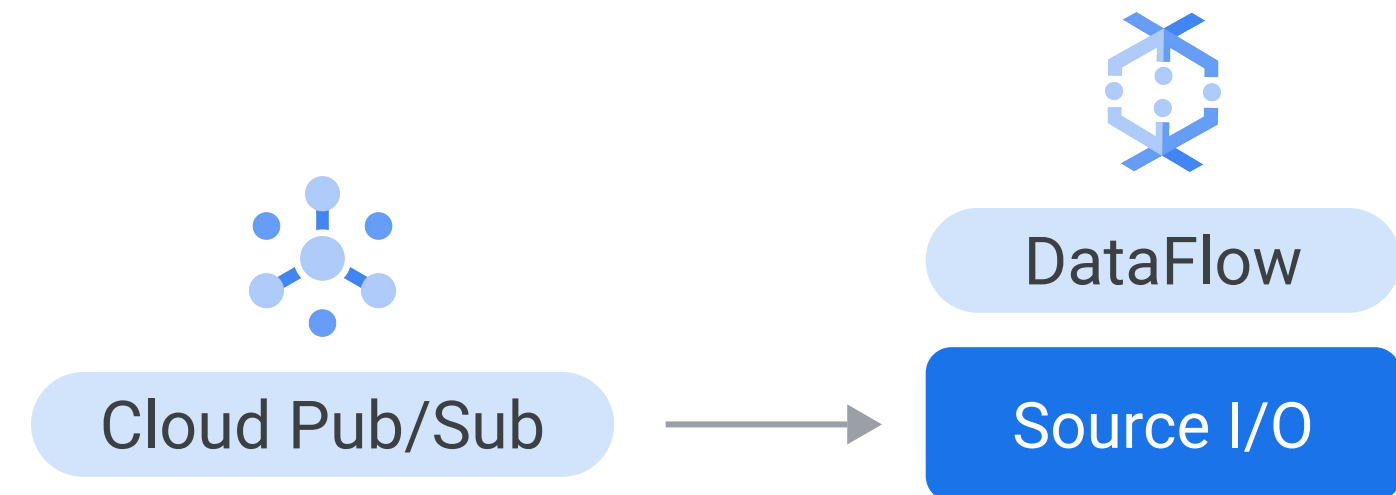


---

# Unbounded sources

Source that reads an unbounded amount of input (e.g., streaming).

- Allowing for the source to not re-read the same data by using **checkpoints**.
- Providing data to the service on what point in time the data is complete by using **watermarks**.

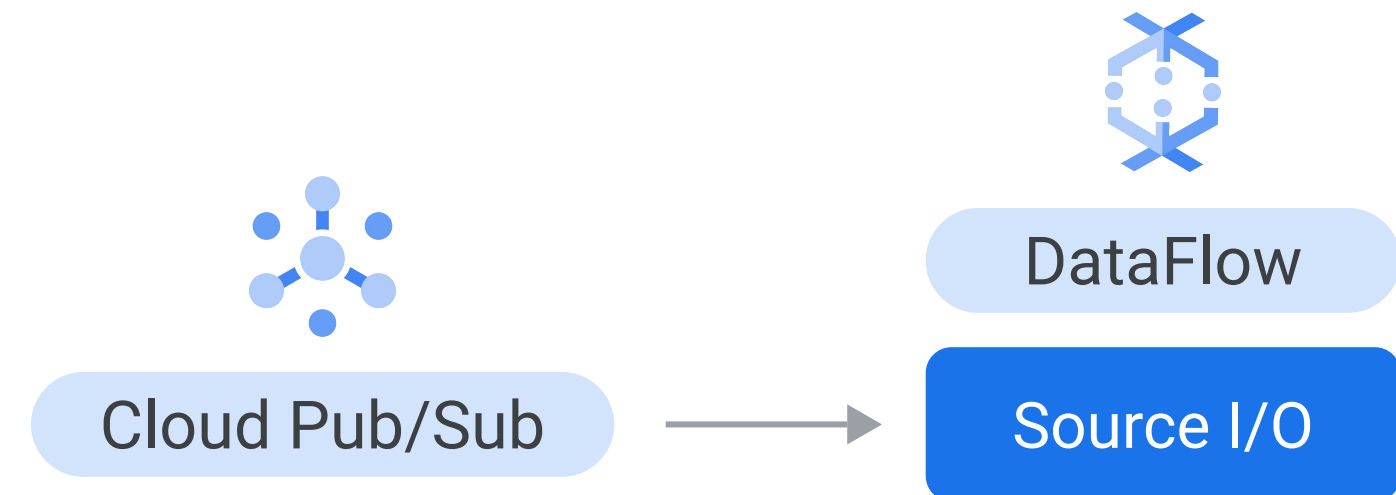


---

# Unbounded sources

Source that reads an unbounded amount of input (e.g., streaming).

- Allowing for the source to not re-read the same data by using **checkpoints**.
- Providing data to the service on what point in time the data is complete by using **watermarks**.
- Deduping the data with the option to make use of **Record IDs** from the unbounded source.



---

# Data sinks

Sinks are often “normal” PTransforms that write data to end systems.

Java

```
@AutoValue
    public abstract static class Write<T> extends
PTransform<PCollection<T>,WriteResult> {
```

Python

```
class WriteToPubSub(PTransform):
```



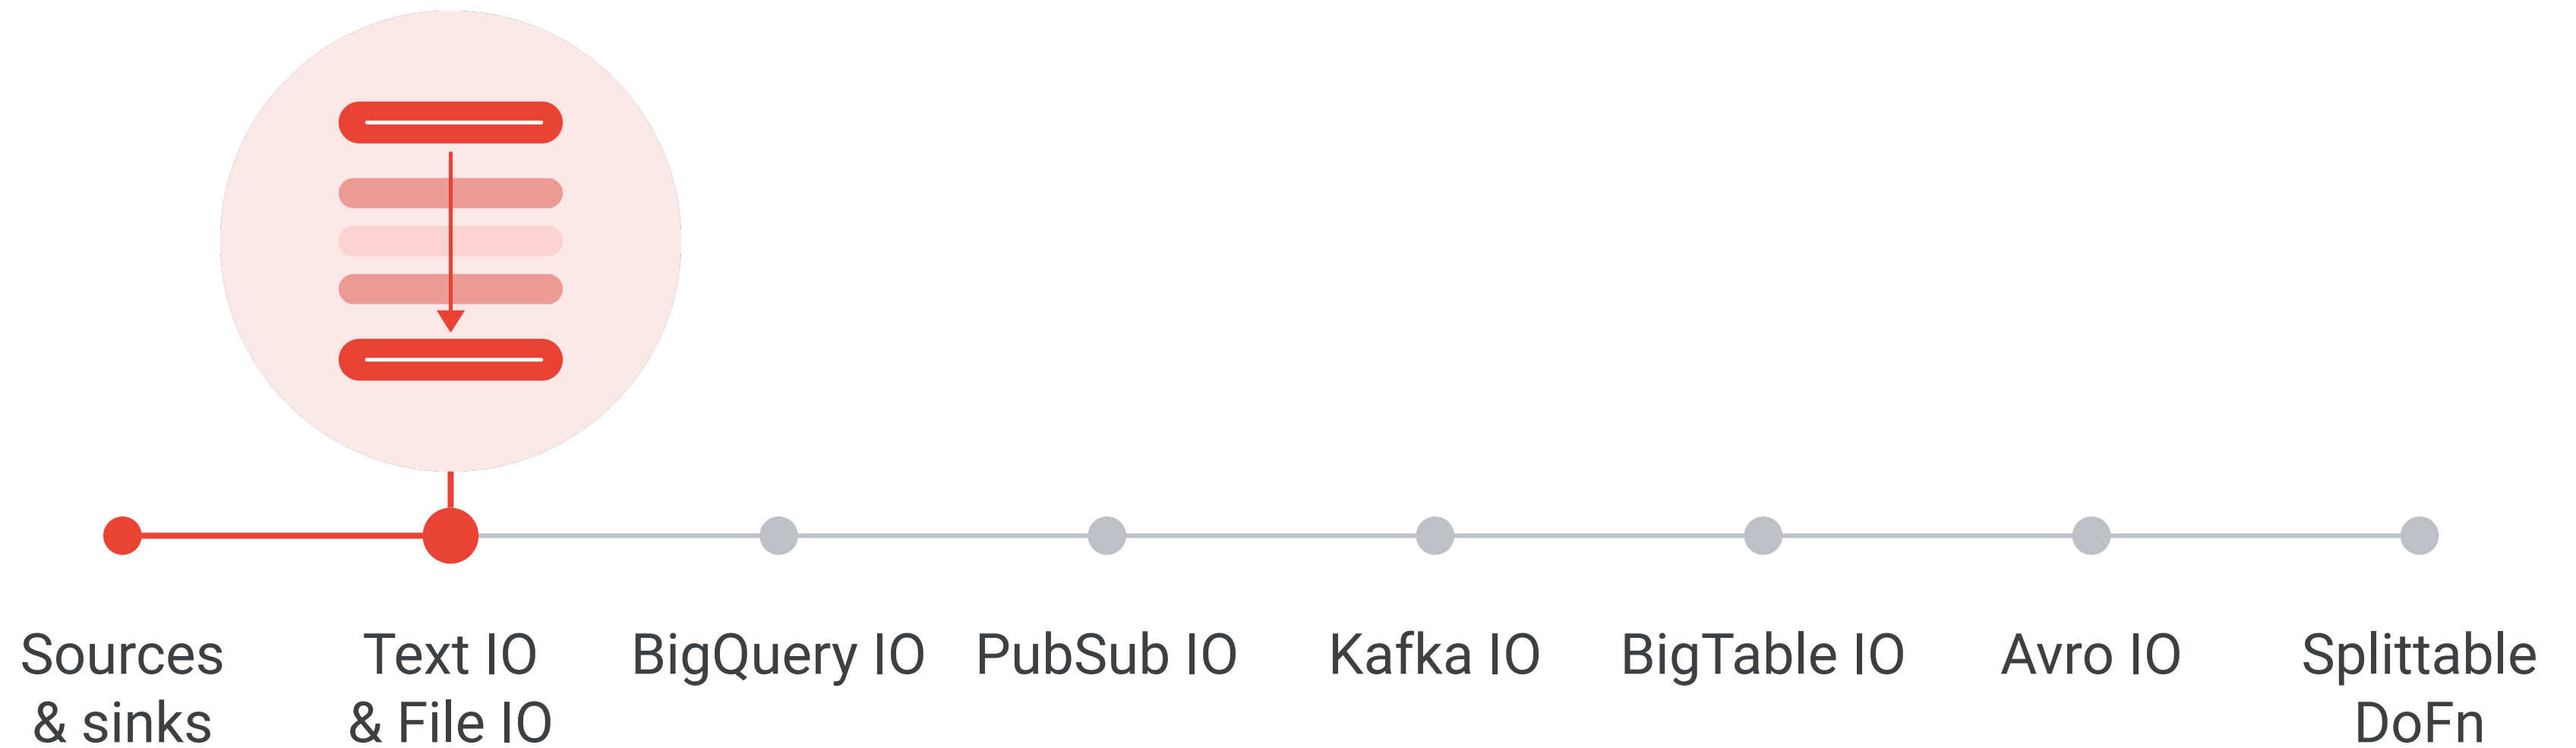
---

Updated list of Apache Beam's IO connectors

[s.apache.org/beam-io](https://s.apache.org/beam-io)

# Sources and Sinks

## Agenda



---

# Text IO reading

Java

```
Pipeline
  .apply(
    "Read from source",
    TextIO
      .read() ← Read method
      .from(options
        .getInputFilePattern()))
```

---

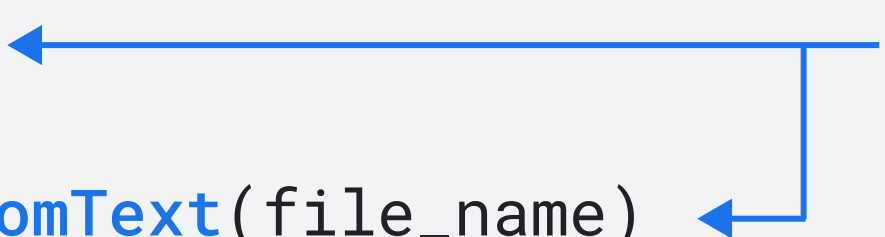
# Text IO reading

Python

```
pcol1 = (pipeline
  | 'Create' >> Create([file_name])
  | 'ReadAll' >> ReadAllFromText())
```

Read method

```
pcol12 = pipeline | 'Read' >> ReadFromText(file_name)
```



# File IO reading with filenames

Java

```
p.apply(
  FileIO
  .match() ← Match file pattern
  .filepath("hdfs://path/to/*.gz") ←
).apply(
  FileIO
  .readMatches().withCompression(Compression.GZIP))
.apply(
  ParDo.of(
    new DoFn<FileIO.ReadableFile, String>() {
      @ProcessElement
      public void process(
        @Element FileIO.ReadableFile file) {
        LOG.info("File Metadata resourceId is {} ",
          file.getMetadata().resourceId()); ← Access file metadata
      }
    }
  )));
```

---

# File IO reading with filenames

Python

```
with beam.Pipeline() as p:
    readable_files = (
        p
        | fileio.MatchFiles ( ' hdfs://path/to/*.txt ' ) ← Match file pattern
        | fileio.ReadMatches()
        | beam.Reshuffle())
    files_and_contents = (
        readable_files
        | beam.Map(lambda x: ( x.metadata.path , ← Access file metadata
                               x.read_utf8() )))
```

---

# File IO processing files as they arrive

Java

```
p.apply(
  FileIO.match()
    .filepath("...")
    . continuously ( ← Continuous file monitoring
      Duration.standardSeconds(30), ← Every 30 seconds
      Watch.Growth.afterTimeSinceNewOutput(
        Duration.standardHours(1)))) ← for 1 hour
```

---

# File IO processing files as they arrive

Python

```
with beam.Pipeline() as p:  
    readable_files = (  
        p  
        | beam.io.ReadFromPubSub(...)   
        ... #<Parse PubSub Message and Yield  
Filename>  
    )
```

Trigger with  
message queue

```
files_and_contents = (  
    readable_files  
    | ReadAllFromText())
```

Used parsed  
filename to read



---

# Contextual Text IO reading

Java

```
PCollection<Row> records =  
    p.apply(ContextualTextIO.read().from("..."));  
  
PCollection<Row> records2 =  
    p.apply(ContextualTextIO.read()  
        .from("/local/path/to/files/*.csv")  
        .withHasMultilineCSVRecords(true));  
  
PCollection<Row> records3 =  
    p.apply(ContextualTextIO.read()  
        .from("/local/path/to/files/*")  
        .watchForNewFiles(  
            Duration.standardMinutes(1),  
            afterTimeSinceNewOutput(  
                Duration.standardHours(1))));
```

---

# Text IO writing

Java

```
csv.apply(  
    "Write to storage",  
    TextIO  
    .write()  
    .to(Options  
        .getTextWritePrefix()  
        .withSuffix(".csv"));
```

Write method



---

# Text IO writing

Python

```
transformed_data  
| 'write' >> WriteToText( ← Write method  
    known_args.output, coder=JsonCoder()))
```

---

# Text IO writing with dynamic destinations

Java

```
PCollection<BankTransaction> transactions = ...;

transactions.apply(FileIO.<TransactionType,
Transaction>writeDynamic() ← Dynamic destination
    .by(Transaction::getTypeName)
    .via(tx -> tx.getTypeName().toFields(tx),
        type -> new CSVSink(type.getFieldNames()))
    .to("../path/to/")
    .withNaming(type -> defaultNaming(
        type + "-transactions", ".csv")) ← Generate dynamic
                                           filename);
```

---


# Text IO writing with dynamic destinations

Python

```
(my_pcollection
| beam.io.fileio.WriteToFiles( ← Dynamic destination
    path='/my/file/path',
    destination=lambda record: 'avro'
    if record['type'] == 'A' else 'csv',
    sink=lambda dest: AvroSink() ← Write to dynamic sink
    if dest == 'avro' else CsvSink(),
    file_naming=beam.io.fileio
    .destination_prefix_naming()))
```

# Sources and Sinks

## Agenda



---

# BigQuery IO reading with query

Java

```
PCollection<Double> maxTemperatures =  
    p.apply(  
        BigQueryIO.read(  
            (SchemaAndRecord elem) -> (Double) ← Map results  
                elem.getRecord()  
                .get("max_temperature"))  
        .fromQuery( ← Source using query  
            "SELECT max_temperature FROM  
            `clouddataflow-readonly.samples.weather_stations`")  
        .usingStandardSql()  
        .withCoder(DoubleCoder.of()));
```

---

# BigQuery IO reading with query

Python

```
max_temperatures = (
    p
    | 'QueryTableStdSQL' >> beam.io.ReadFromBigQuery( ← Map results
        query='SELECT max_temperature FROM '\
        '`clouddataflow-readonly.samples.weather_stations`',
        use_standard_sql=True)
    | beam.Map(lambda elem: elem['max_temperature'])) ← Source using query
```



# BigQuery IO reading with BigQuery Storage API

Java

```
PCollection<MyData> rows =  
    pipeline.apply("Read from BigQuery table",  
        BigQueryIO.readTableRows()  
            .from(  
                String.format("%s:%s.%s",  
                    project, dataset, table))  
            .withMethod(Method.DIRECT_READ)  
            // .withRowRestriction  
            .withSelectedFields(  
                Arrays.asList(..."string_...", "Int64..."))  
            .apply("TableRows to MyData",  
                MapElements.into(  
                    TypeDescriptor.of(MyData.class))  
                .via(MyData::fromTableRow))
```

Storage API  
read method

Utilizes predicate  
filtering

Utilizes column  
projection

---

# BigQuery IO writing with dynamic destinations

Java

```
pc.apply(BigQueryIO.<Purchase>write(tableSpec)
    .useBeamSchema() ← Schema definition
    .to((ValueInSingleWindow<Purchase> purchase) -> {
        return new TableDestination(
            "project:dataset-" +
                purchase.getValue().getUser() + ← Dynamic destination
                ":purchases", "");
    });
```

---

# BigQuery IO writing with dynamic destinations

Python

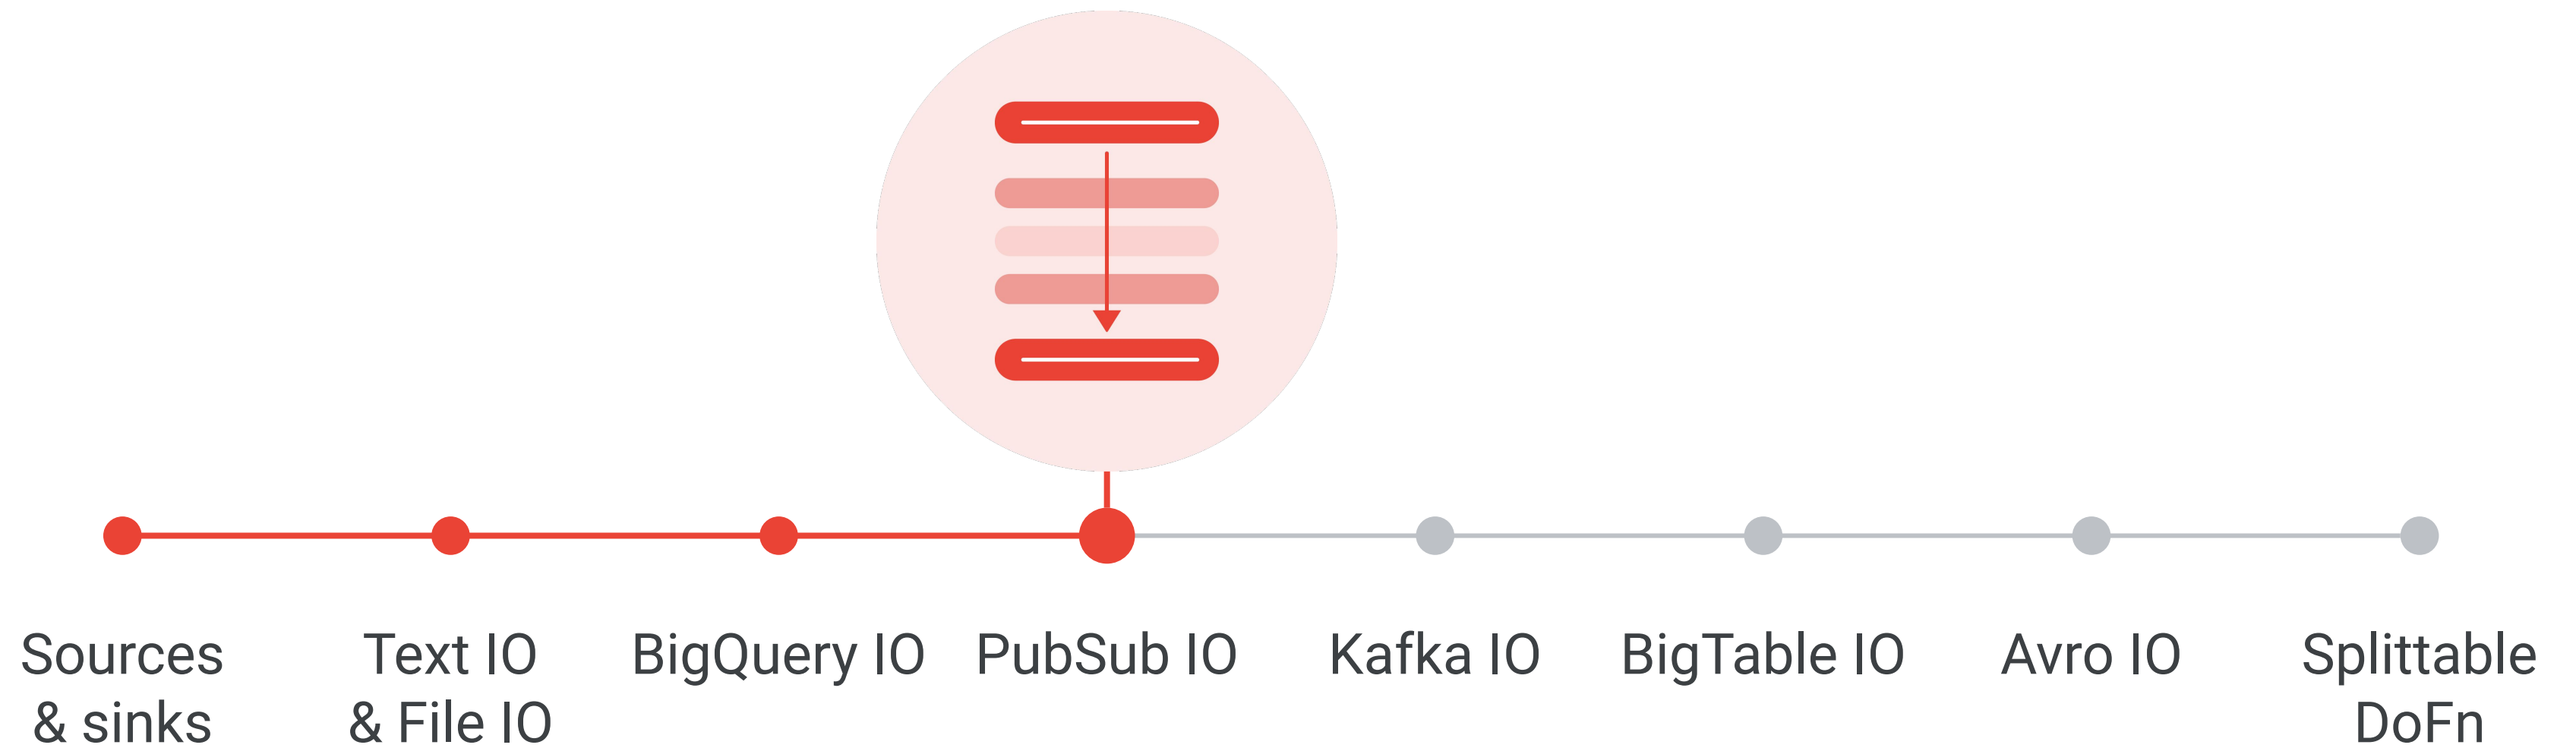
```
def table_fn(element, fictional_characters):  
    if element in fictional_characters:  
        return 'my_dataset.fictional_quotes'  
    else:  
        return 'my_dataset.real_quotes'  
  
quotes | 'WriteWithDynamicDestination' >>  
beam.io.WriteToBigQuery(  
    table_fn,  
    schema=table_schema,  
    table_side_inputs=(fictional_characters_view, ),  
    ...)
```

Dynamic destination

Schema destination

# Sources and Sinks

## Agenda



---

# PubSub IO reading

Java

```
pipeline
  .apply("Read PubSub Messages",
    PubsubIO
      .readStrings()
      .fromTopic(options.getInputTopic()))
  .apply(
    Window.into( ←
      FixedWindows.of(
        Duration.standardMinutes(
          options.getWindowSize()))));
```

Windowing using  
message timestamps

---

# PubSub IO reading

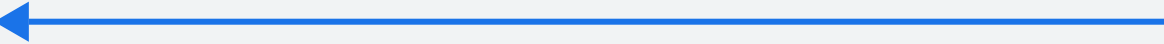
Java

```
class GroupWindowsIntoBatches(beam.PTransform):
```

```
...
```

```
    >> beam.WindowInto(  
        window.FixedWindows(self.window_size))
```

Windowing using  
message timestamps



```
pipeline
```

```
    | "Read PubSub Messages"
```

```
    >> beam.io.ReadFromPubSub(topic=input_topic)
```

```
    | "Window into"
```

```
    >> GroupWindowsIntoBatches(window_size)
```

# PubSub IO reading

## Java

```
appliedUdf
    .get(KafkaPubsubConstants.UDF_OUT)
    .apply("getSuccessUDFOutElements",
        MapElements.into(stringTypeDescriptor).via(FailsafeElement::getPayload))
    .setCoder(NullableCoder.of(StringUtf8Coder.of()))
    .apply("writeSuccessMessages",
        PubsubIO.writeStrings().to(options.getOutputTopic()));

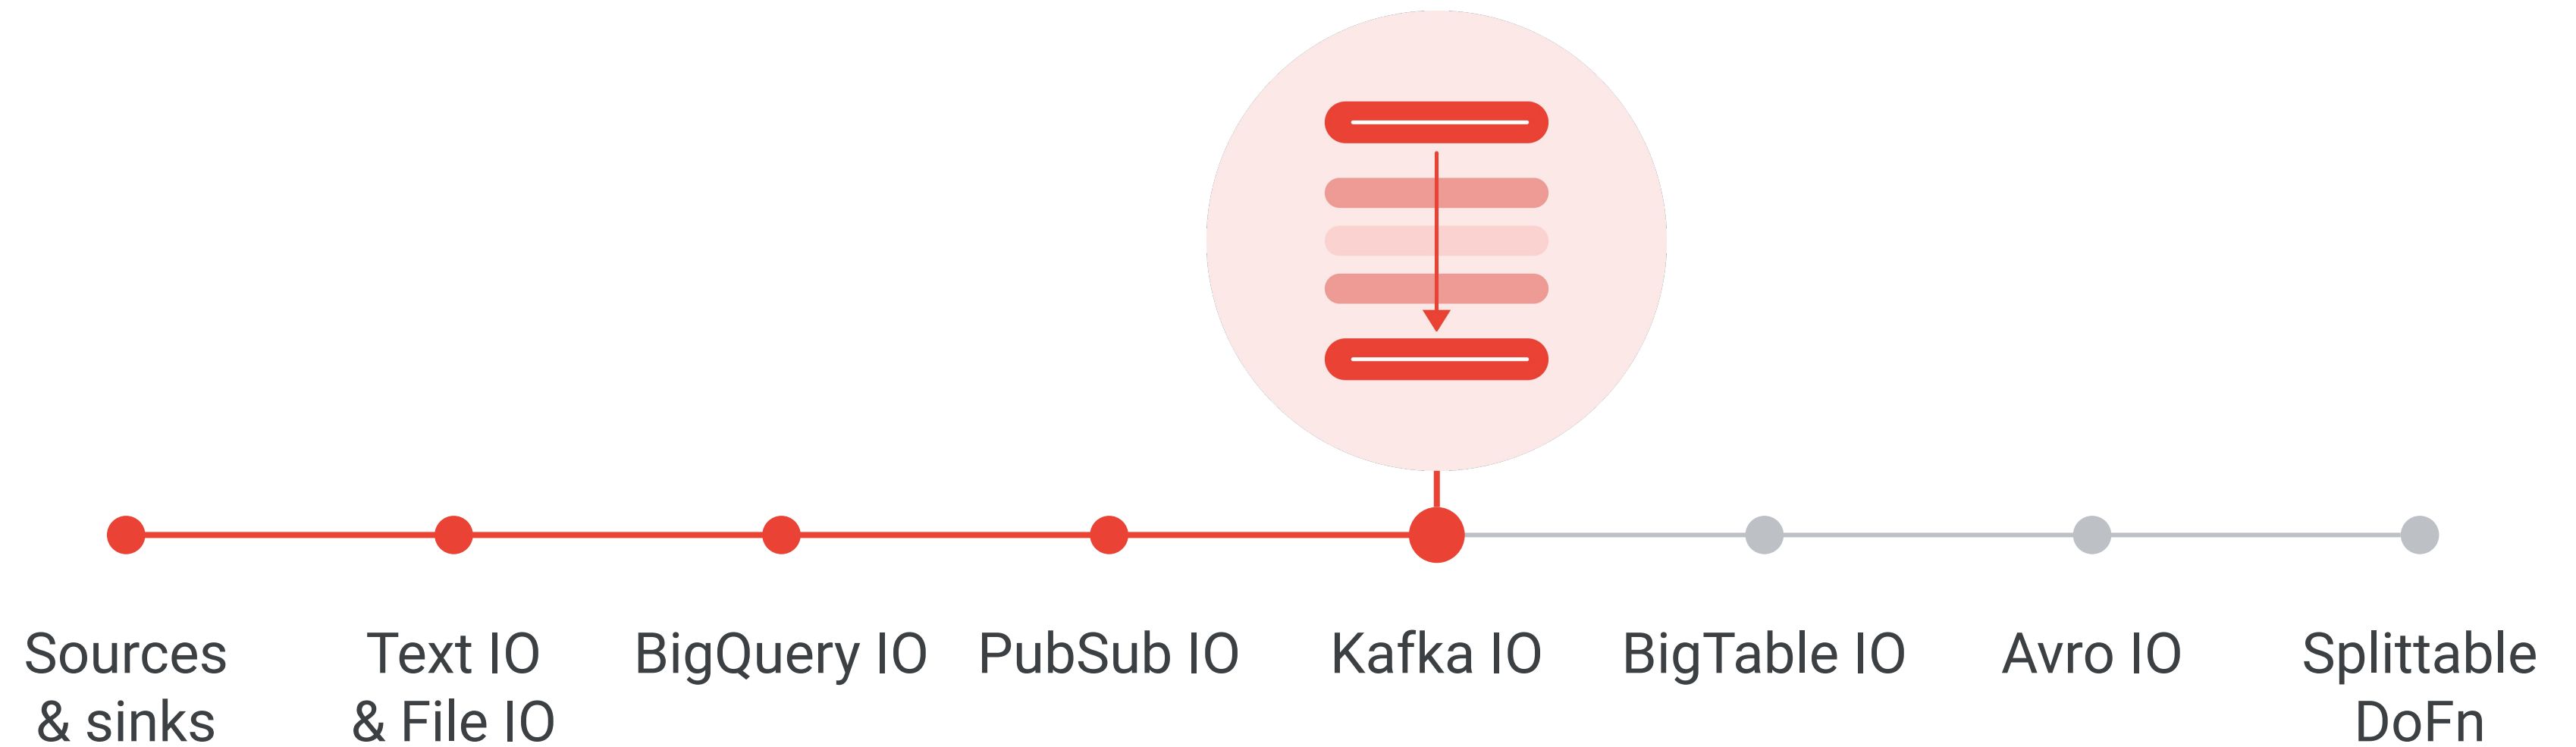
if (options.getOutputDeadLetterTopic() != null) {
    appliedUdf.get(KafkaPubsubConstants.UDF_DEADLETTER_OUT)
        .apply("getFailedMessages",
            MapElements.into(
                TypeDescriptors.kvs(stringTypeDescriptor, stringTypeDescriptor))
            .via(FailsafeElement::getOriginalPayload))
        .apply("extractMessageValues",
            MapElements.into(stringTypeDescriptor).via(KV<String, String>::getValue))
        .setCoder(NullableCoder.of(StringUtf8Coder.of()))
        .apply("writeFailureMessages",
            PubsubIO.writeStrings().to(options.getOutputDeadLetterTopic()));
}
```

Write message that  
passes UDF criteria

Dead-letter message  
that fails UDF criteria

# Sources and Sinks

## Agenda





---

# Kafka IO reading

Java

```
PCollection<KV<String, String>> records =  
pipeline  
    .apply("Read From Kafka",  
        KafkaIO.<String, String>read()  
            .withConsumerConfigUpdates(ImmutableMap.of(  
                ConsumerConfig  
                    .AUTO_OFFSET_RESET_CONFIG, "earliest"))  
    .withBootstrapServers(options.getBootstrapServers())  
        .withTopics(<...list...>) ← Topic selection  
        .withKeyDeserializerAndCoder(...))  
        .withValueDeserializerAndCoder(...)  
        .withoutMetadata())
```

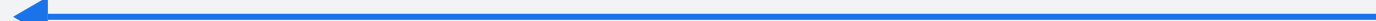
---

# Kafka IO reading

Python


```
pipeline
| ReadFromKafka(
  consumer_config={
    'bootstrap.servers': bootstrap_servers,
    topics=[topic])
```

Cross-language  
transforms



# Sources and Sinks

## Agenda

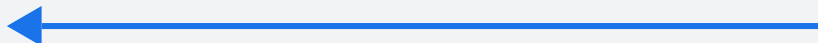


---

# BigTable IO reading with row filters

Java

```
p.apply("filtered read",  
    BigtableIO.read()  
        .withProjectId(projectId)  
        .withInstanceId(instanceId)  
        .withTableId("table")  
        .withRowFilter(filter));
```



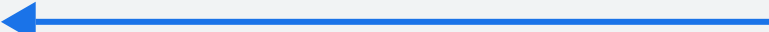
Row filtered read

---

# BigTable IO reading with prefix scan

Java

```
ByteKeyRange keyRange = ...;  
p.apply("read",  
    BigtableIO.read()  
        .withProjectId(projectId)  
        .withInstanceId(instanceId)  
        .withTableId("table")  
        .withKeyRange(keyRange));
```



Prefix scan

---

# BigTable IO writing with additional actions

Java

```
PCollection<KV<..., Iterable<Mutation>>> data = ...;
```


```
PCollection<BigtableWriteResult> writeResults =  
    data.apply("write", BigtableIO.write()  
        .withProjectId("project")  
        .withInstanceId("instance")  
        .withTableId("table"))  
        .withWriteResults();
```

```
PCollection<...> moreData = ...;
```

moreData


```
    .apply("wait for writes", Wait.on(writeResults))  
    .apply("do something", ParDo.of(...))
```

Additional actions  
triggered after write  
completion



# Sources and Sinks

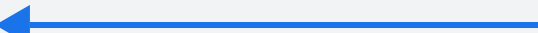
## Agenda

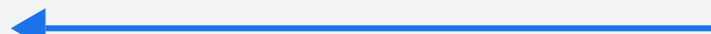


---

# Avro IO reading with known schema

Java

```
PCollection<AvroAutoGenClass> records =  Read Avro schema  
p.apply(AvroIO.read(AvroAutoGenClass.class)  
    .from("gs:...*.avro"));
```

```
Schema schema = new Schema.Parser()  Read Avro schema  
    .parse(new File("schema.avsc"));  
PCollection<GenericRecord> records =  
p.apply(AvroIO.readGenericRecords(schema)  
    .from("gs:...-*.avro"));
```



---


# Avro IO reading multiple files

Python

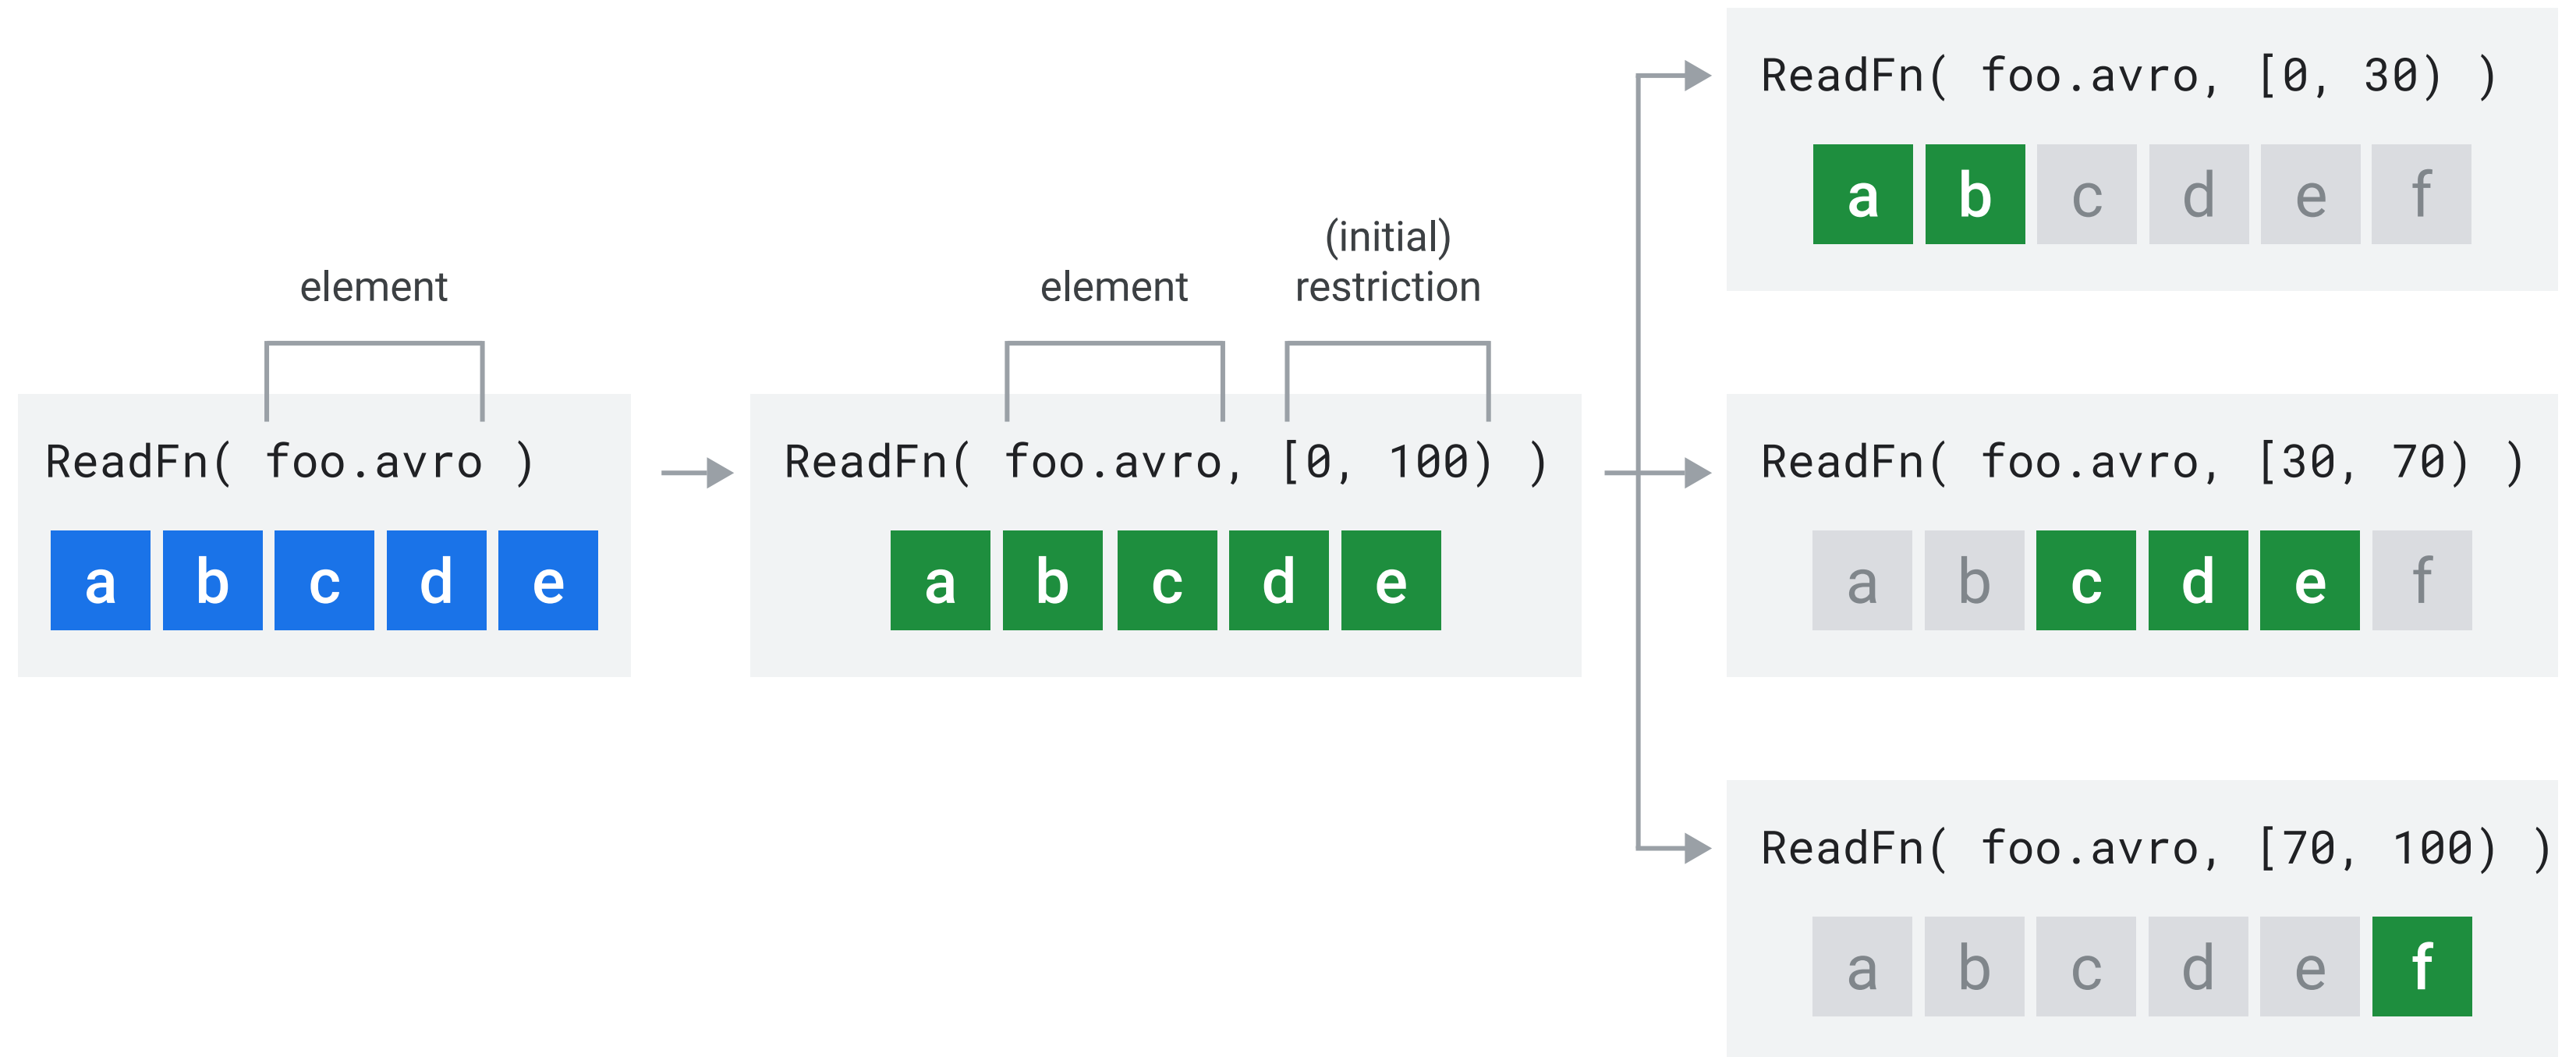
```
with beam.Pipeline() as p:  
    records = p | 'Read' >> beam.io.ReadFromAvro(  
        '/mypath/myavrofiles*')  
    # Read Avro schema
```

# Sources and Sinks

## Agenda



# Splittable DoFN



---

# Splittable DoFn custom source

Java

```
@BoundedPerElement
private static class FileToWordsFn extends DoFn<String,
Integer> {
    @GetInitialRestriction
    public OffsetRange getInitialRestriction(
        @Element String fileName) throws IOException {
        return new OffsetRange(0,
            new File(fileName).length());
    }

    @ProcessElement
    public void processElement(
        @Element String fileName,
        RestrictionTracker<OffsetRange, Long> tracker,
        OutputReceiver<Integer> outputReceiver){...}
```

← Tracking subset of  
restriction completed

---

# Splittable DoFn custom source

Python

```
class FileToWordsRestrictionProvider(
    beam.io.RestrictionProvider):
    def initial_restriction(self, file_name):
        return OffsetRange(0,
os.stat(file_name).st_size)

    def create_tracker(self, restriction):
        return beam.io.restriction_trackers
            .OffsetRestrictionTracker()

class FileToWordsFn(beam.DoFn):
    def process(...)
```

Initial restriction

Tracking subset of restriction completed

---

# Dataflow best practices

## Leverage templates

Use the existing templates or use the open source code as a basis for your code to accelerate your pipeline development.

[!\[\]\(eafc244b53721dd1ec133f0772f70fc7\_img.jpg\) Python Beam examples](#)

[!\[\]\(d3fb9f94af8b26d1c844efa9a98805b0\_img.jpg\) Java Dataflow templates](#)

