



State and Timers

Israel Herraiz

Strategic Cloud Engineer,
Google Cloud





Agenda

Course Intro

Beam Concepts Review

Windows, Watermarks, and Triggers

Sources and Sinks

Schemas

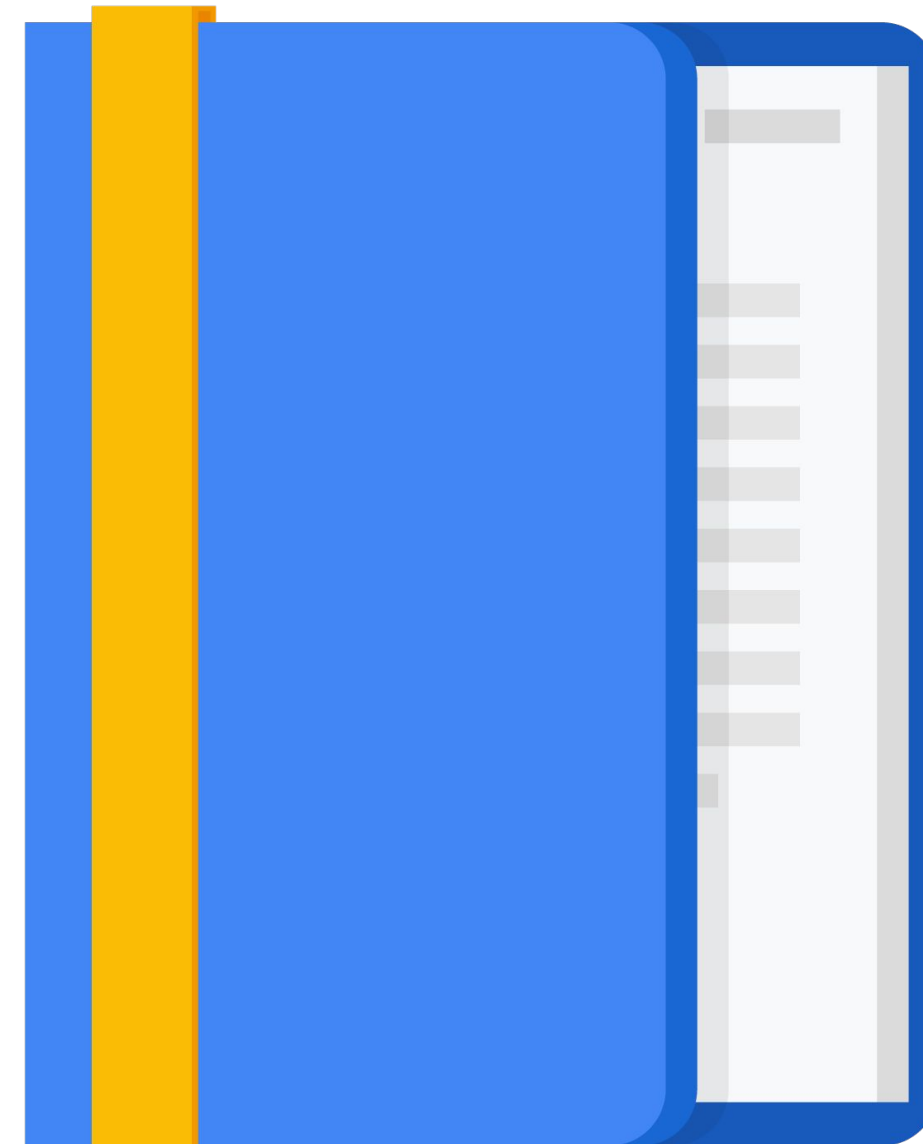
State and Timers

Best Practices

Dataflow SQL and DataFrames

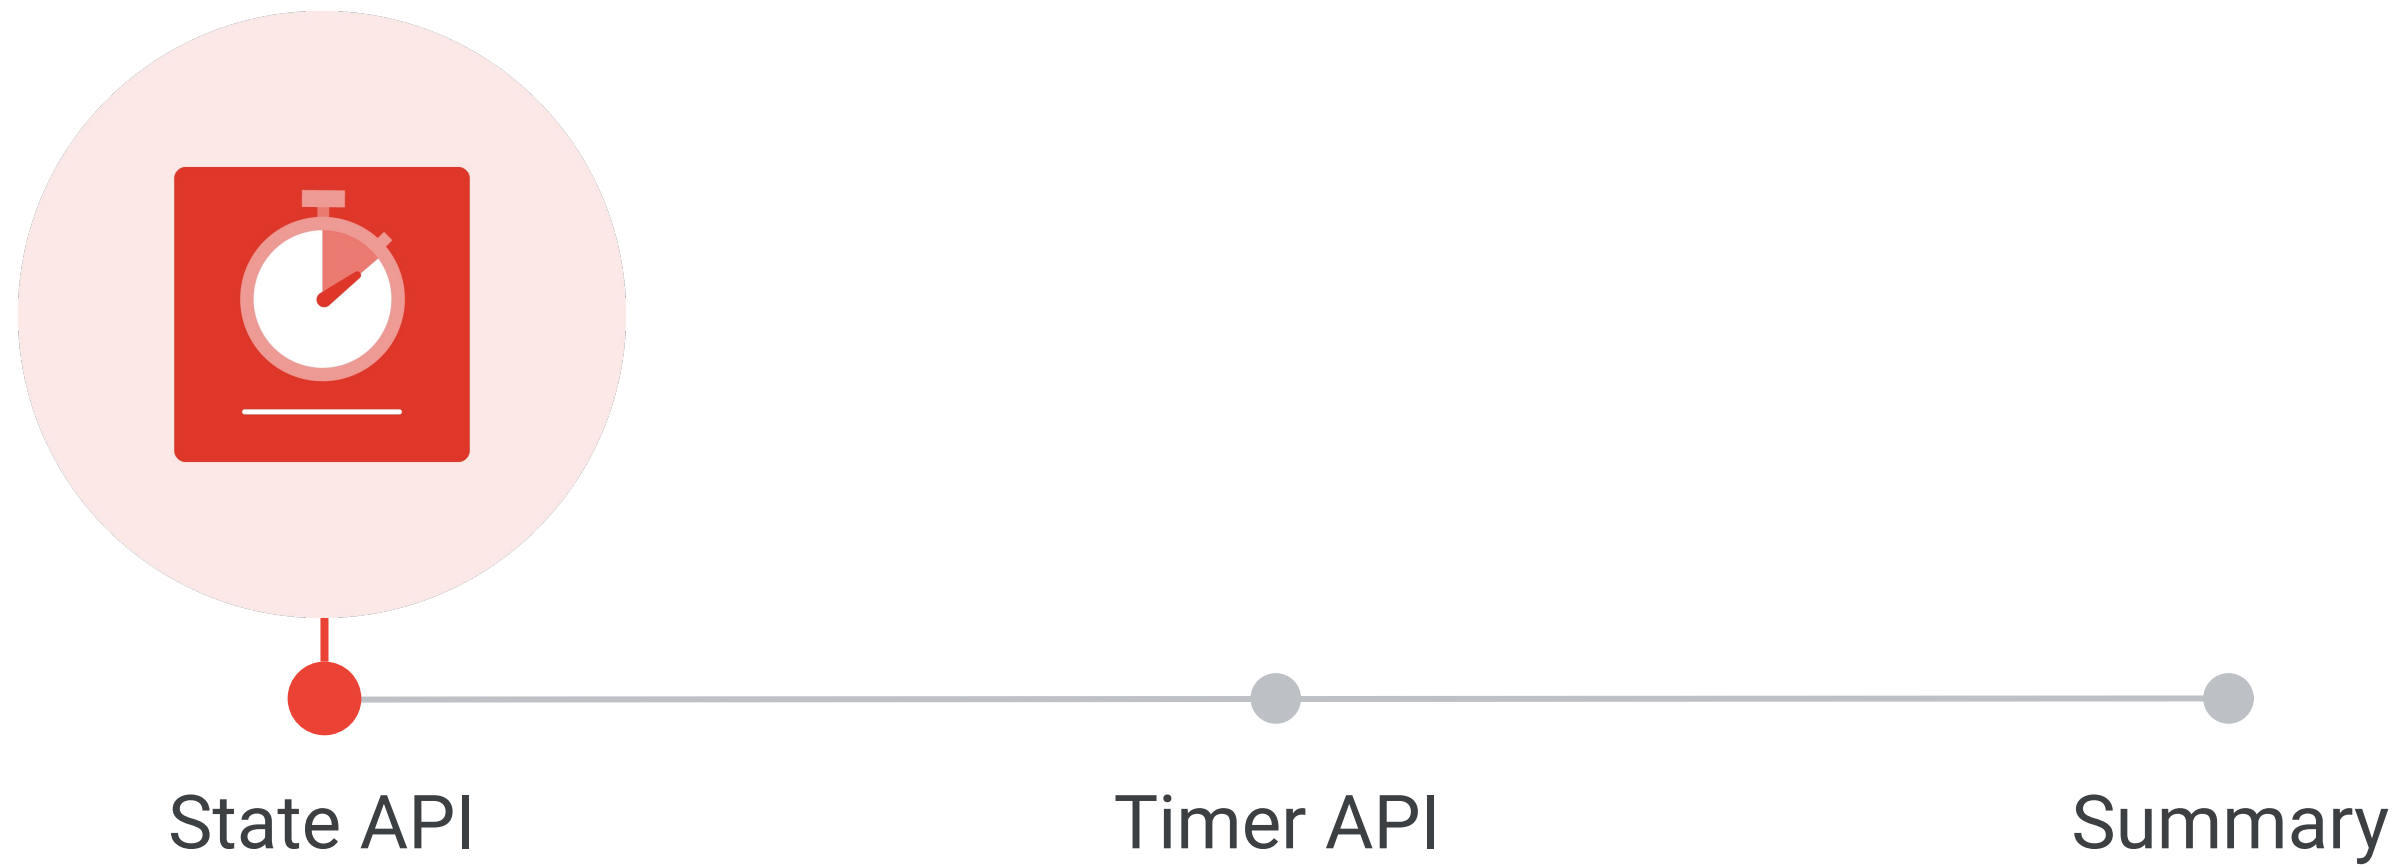
Beam Notebooks

Summary

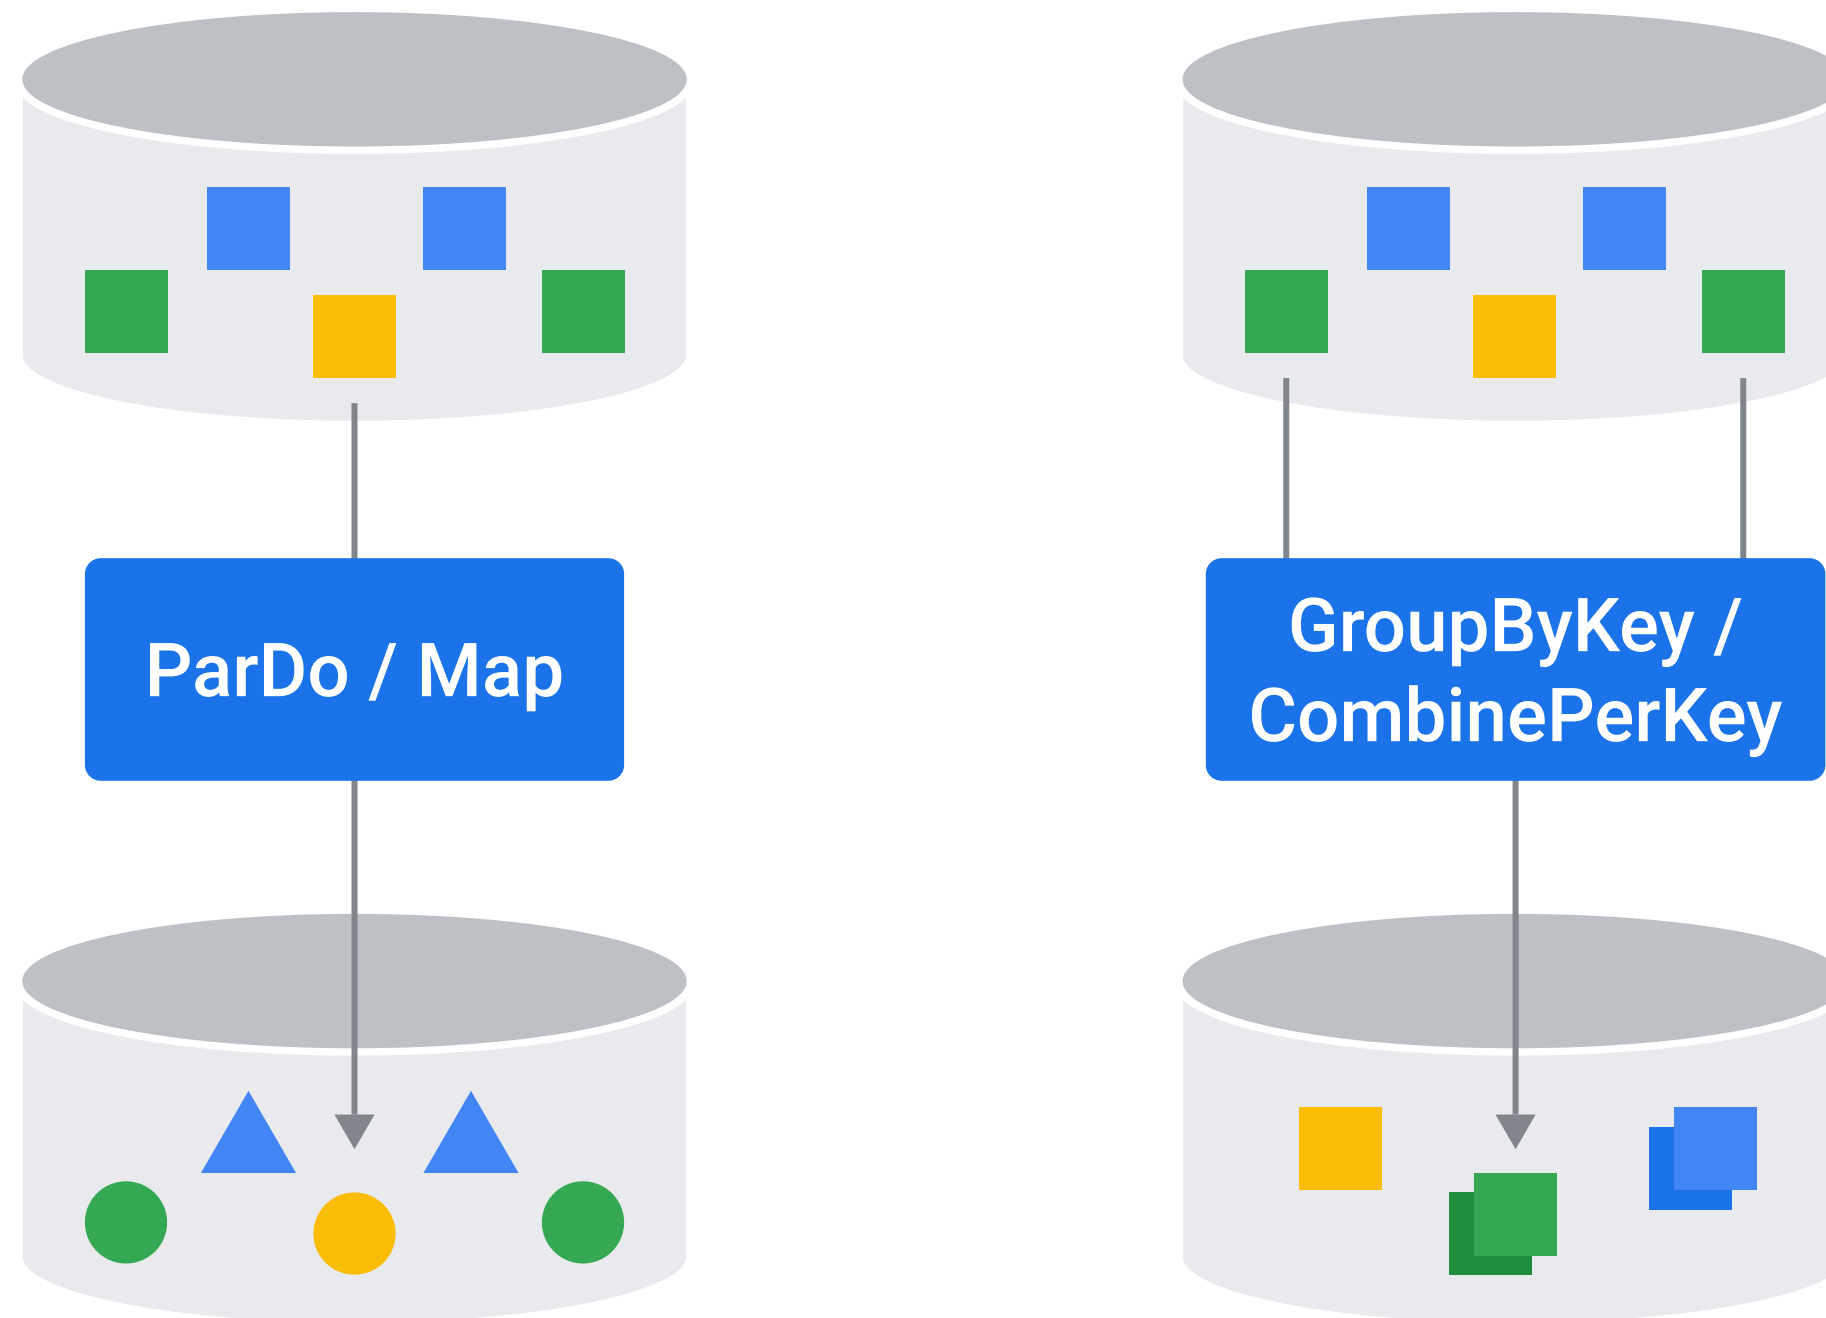


State and Timers

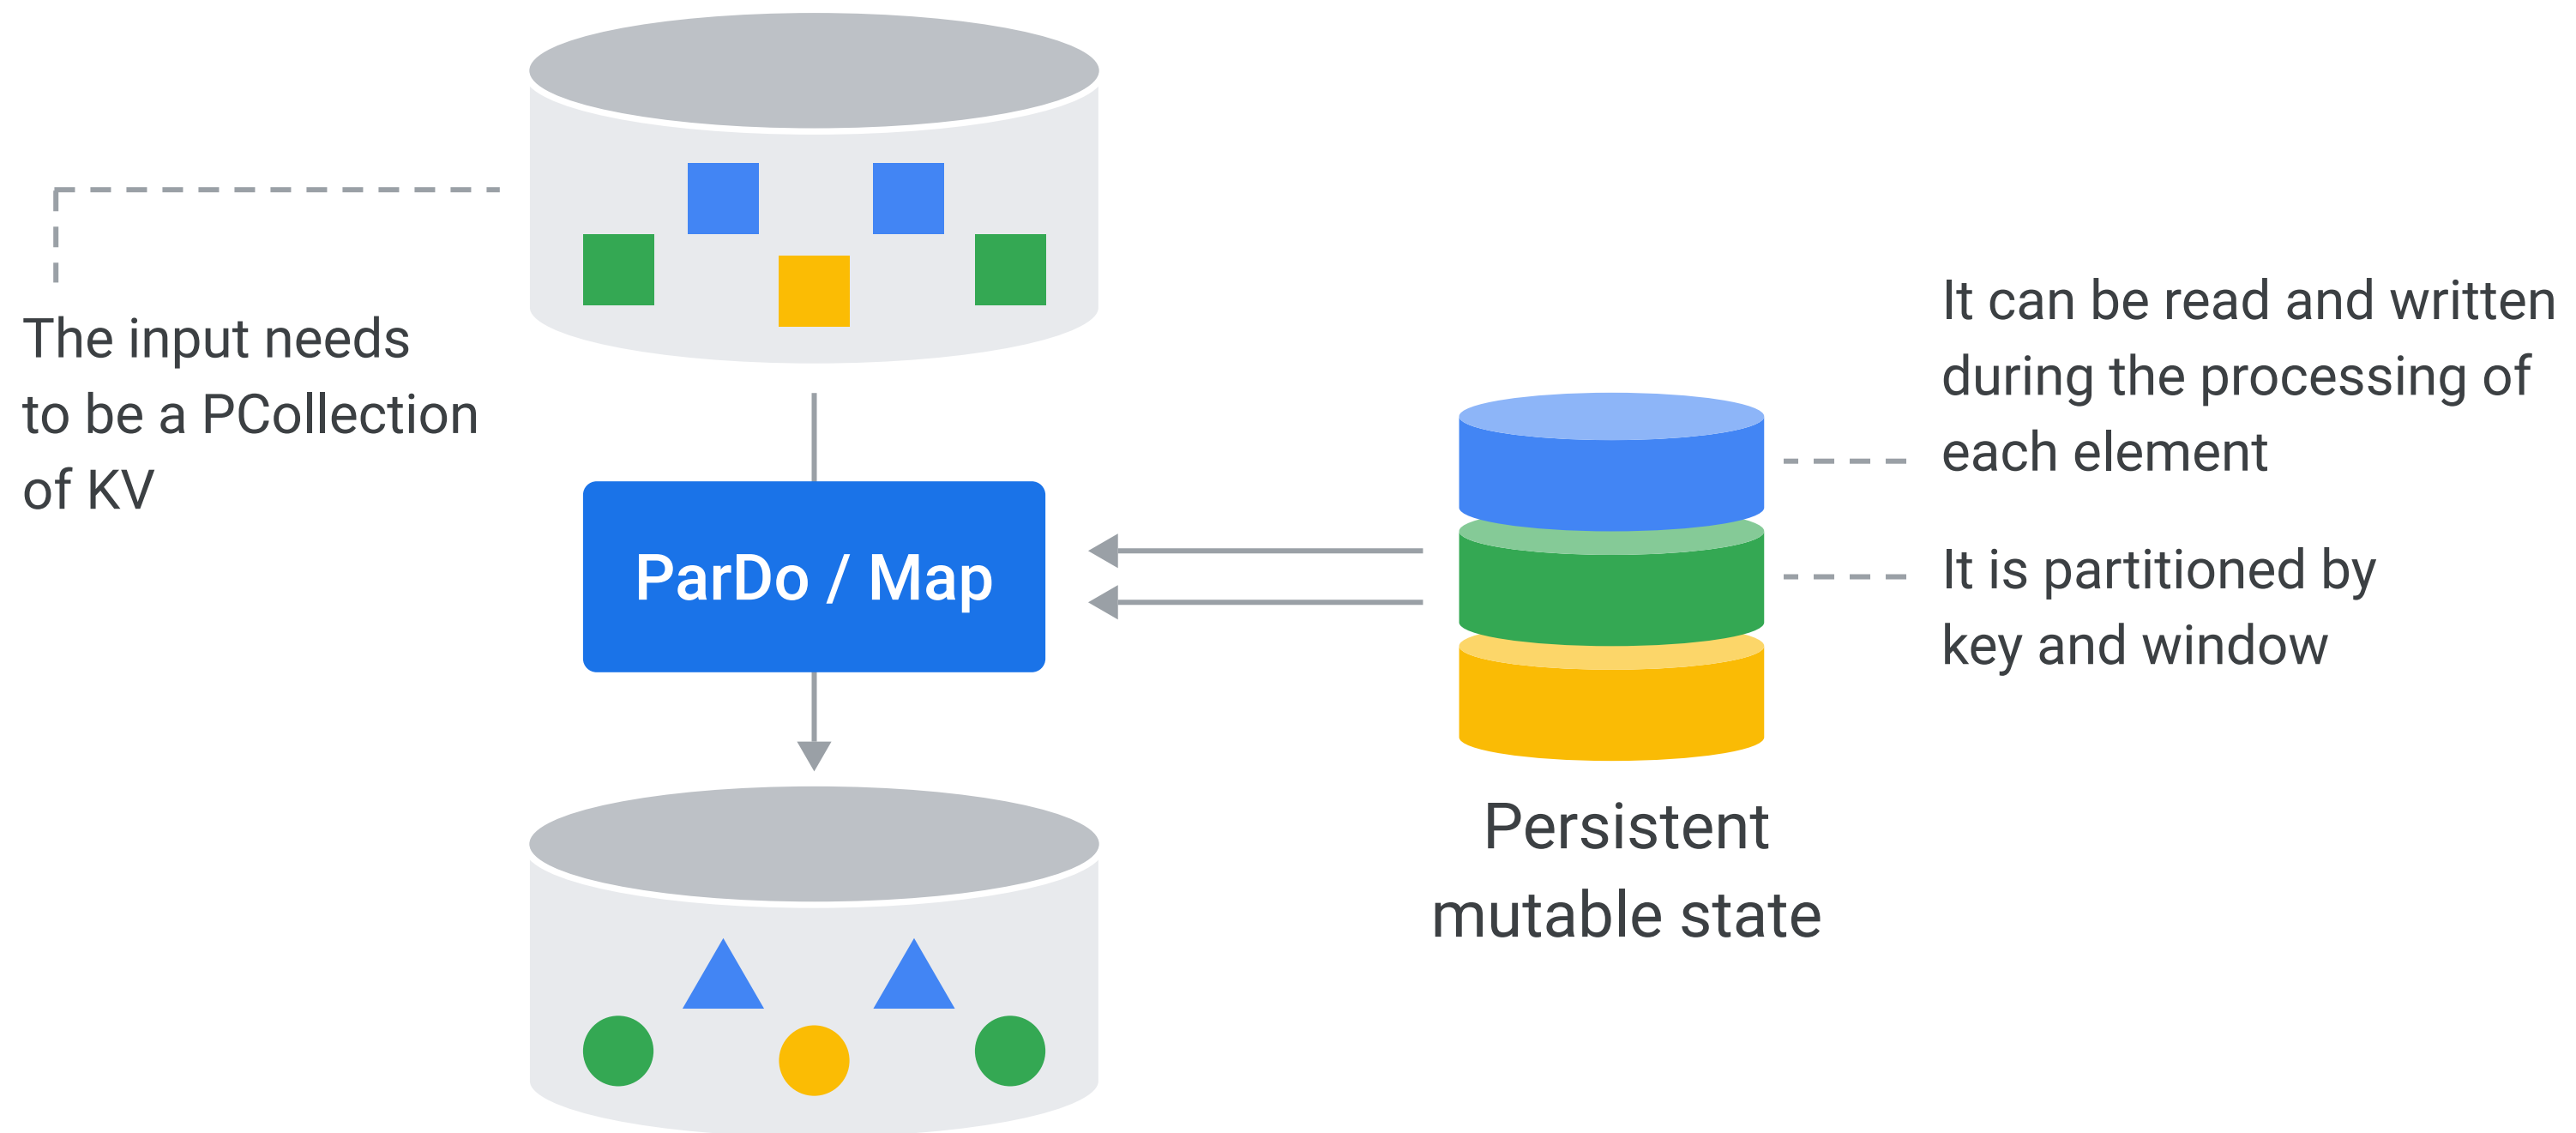
Agenda



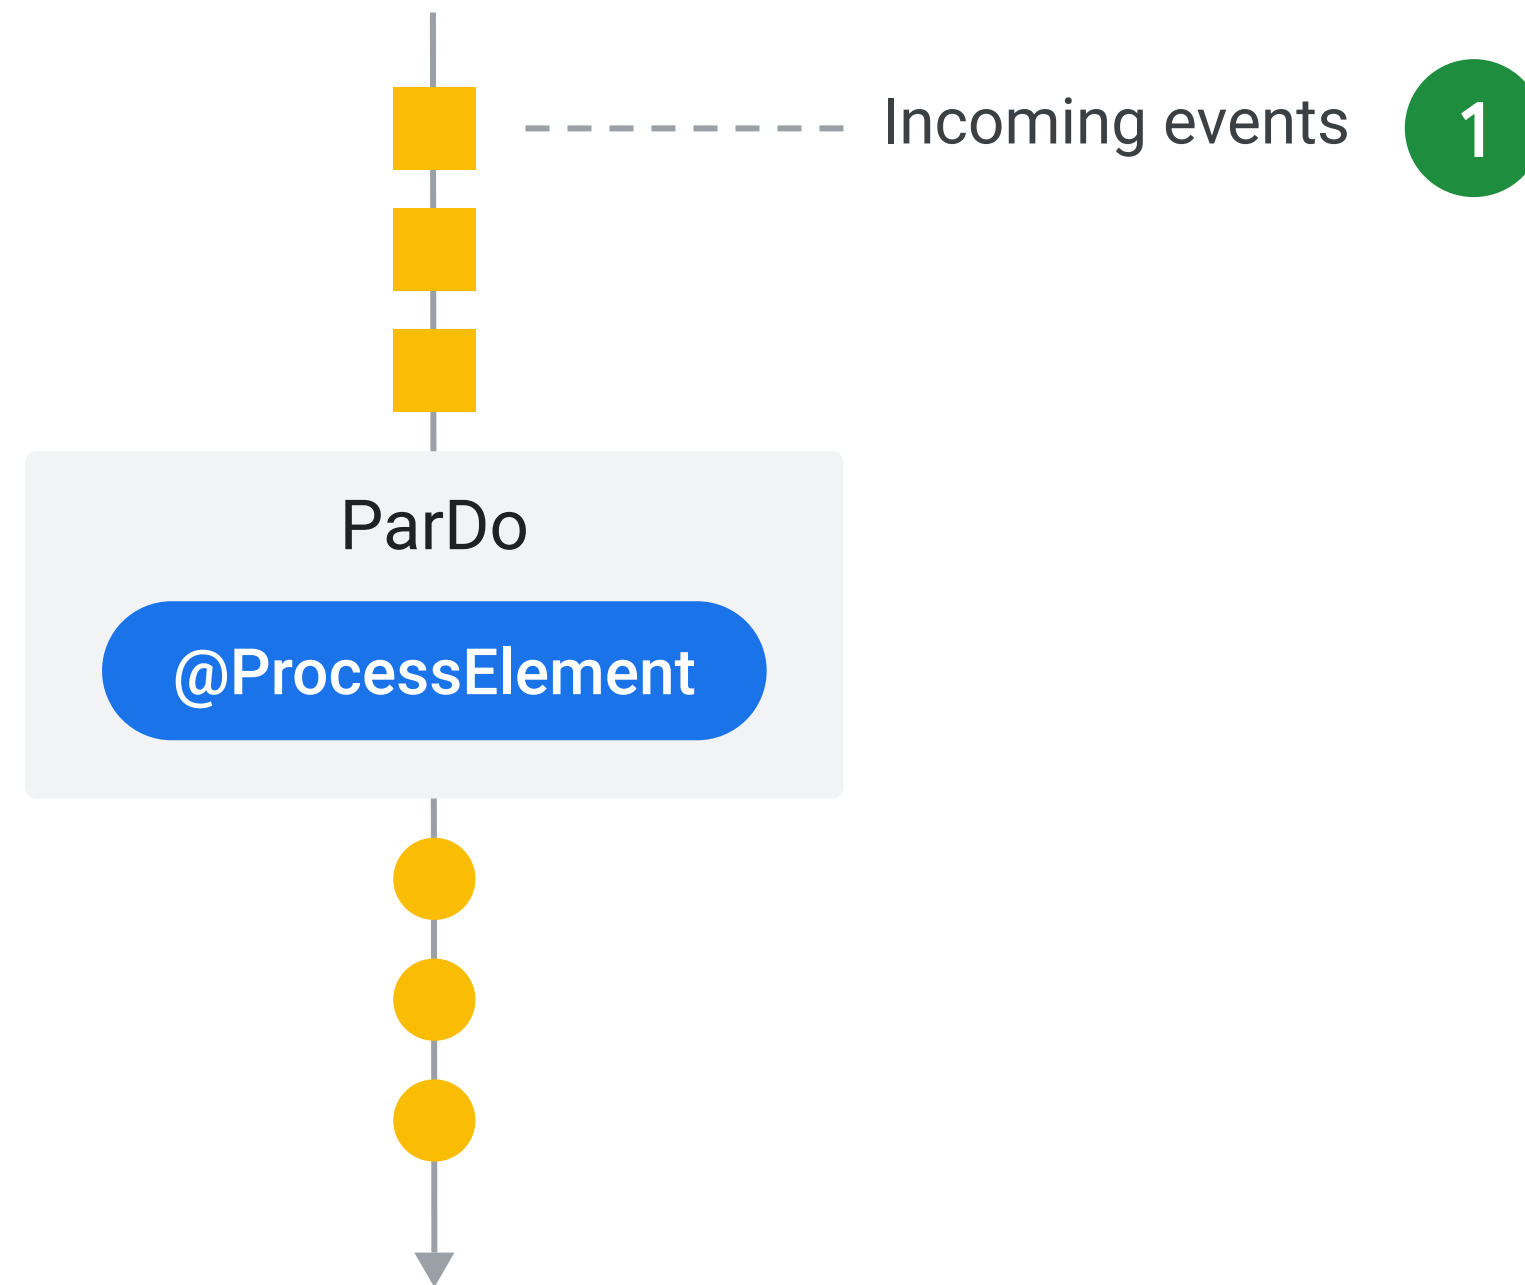
How to aggregate data in Apache Beam



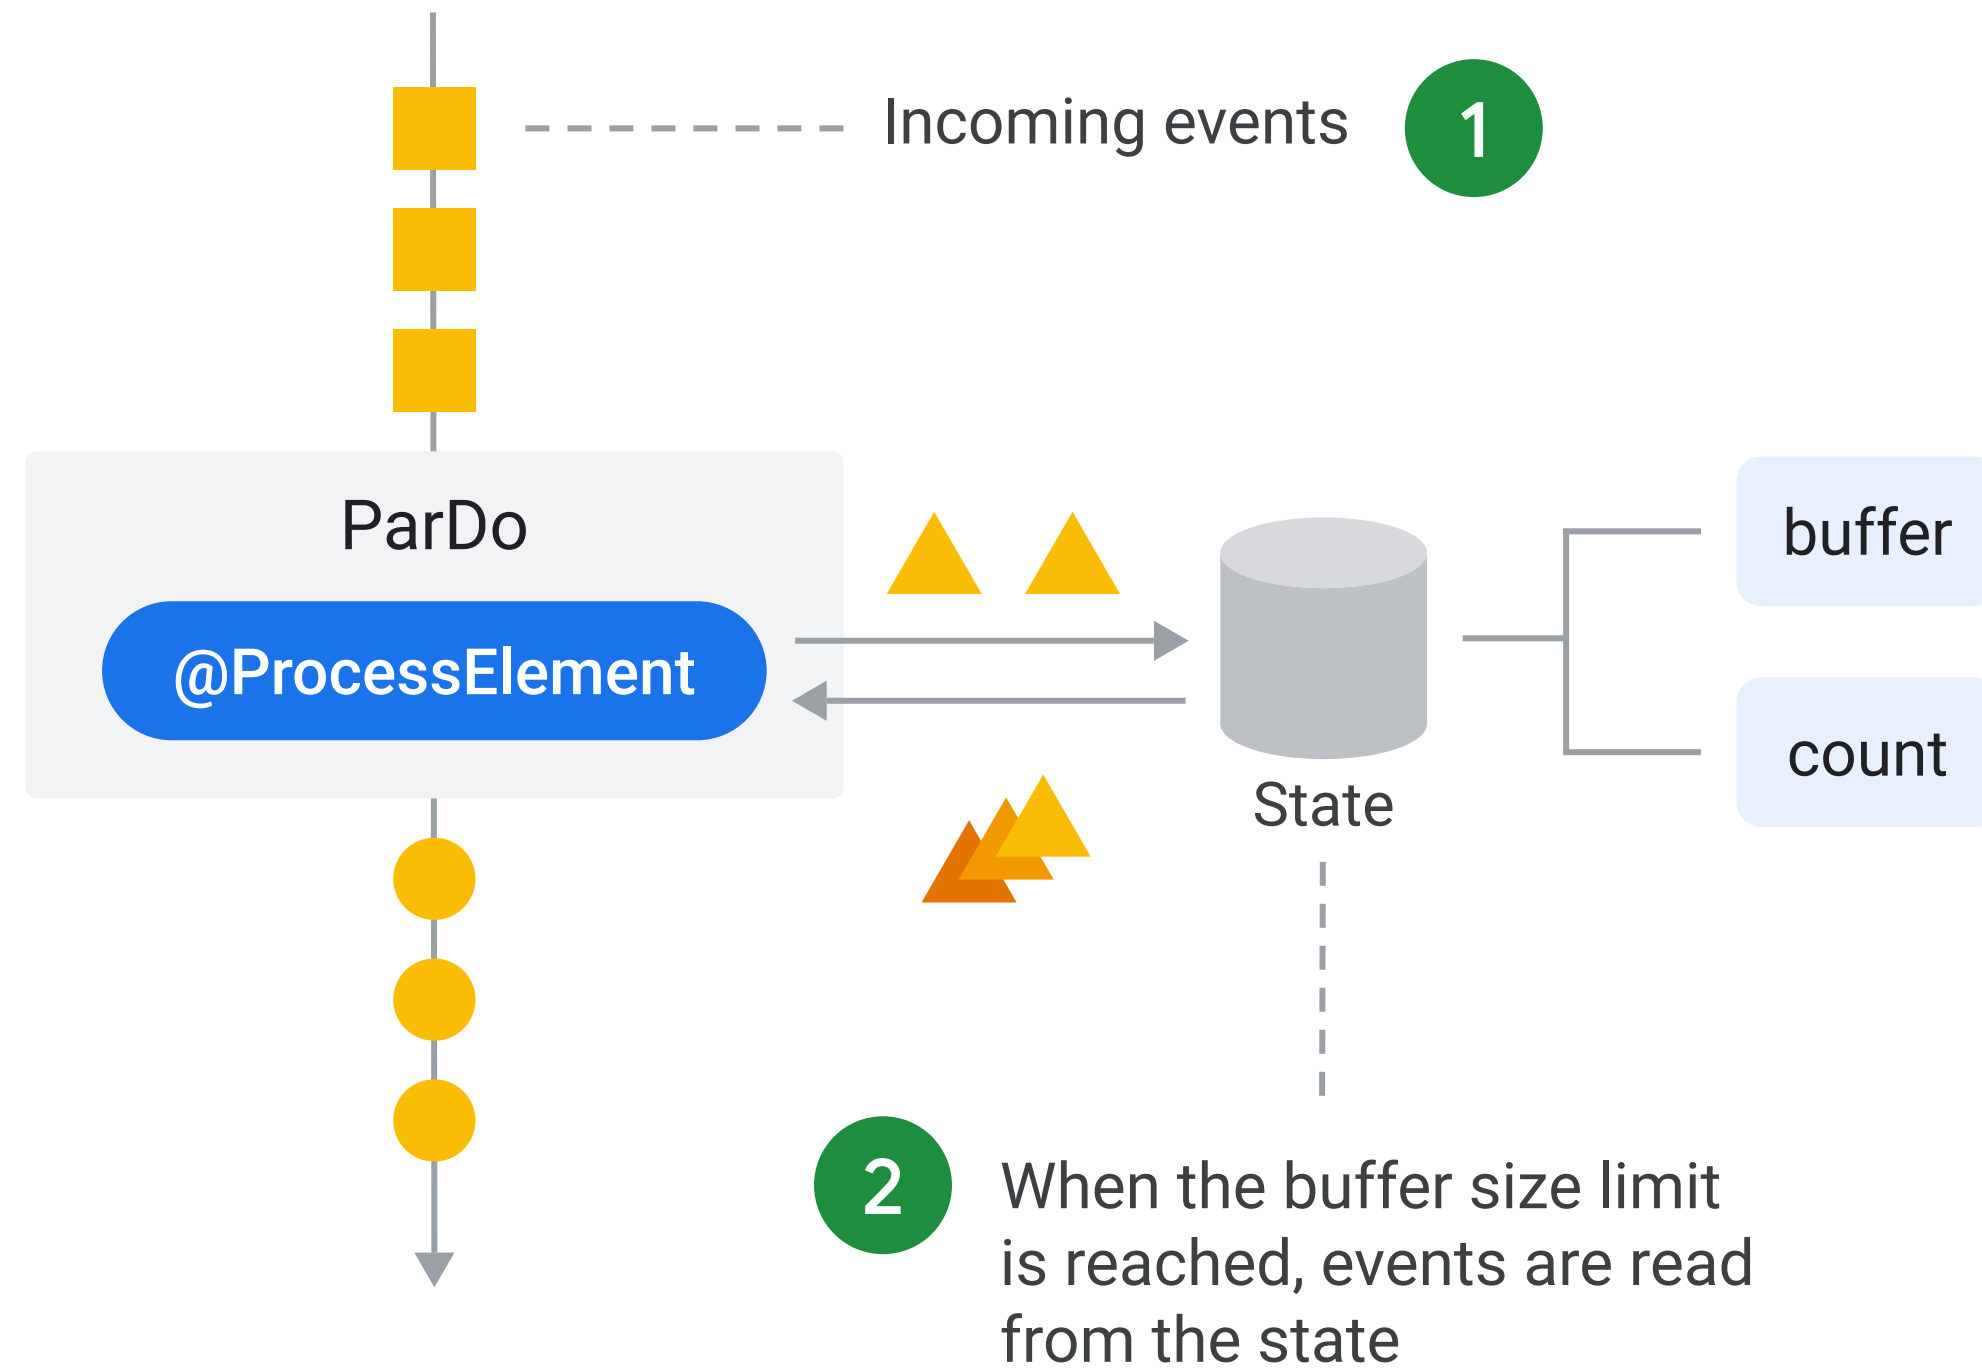
Introducing stateful ParDo



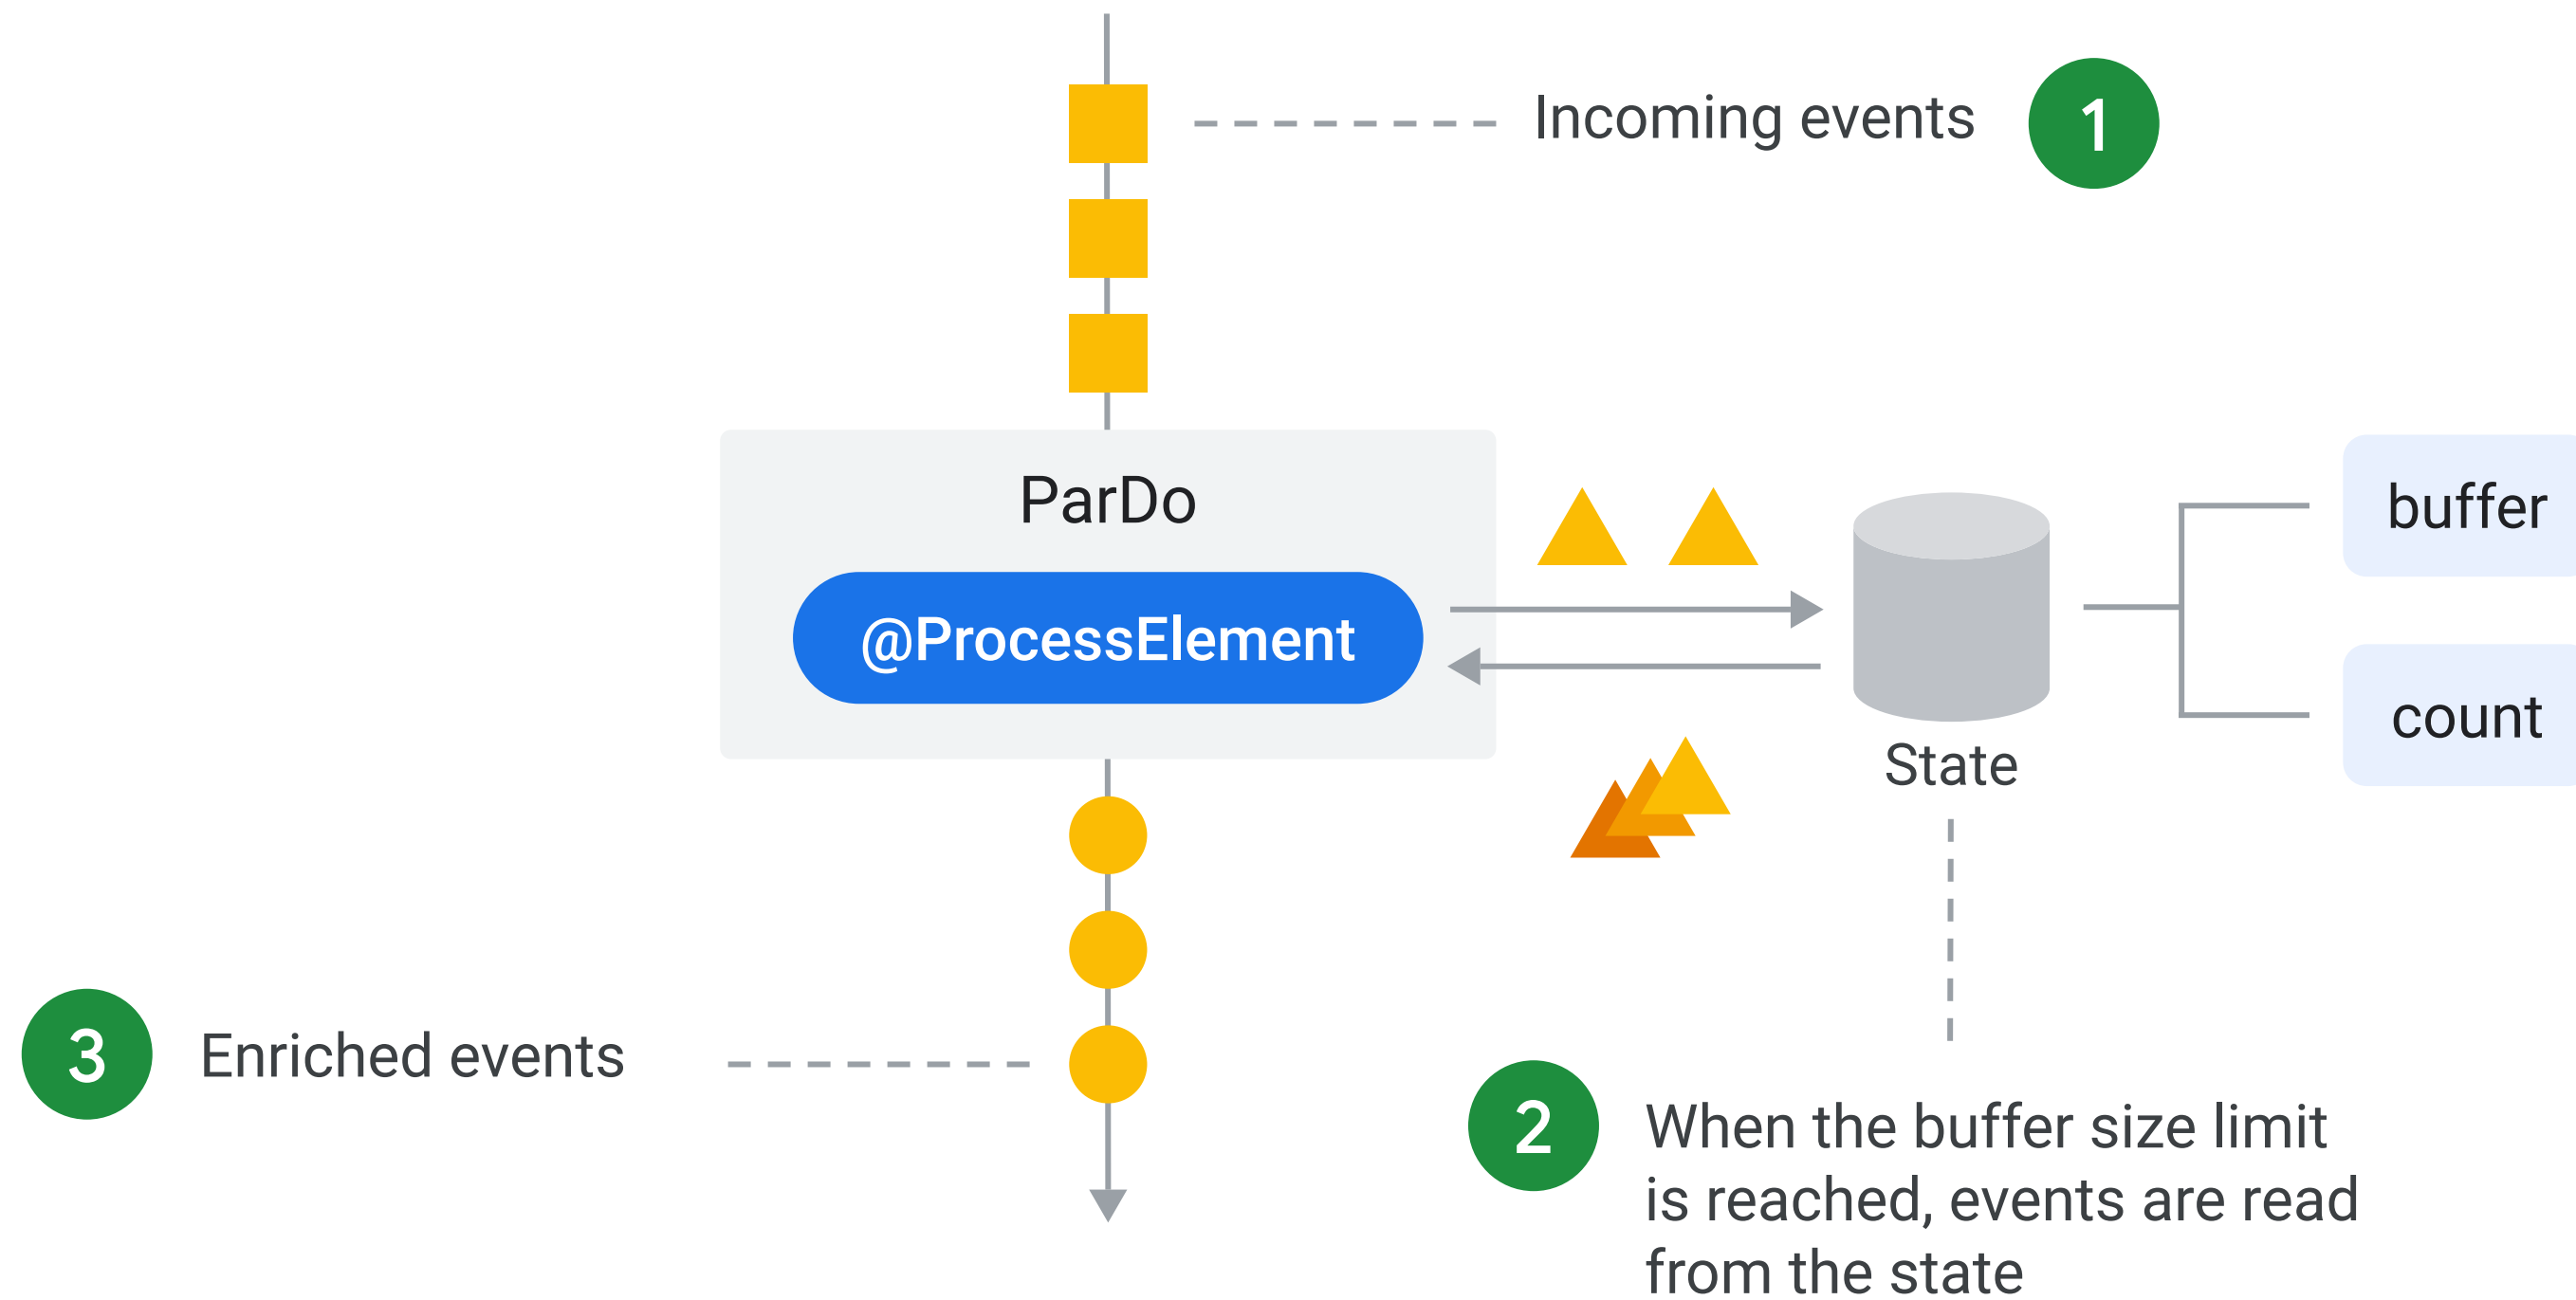
How does it work?



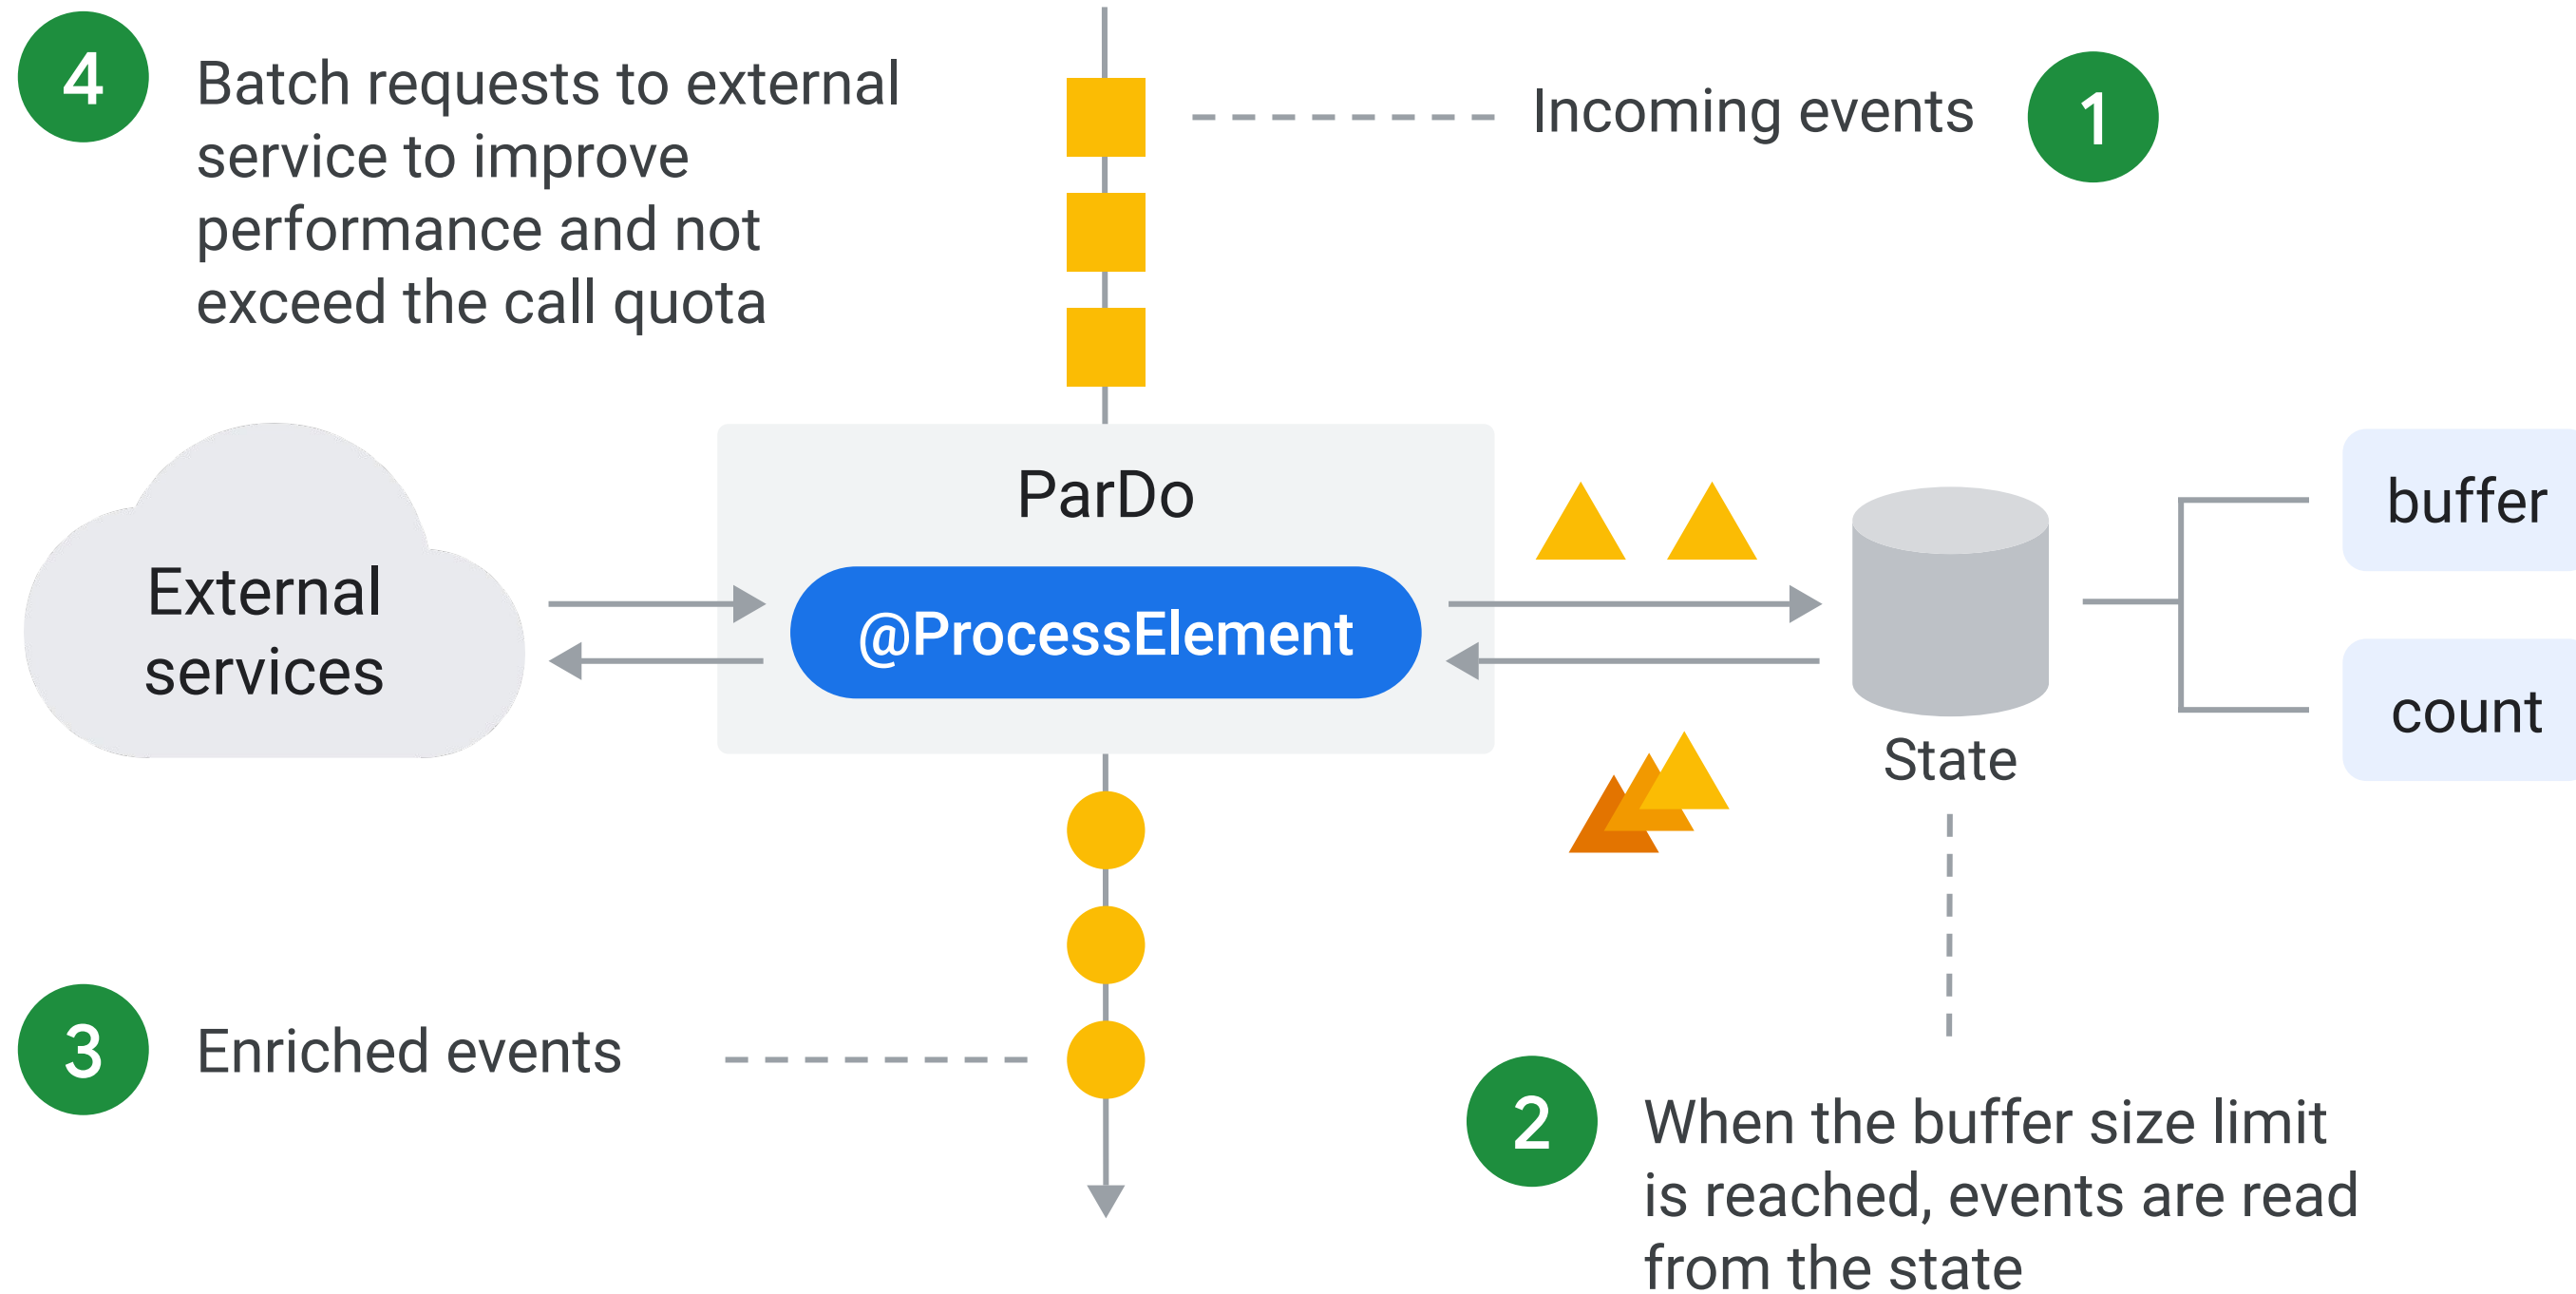
How does it work?



How does it work?



What is it useful for?



Python

```
class StatefulBufferingFn(beam.DoFn):
    MAX_BUFFER_SIZE = 500;
    BUFFER_STATE = BagStateSpec('buffer', EventCoder())
    COUNT_STATE = CombiningValueStateSpec('count',
                                           VarIntCoder(),
                                           combiners.SumCombineFn())

    def process(self, element,
               buffer_state=beam.DoFn.StateParam(BUFFER_STATE),
               count_state=beam.DoFn.StateParam(COUNT_STATE)):
        buffer_state.add(element)
        count_state.add(1)
        count = count_state.read()
        if count >= MAX_BUFFER_SIZE:
            for event in buffer_state.read():
                yield event
            count_state.clear()
            buffer_state.clear()
```

Increment count and
add element to buffer

When buffer size limit
is reached, a request is
sent to the external
service

Java

```
new DoFn<Event, EnrichedEvent>() {
    private static final int MAX_BUFFER_SIZE = 500;
    @StateId("buffer") private final StateSpec<BagState<Event>> bufferedEvents = StateSpecs.bag();
    @StateId("count") private final StateSpec<ValueState<Integer>> countState = StateSpecs.value();

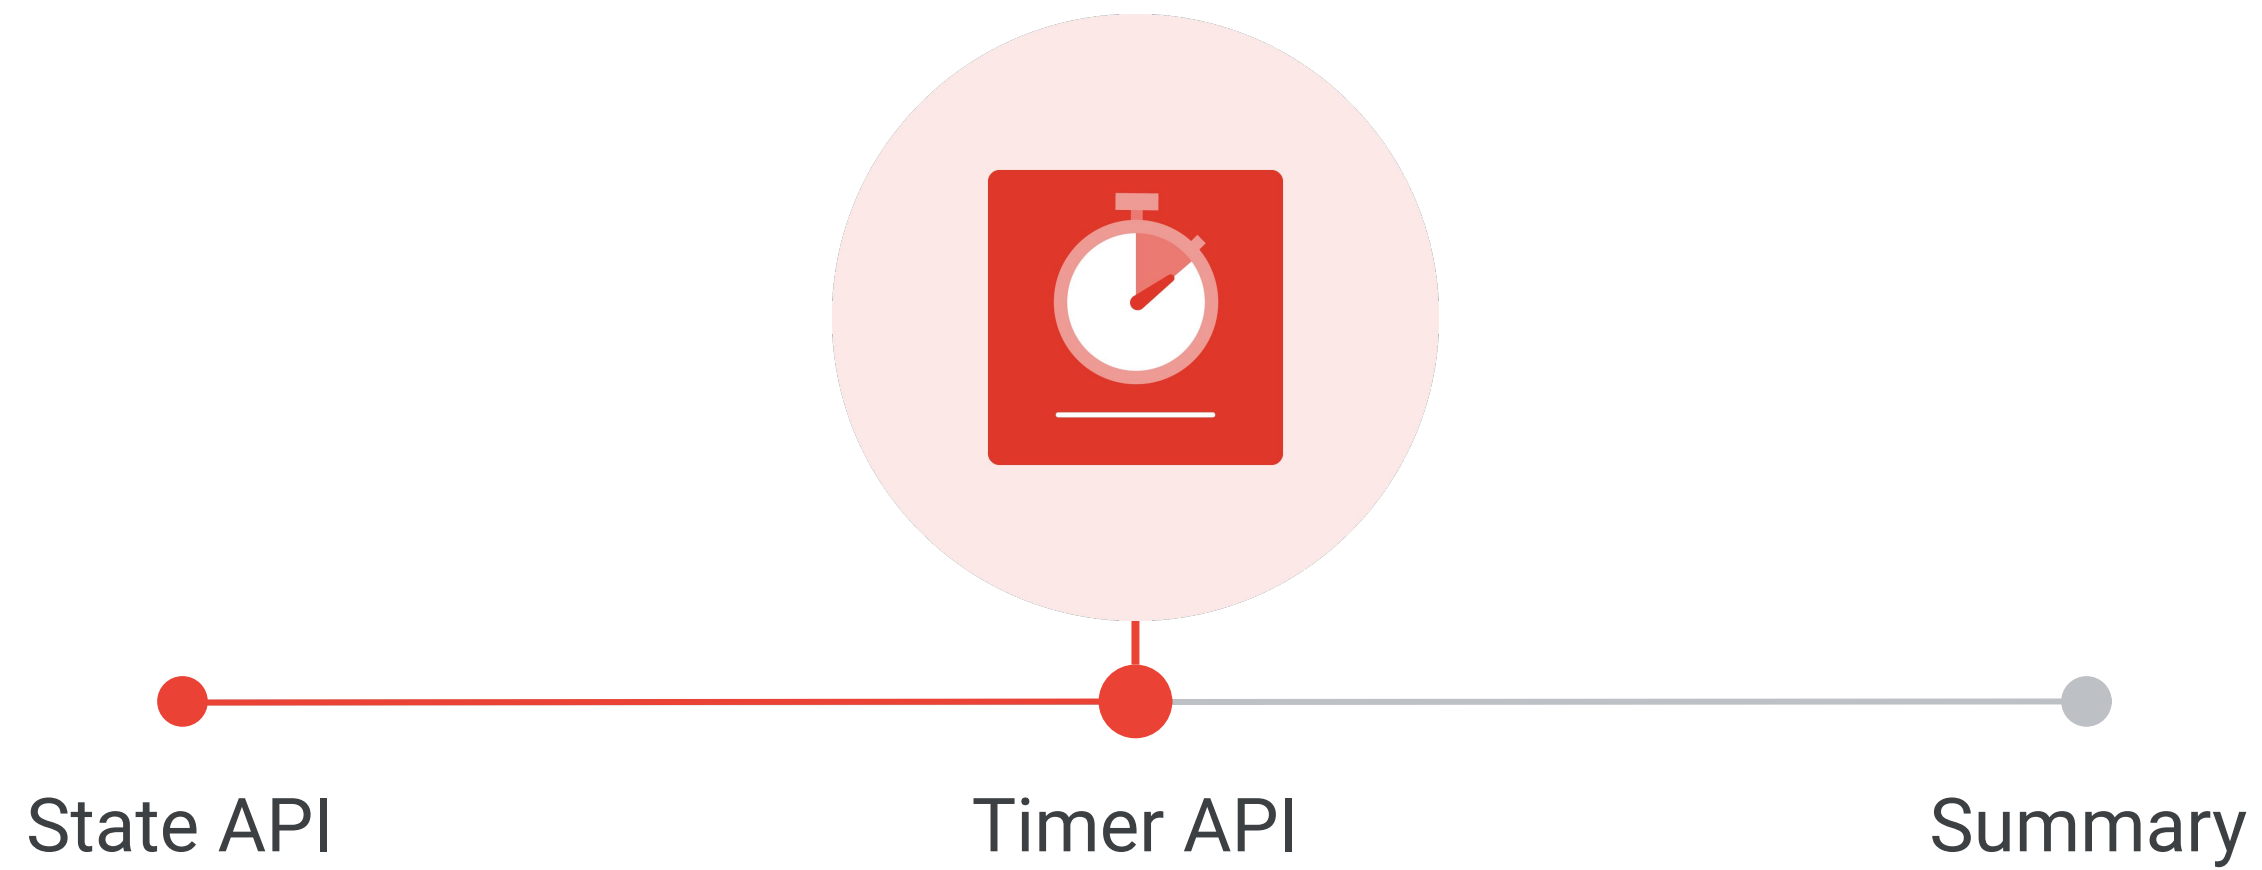
    @ProcessElement
    public void process(ProcessContext context,
        @StateId("buffer") BagState<Event> bufferState,
        @StateId("count") ValueState<Integer> countState) {
        int count = firstNonNull(countState.read(), 0);
        count = count + 1; countState.write(count);
        bufferState.add(context.element());
        if (count >= MAX_BUFFER_SIZE) {
            for (EnrichedEvent enrichedEvent : enrichEvents(bufferState.read())) {
                context.output(enrichedEvent);
            }
            bufferState.clear();
            countState.clear();
        }
    }
    ...
}
```

Increment count and
add element to buffer

When buffer size limit
is reached a request is
sent to the external
service

State and Timers

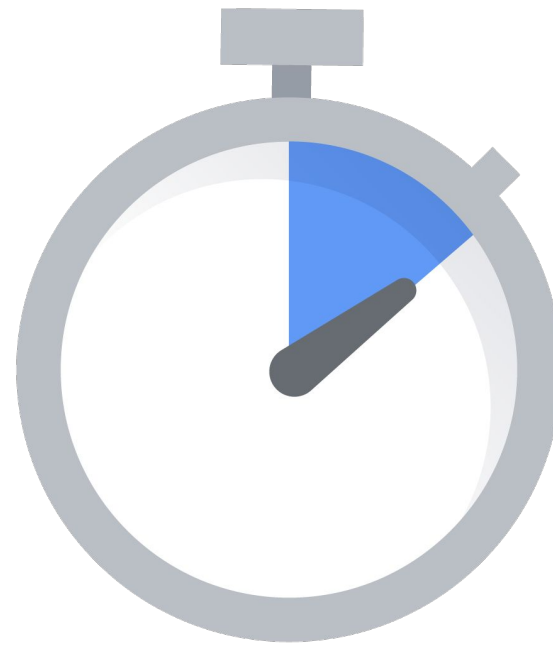
Agenda



Timers

Event-time timers

Callback when the watermark reaches some threshold.



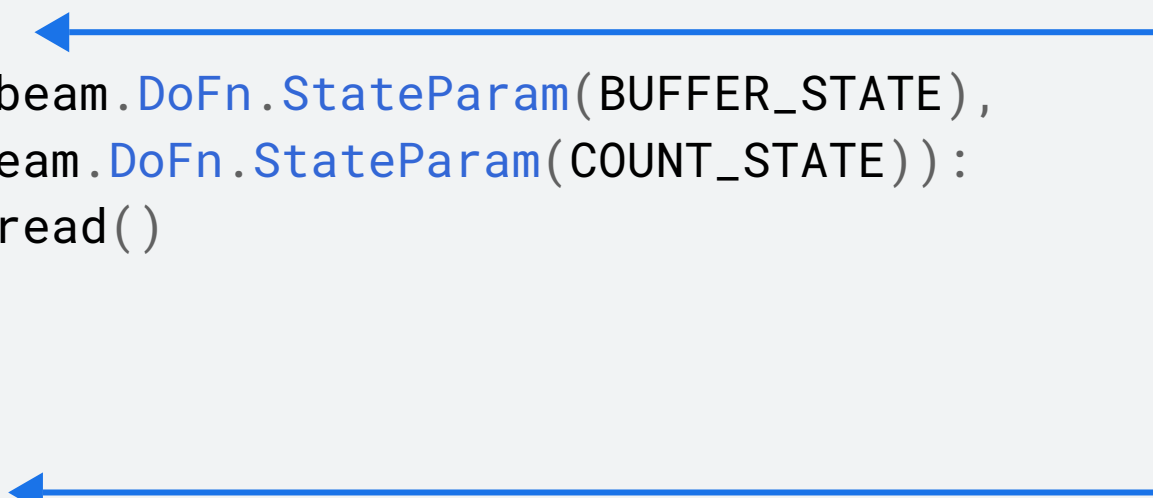
Processing-time timers

Callback after a certain amount of time has elapsed.

Python

```
class StatefulBufferingFn(beam.DoFn):
    ...
    EXPIRY_TIMER = TimerSpec('expiry', TimeDomain.WATERMARK)
    def process(self, element,
                w=beam.DoFn.WindowParam,
                ...
                expiry_timer=beam.DoFn.TimerParam(EXPIRY_TIMER)):
        expiry_timer.set(w.end + ALLOWED_LATENESS)
        ... same logic as on previous code slide...

    @on_timer(EXPIRY_TIMER)
    def expiry(self,
               buffer_state=beam.DoFn.StateParam(BUFFER_STATE),
               count_state=beam.DoFn.StateParam(COUNT_STATE)):
        events = buffer_state.read()
        for event in events:
            yield event
        buffer_state.clear()
        count_state.clear()
```

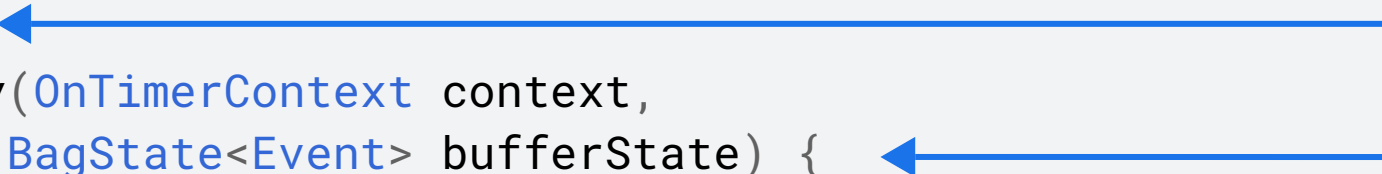


Added an event time timer so that when the window expires, any events remaining in the buffer are processed.

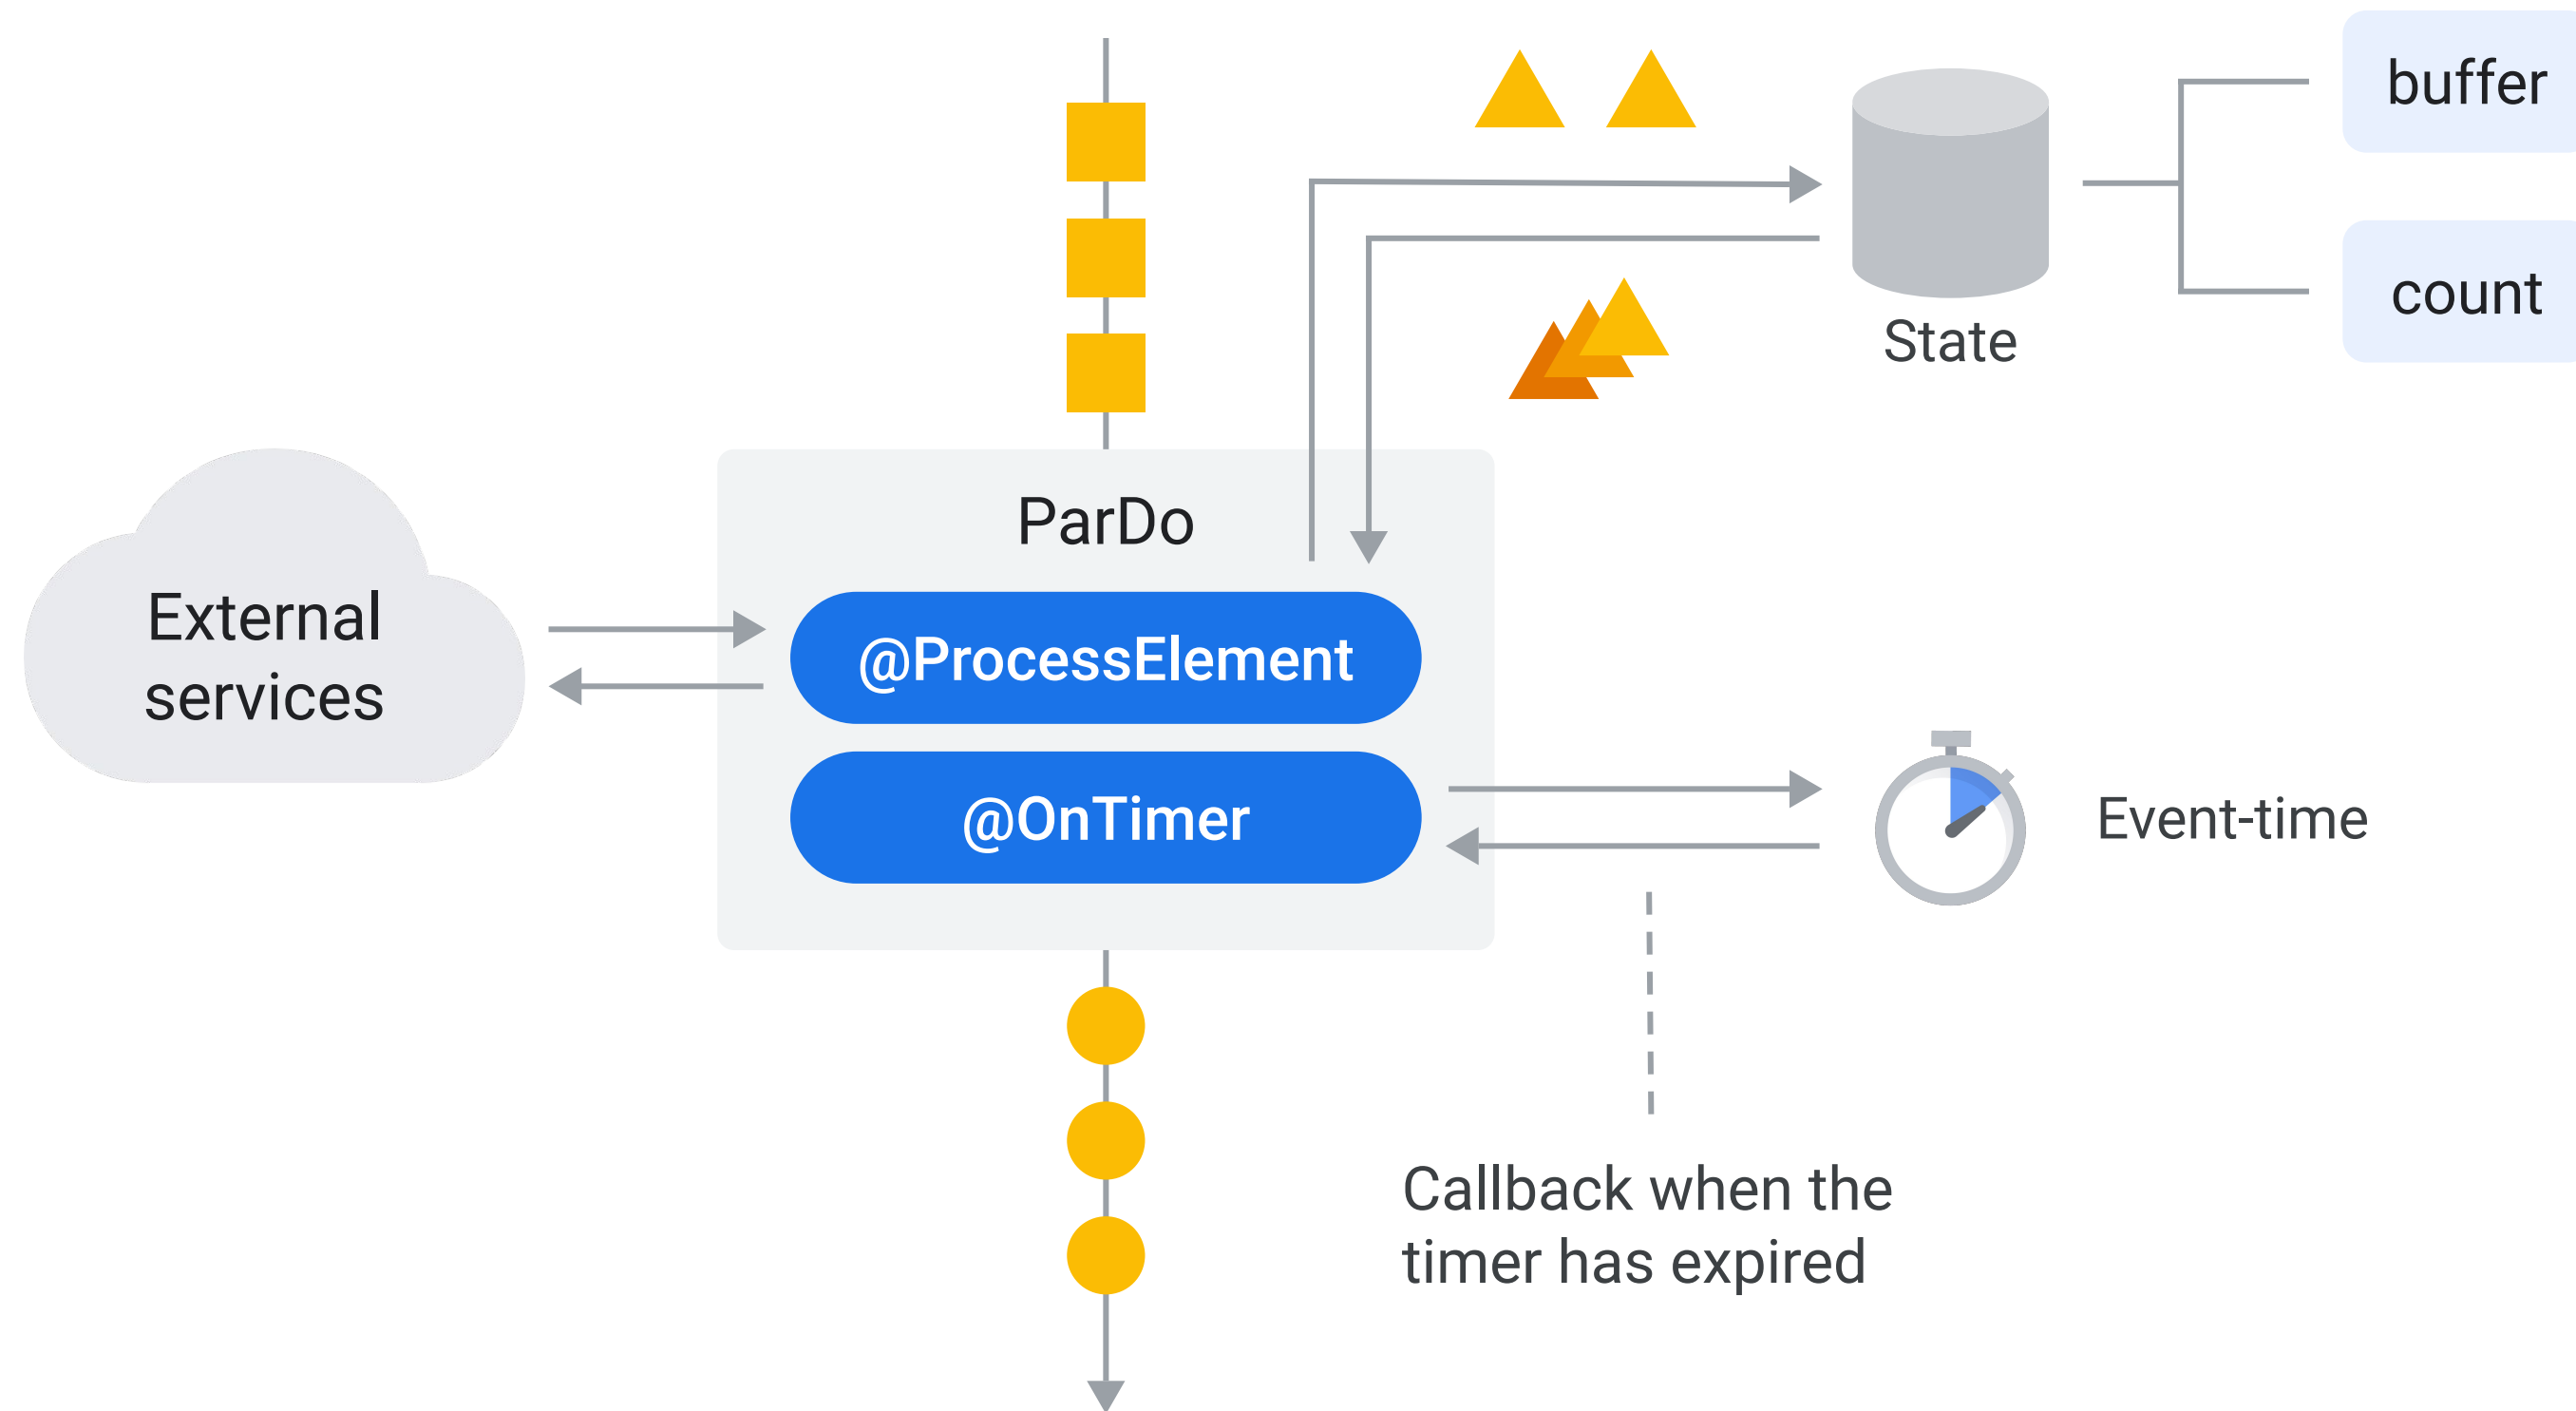
Java

```
new DoFn<Event, EnrichedEvent>() {  
    ...  
    @TimerId("expiry") private final TimerSpec expirySpec = TimerSpecs.timer(TimeDomain.EVENT_TIME);  
    ...  
    @ProcessElement  
    public void process(ProcessContext context, BoundedWindow window, ...  
        @TimerId("expiry") Timer expiryTimer) {  
        expiryTimer.set(window.maxTimestamp().plus(allowedLateness));  
        ... same logic as on previous code slide...  
    }  
  
    @OnTimer("expiry")  
    public void onExpiry(OnTimerContext context,  
        @StateId("buffer") BagState<Event> bufferState) {  
        if (!bufferState.isEmpty().read()) {  
            for (EnrichedEvent enrichedEvent : enrichEvents(bufferState.read())) {  
                context.output(enrichedEvent);  
            }  
            bufferState.clear();  
        }  
    }  
}
```

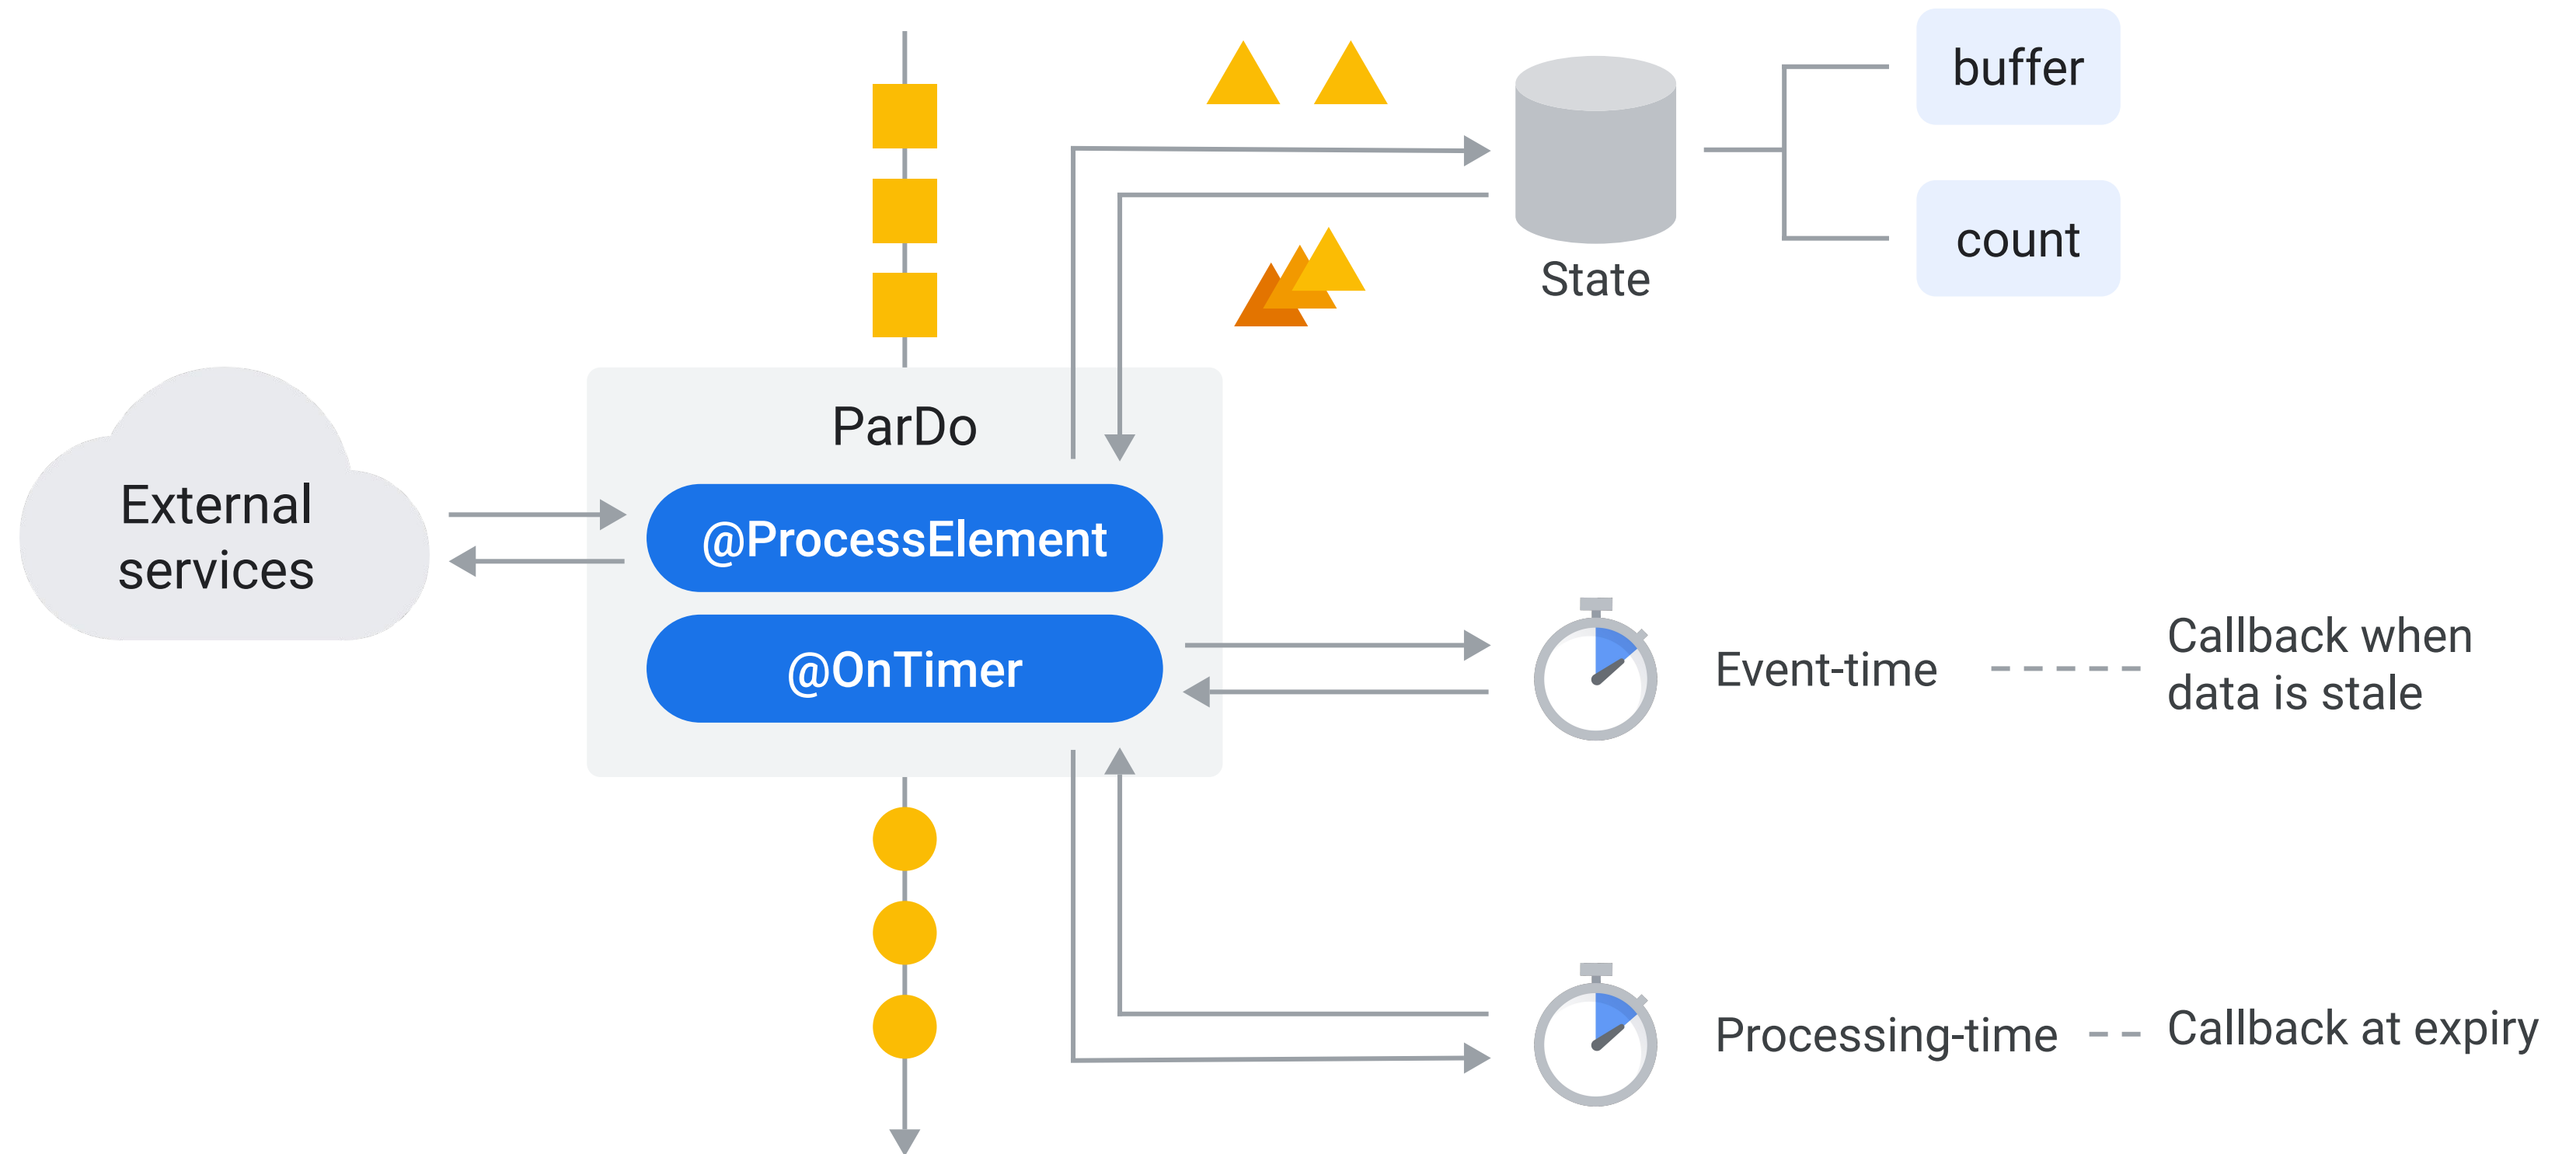
Added an event-time timer so that when the window expires, any events remaining in the buffer are processed.



Stateful DoFn with state and timers

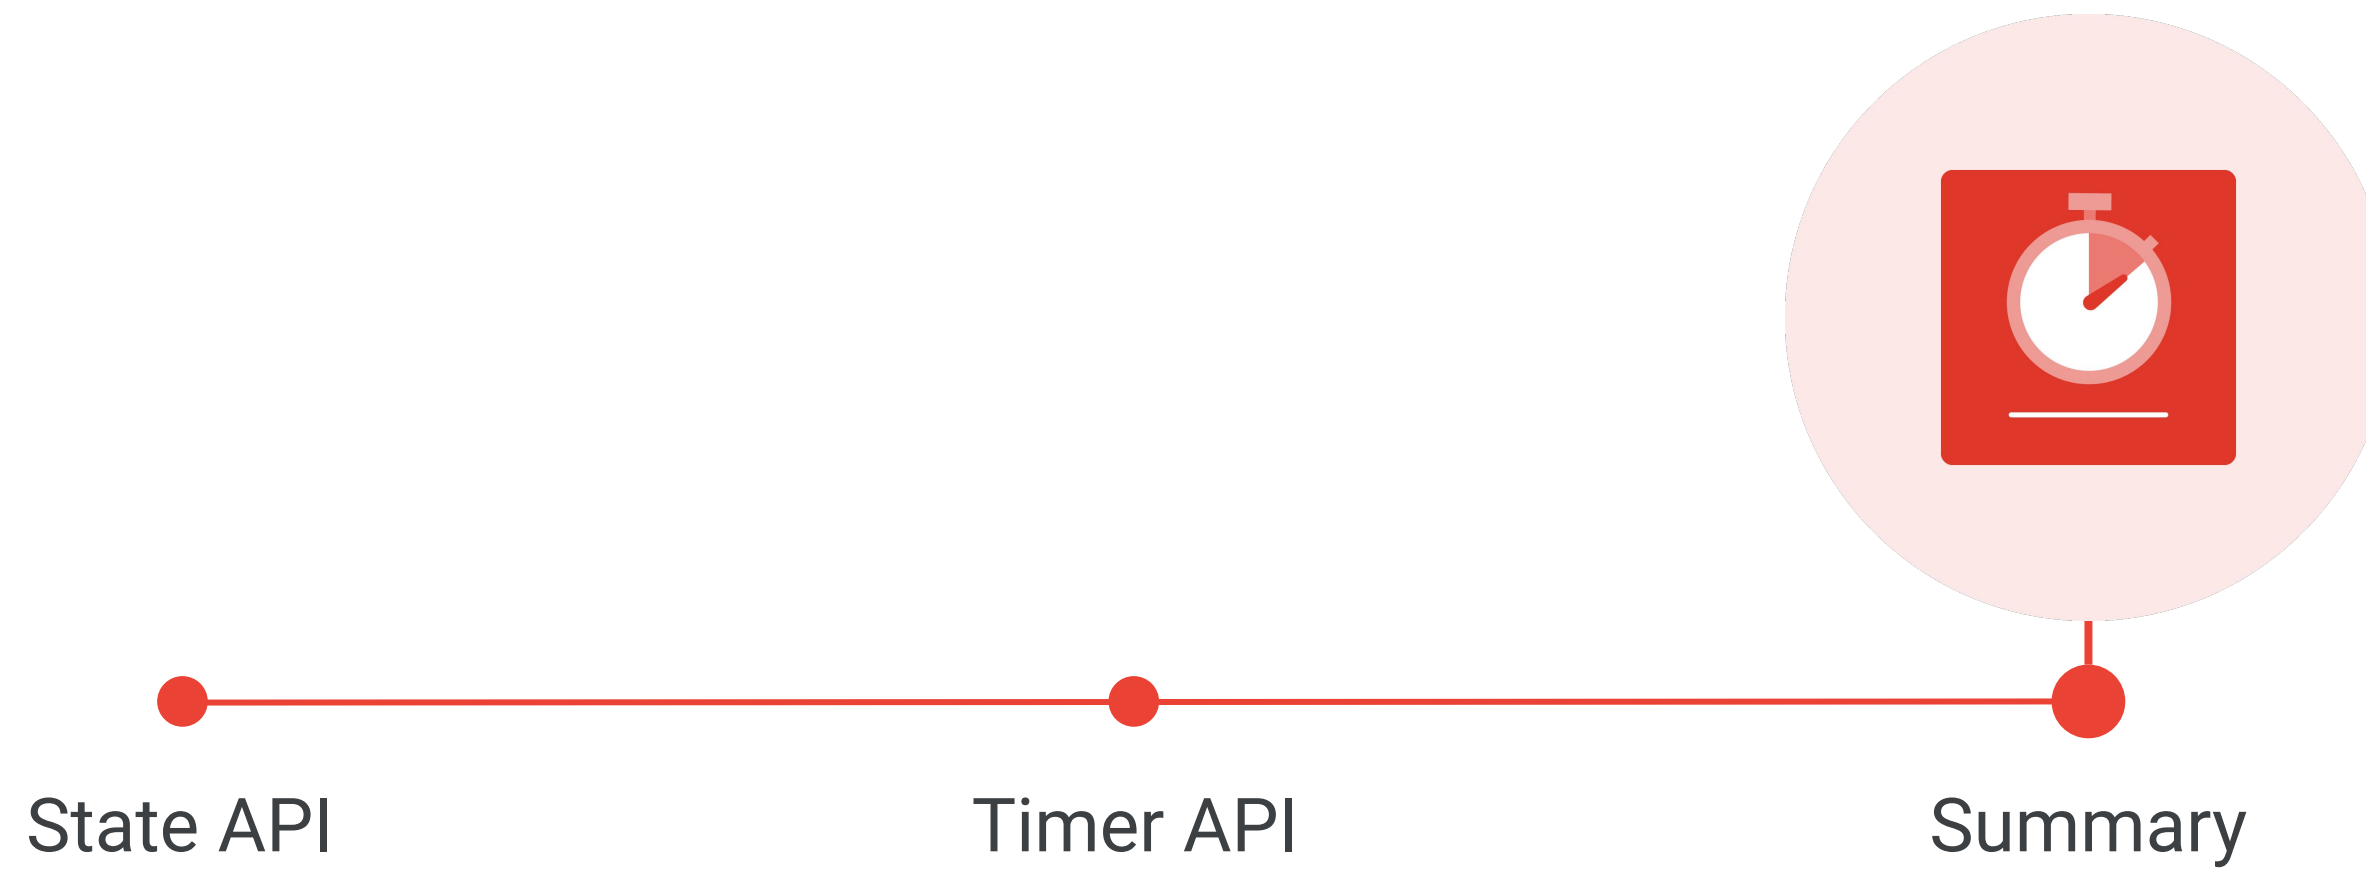


Stateful DoFn with state and timers

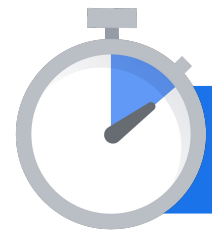


State and Timers

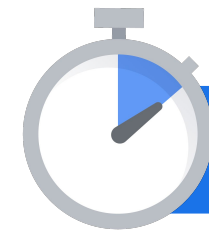
Agenda



Summary: Types of timers

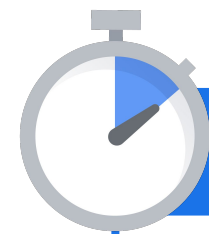


Processing time



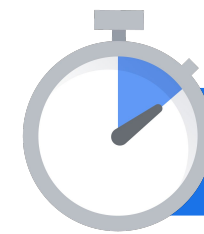
Event time

Summary: Types of timers



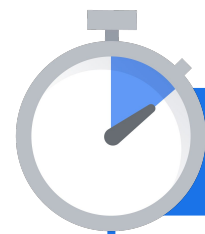
Processing time

- Timeouts
- Relative times ("in 5 minutes")
- Periodic output based on state



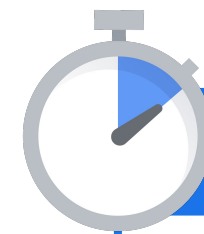
Event time

Summary: Types of timers



Processing time

- Timeouts
- Relative times ("in 5 minutes")
- Periodic output based on state




Event time



- Output based on completeness of input data
- Absolute times ("when the data is complete up to 5:00am")
- Final/authoritative outputs
- Don't leave data behind in state!

Summary: Types of state variables




Summary: Types of state variables

Type	Strength	Dataflow runner
Value	Read/write any value (but always the whole value)	





Summary: Types of state variables

Type	Strength	Dataflow runner
Value	Read/write any value (but always the whole value)	
Bag	Cheap append. No ordering on read	






Summary: Types of state variables

Type	Strength	Dataflow runner
Value	Read/write any value (but always the whole value)	
Bag	Cheap append No ordering on read	
Combining	Associative/commutative compaction	

Summary: Types of state variables

Type	Strength	Dataflow runner
Value	Read/write any value (but always the whole value)	
Bag	Cheap append No ordering on read	
Combining	Associative/commutative compaction	
Map	Read/write just keys you specify	

Summary: Types of state variables

Type	Strength	Dataflow runner
Value	Read/write any value (but always the whole value)	
Bag	Cheap append No ordering on read	
Combining	Associative/commutative compaction	
Map	Read/write just keys you specify	
Set	Membership checking	

What can you do with state and timers?

- **Domain-specific triggering** ("output when five people who live in Seattle have checked in")
- **Slowly changing dimensions** ("update FX rates for currency ABC")
- **Stream joins** ("join-matrix" / "join-biclique")
- **Fine-grained aggregation** ("add odd elements to accumulator A and event elements to accumulator B")
- **Per-key workflows** (like user sign up flow w/ reminders & expiration)

