Beam Notebooks

Reza Rokni

Developer Advocate,
Google Cloud

Agenda

Course Intro

Beam Concepts Review

Windows, Watermarks, and Triggers

Sources and Sinks

Schemas

State and Timers

Best Practices

SQL and DataFrames

Beam Notebooks

Summary

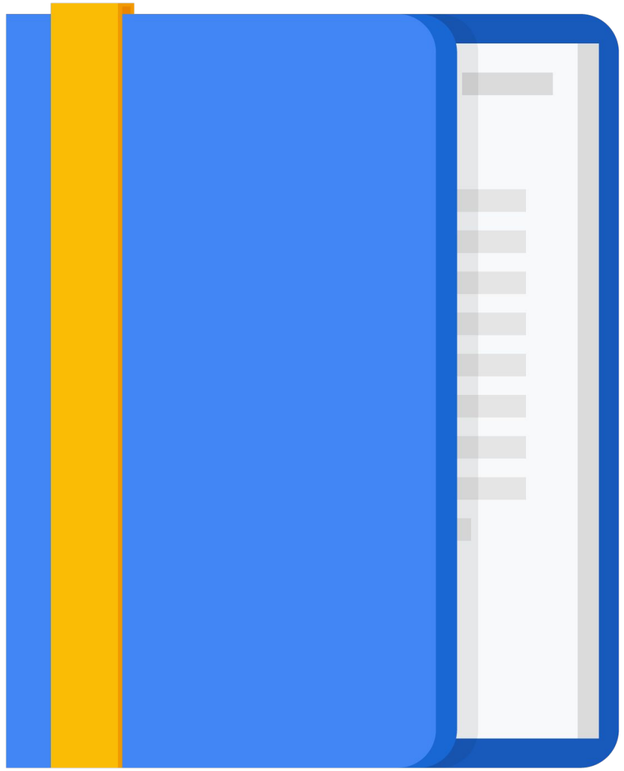

Apache Beam and interactive development

```
storeSales = p  |
beam.io.ReadFromText(“purchases-store”)

                | beam.Map(lambda s: ...)

onlineSales = p |
beam.io.ReadFromText(“purchases-online”)

                | beam.Map(lambda s: ...)

topSales = (storeSales, onlineSales)

                | beam.Flatten()

                | beam.Combiners.Count.perKey()

                | beam.Combiners.Top.of(10, key = lambda
x: x[1])

topSales        | beam.io.WriteToBigQuery(topSales)
```

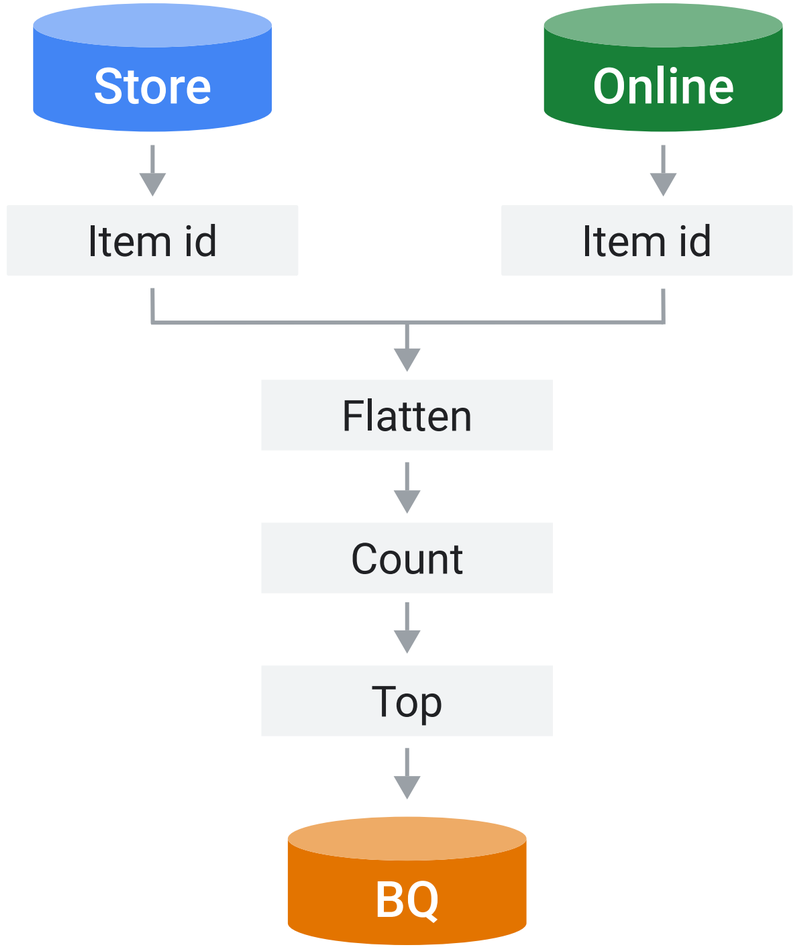

Apache Beam interactive runner

The interactive runner module allows

1. Interactive development

Apache Beam interactive runner

The interactive runner module allows

1. Interactive development
2. Access to intermediate results

Apache Beam interactive runner

The interactive runner module allows

1. Interactive development
2. Access to intermediate results
3. Stream or batch sources

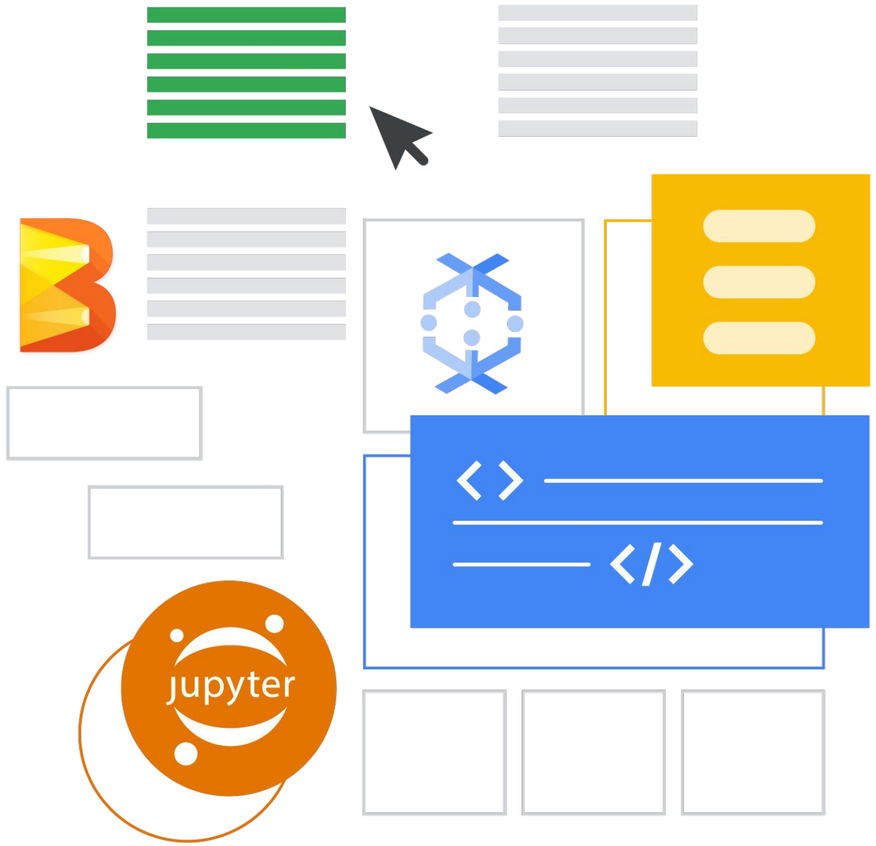

Beam Notebooks

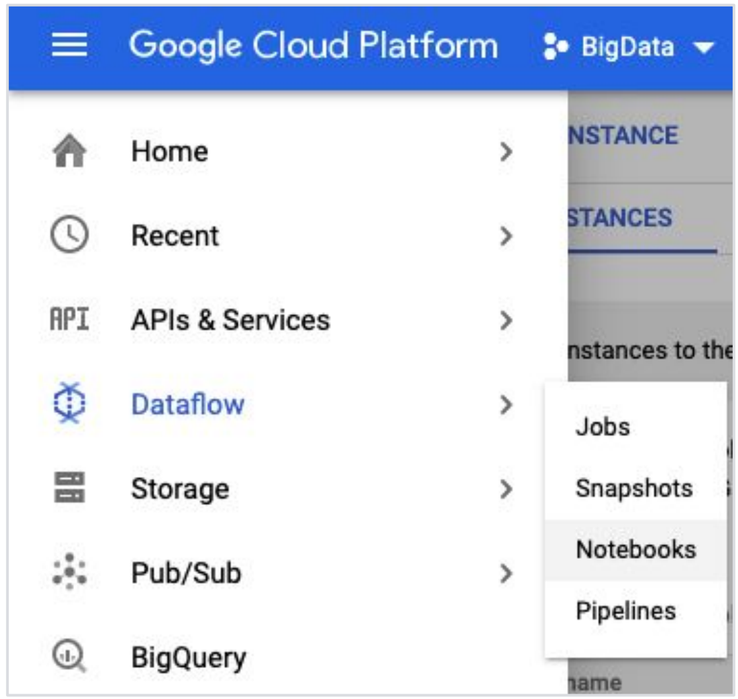

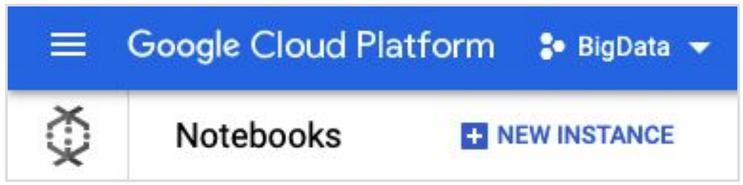

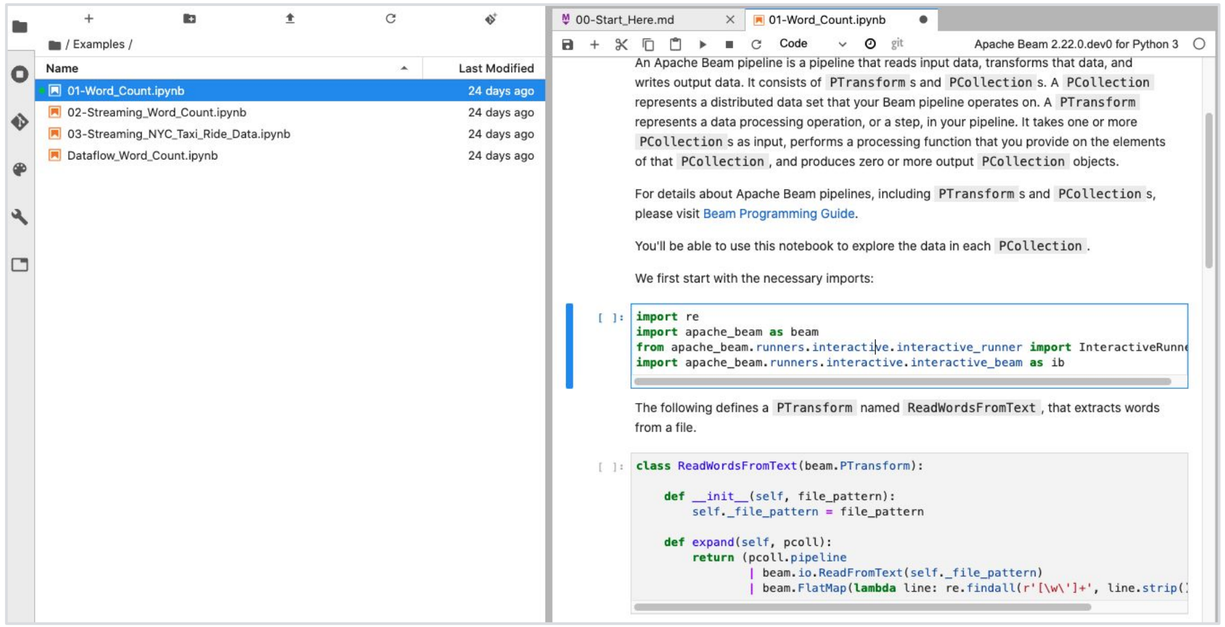

Add a transform

```
words = p | "read" >> beam.io.ReadFromPubSub(topic=topic)

windowed_words = (words
    | "window" >>
beam.WindowInto(beam.window.FixedWindows(10)))

windowed_words_counts = (windowed_words
    | "count" >> beam.combiners.Count.PerElement())
```

Set interactivity options before we run the cell

ib.options.recording_duration

Sets the amount of time the InteractiveRunner records data from an unbounded source

```
# Set the recording duration to 10 min
ib.options.recording_duration = '10m'
```

ib.options.recording_size_limit

Sets the amount of data the InteractiveRunner records (in bytes) from an unbounded source

```
# Set the recording size limit to 1 GB
ib.options.recording_size_limit = 1e9
```

Access Transform output

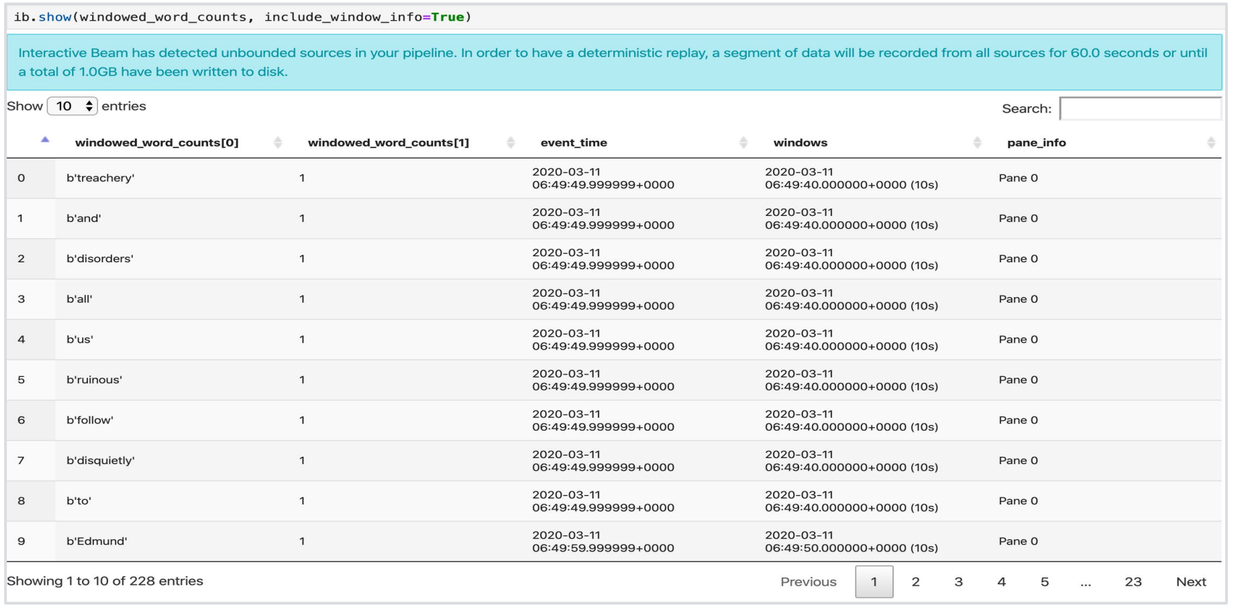

```
ib.show(windowed_word_counts, include_window_info=True)
```

Interactive Beam has detected unbounded sources in your pipeline. In order to have a deterministic replay, a segment of data will be recorded from all sources for 60.0 seconds or until a total of 1.0GB have been written to disk.

Show 10 entries Search:

	windowed_word_counts[0]	windowed_word_counts[1]	event_time	windows	pane_info
0	b'treachery'	1	2020-03-11 06:49:49.999999+0000	2020-03-11 06:49:40.000000+0000 (10s)	Pane 0
1	b'and'	1	2020-03-11 06:49:49.999999+0000	2020-03-11 06:49:40.000000+0000 (10s)	Pane 0
2	b'disorders'	1	2020-03-11 06:49:49.999999+0000	2020-03-11 06:49:40.000000+0000 (10s)	Pane 0
3	b'all'	1	2020-03-11 06:49:49.999999+0000	2020-03-11 06:49:40.000000+0000 (10s)	Pane 0
4	b'us'	1	2020-03-11 06:49:49.999999+0000	2020-03-11 06:49:40.000000+0000 (10s)	Pane 0
5	b'ruinous'	1	2020-03-11 06:49:49.999999+0000	2020-03-11 06:49:40.000000+0000 (10s)	Pane 0
6	b'follow'	1	2020-03-11 06:49:49.999999+0000	2020-03-11 06:49:40.000000+0000 (10s)	Pane 0
7	b'disquietly'	1	2020-03-11 06:49:49.999999+0000	2020-03-11 06:49:40.000000+0000 (10s)	Pane 0
8	b'to'	1	2020-03-11 06:49:49.999999+0000	2020-03-11 06:49:40.000000+0000 (10s)	Pane 0
9	b'Edmund'	1	2020-03-11 06:49:59.999999+0000	2020-03-11 06:49:50.000000+0000 (10s)	Pane 0

Showing 1 to 10 of 228 entries Previous 1 2 3 4 5 ... 23 Next

```
# Materializes the resulting PCollection in a table
ib.show(windowed_word_counts, include_window_info=True)

# Load the output in a Pandas DataFrame
ib.collect(windowed_word_counts, include_window_info=True)
```

Visualize Transform output

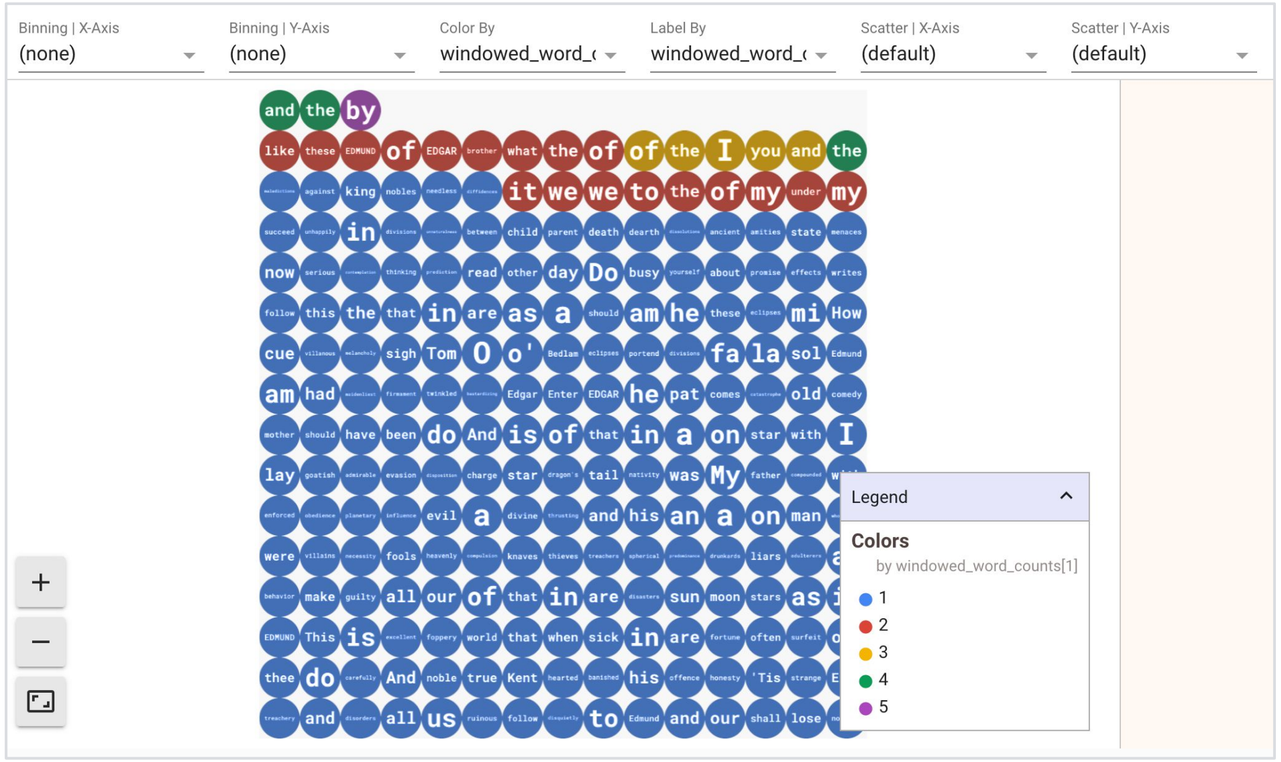


```
# Visualize the data in the notebook
ib.show(windowed_word_counts, include_window_info=True, visualize_data=True)
```

Going from development to production

```
# Import the production Dataflow runner
from apache_beam.runners import DataflowRunner

# Set up Apache Beam pipeline options
options = pipeline_options.PipelineOptions()

# Run the pipeline
runner = DataflowRunner()
runner.run_pipeline(p, options=options)
```